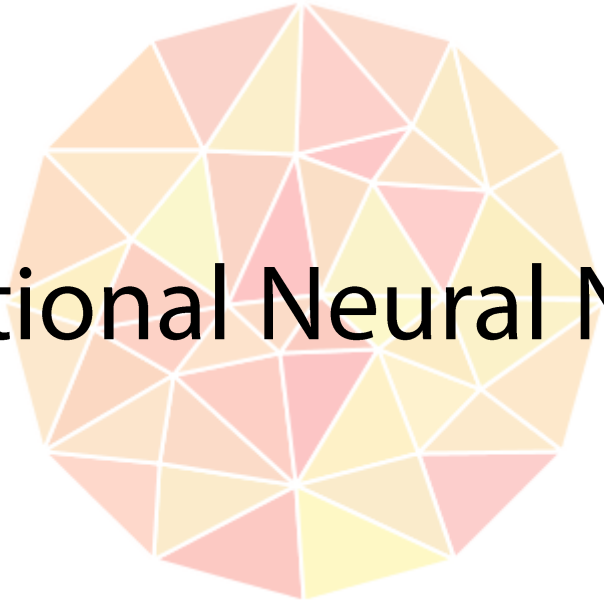

Convolutional Neural Networks

Stephen Baek

Fully Connected Layer

- 28 x 28 image → stretch to 784 x 1
- 64 x 64 x 3 image → stretch to 12288 x 1
- …

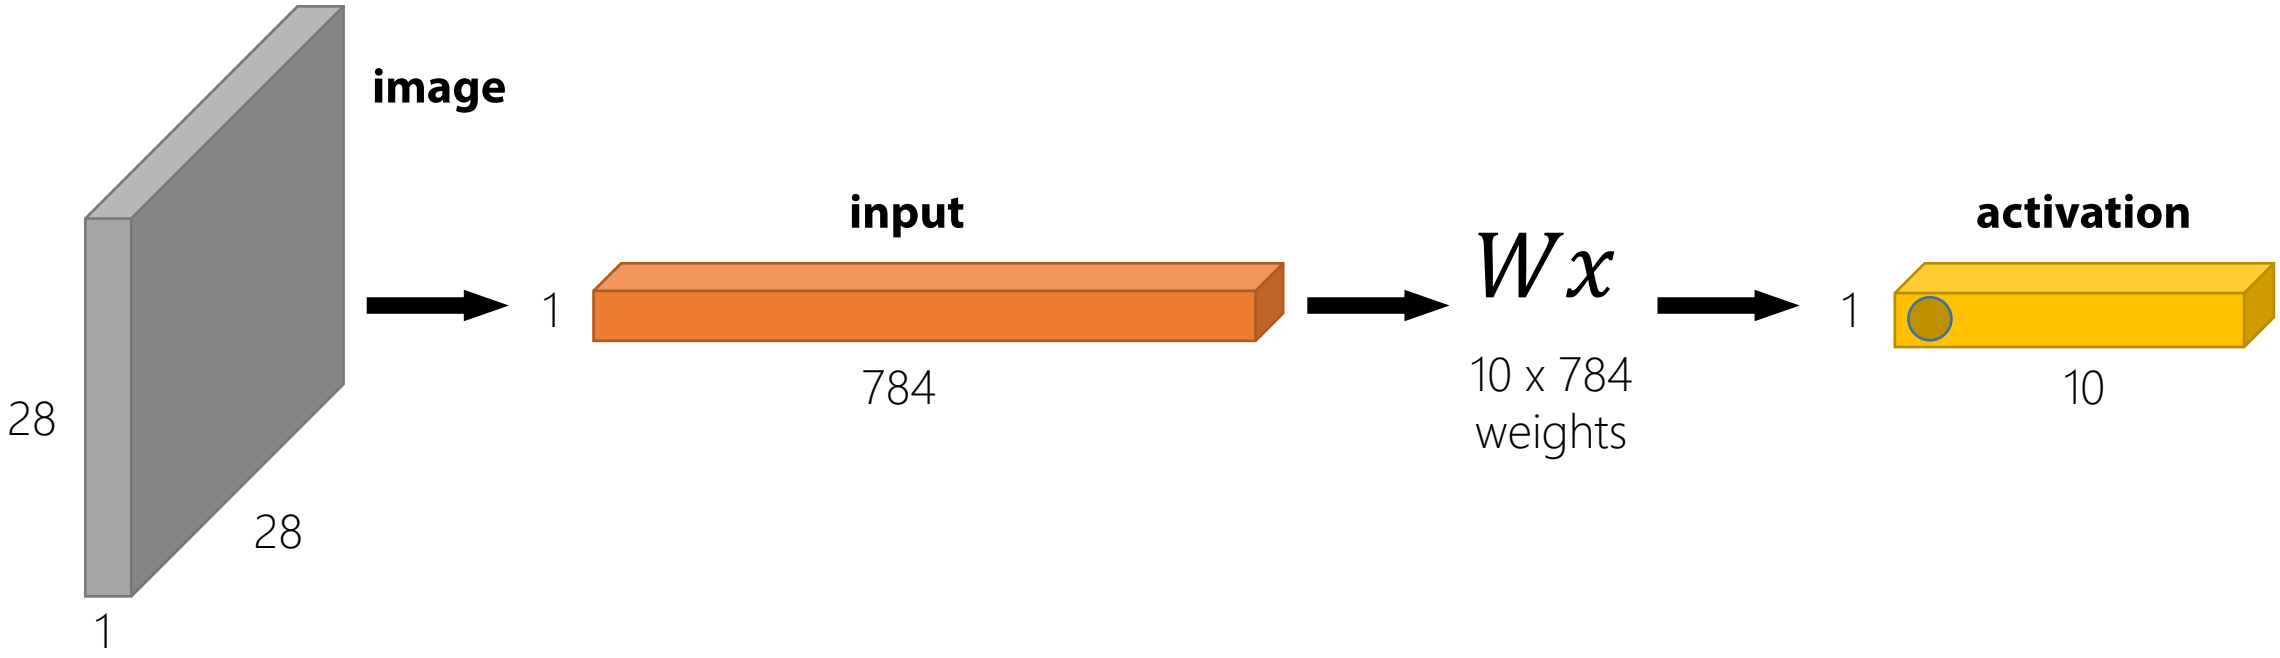

Fully Connected Layer

- 28 x 28 image → stretch to 784 x 1
- 64 x 64 x 3 image → stretch to 12288 x 1
- …

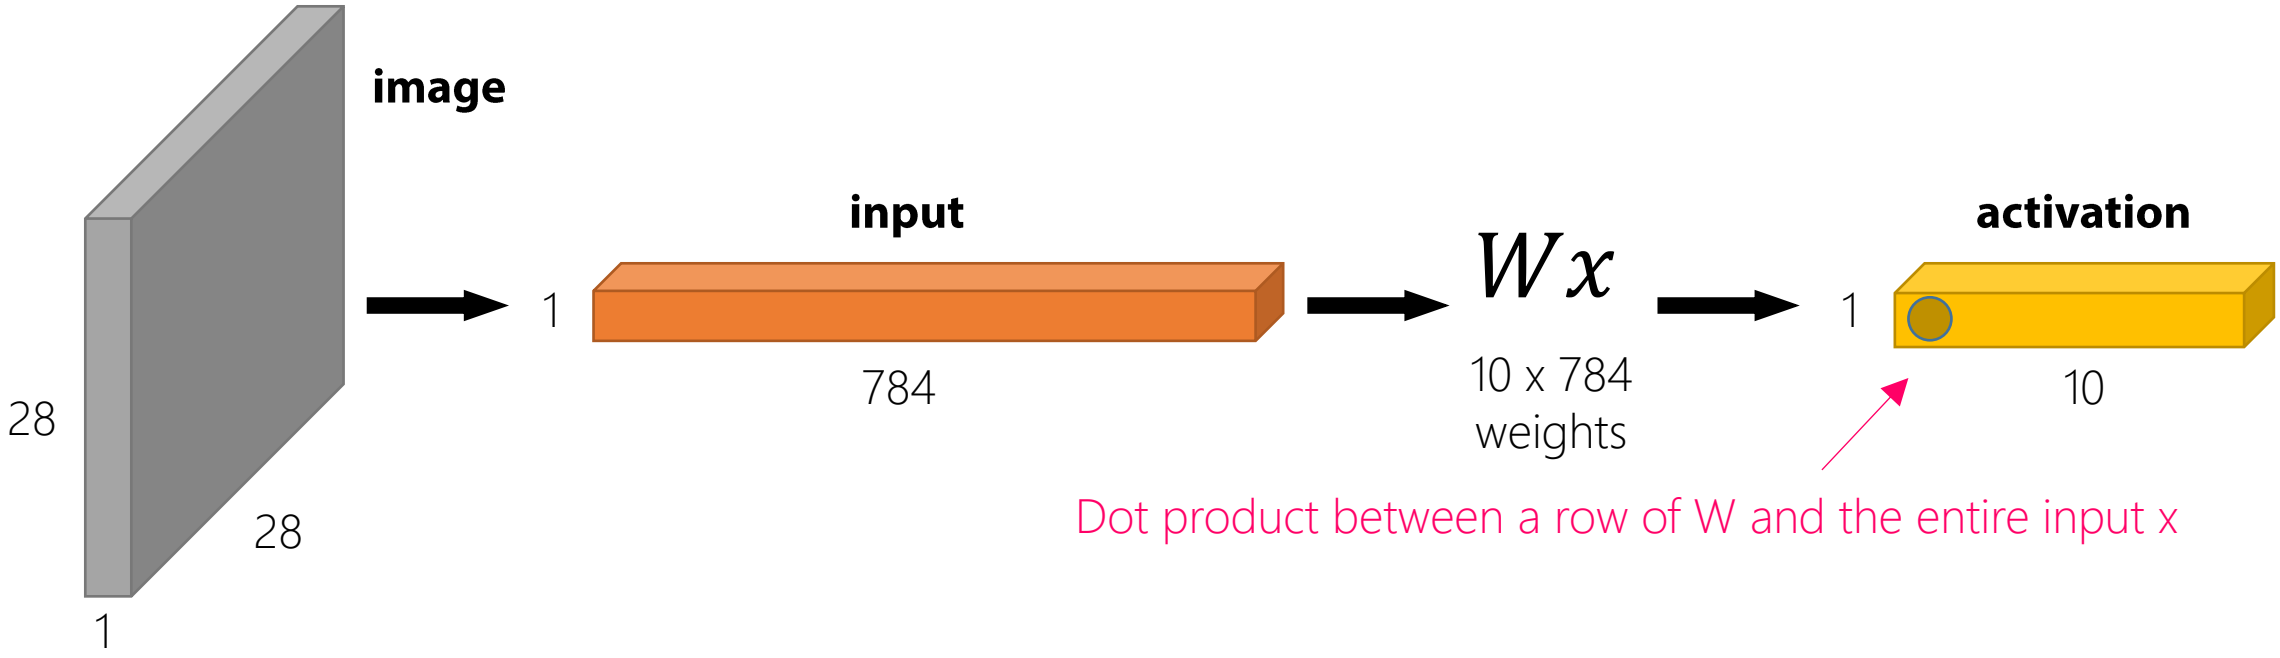

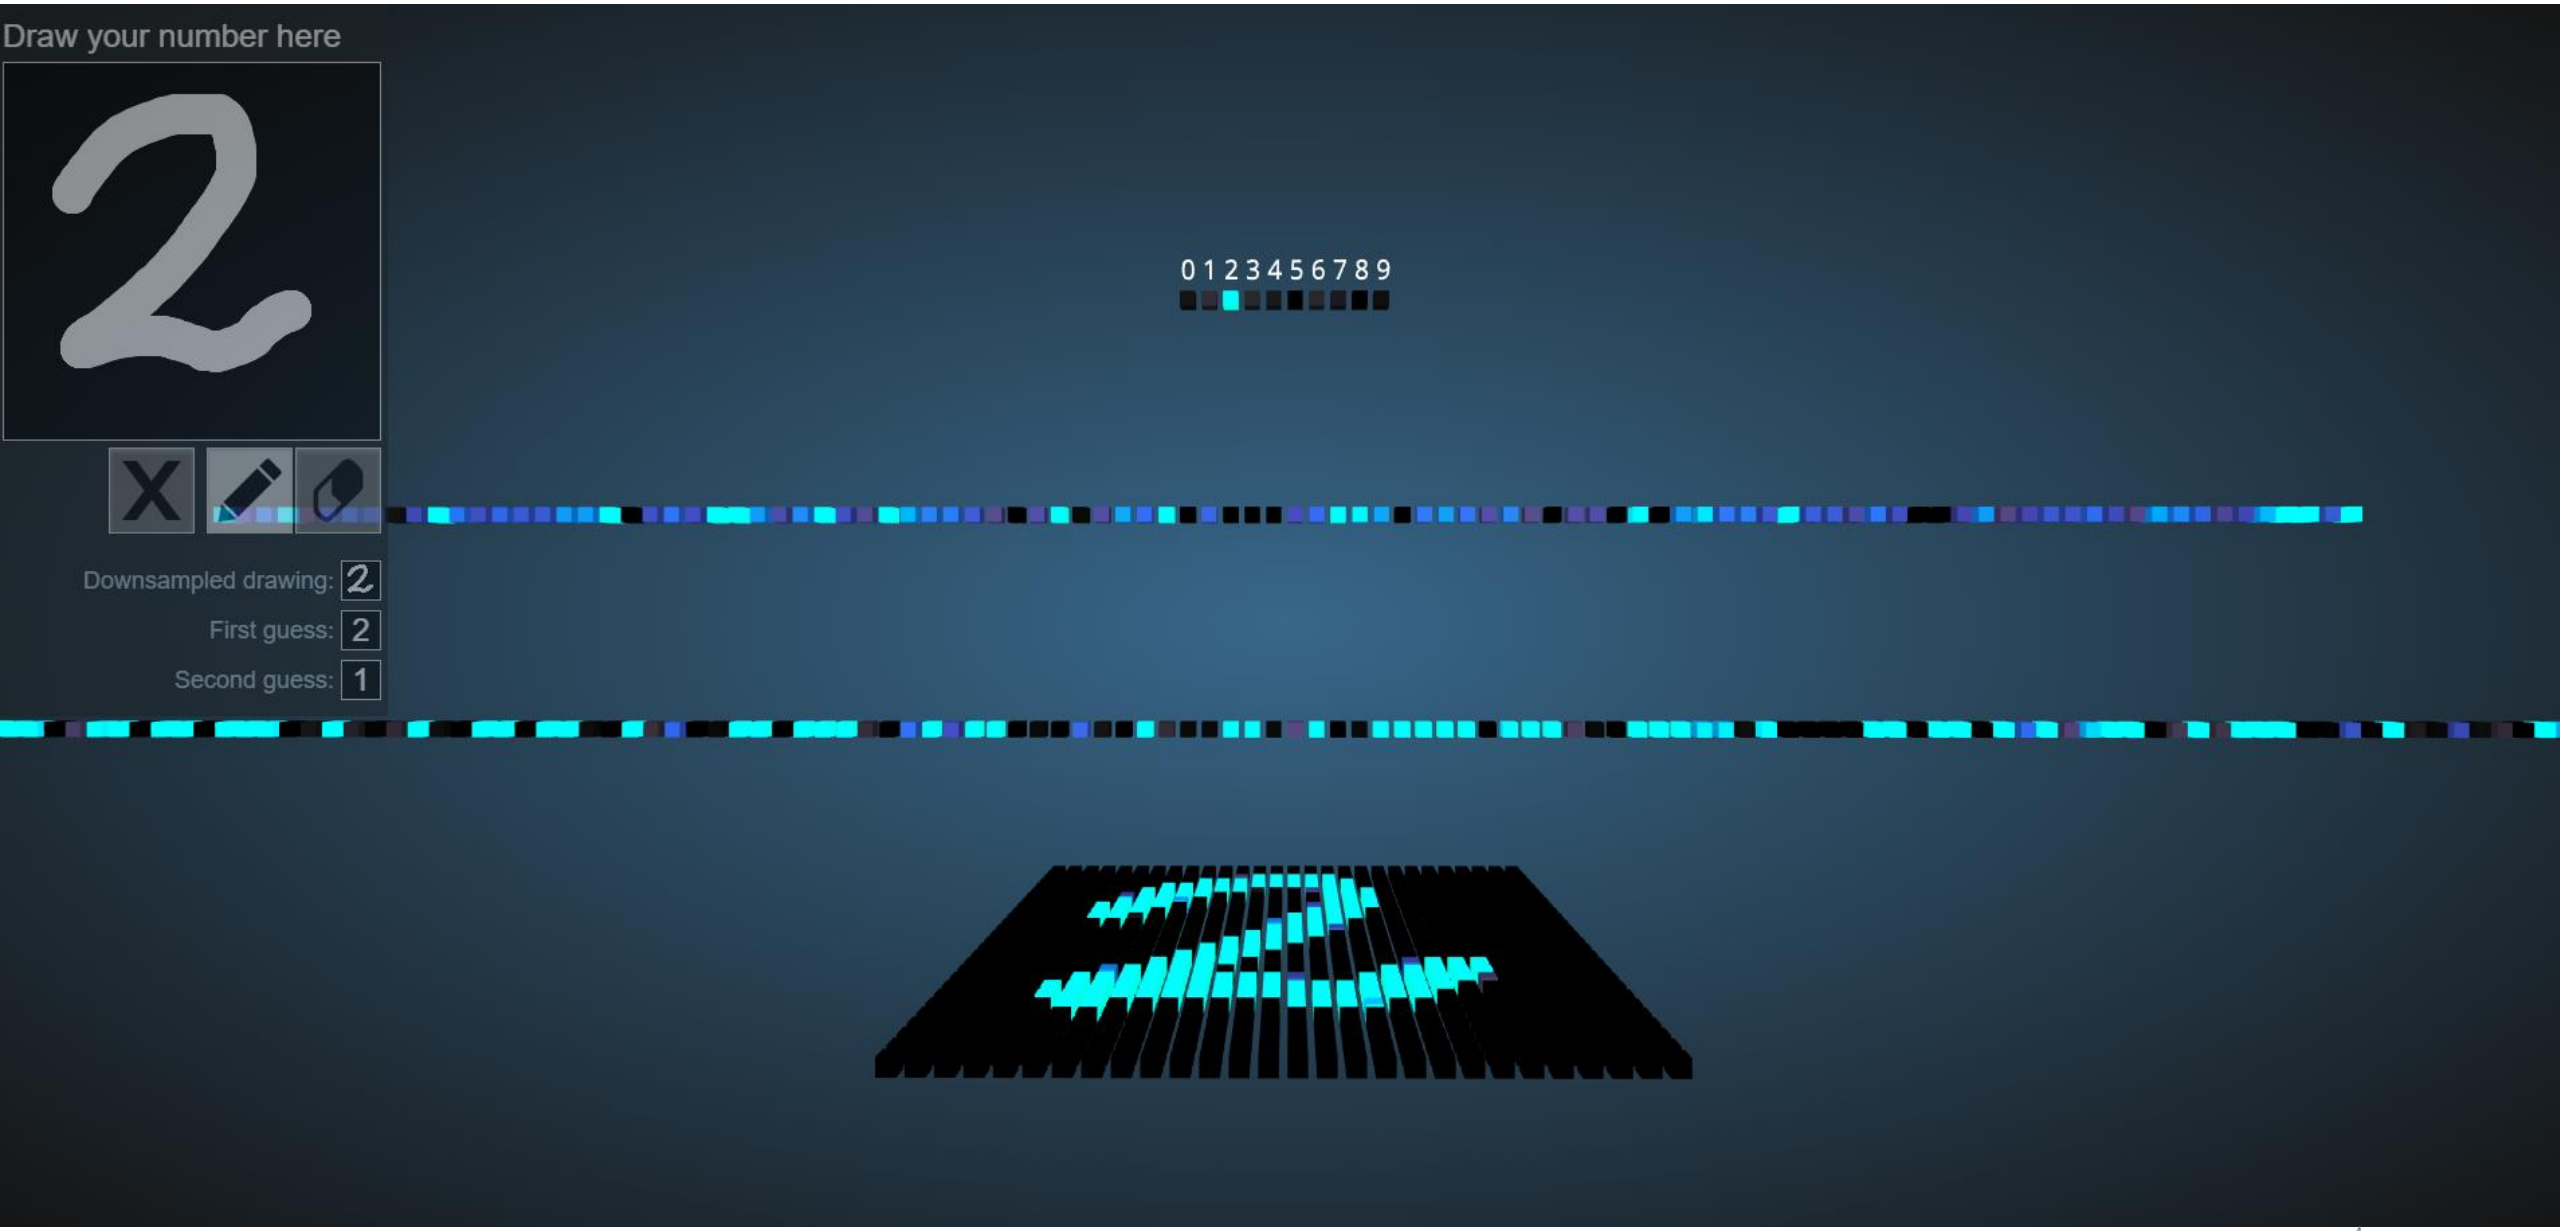

http://scs.ryerson.ca/~aharley/vis/fc/

Hubel & Wiesel (1959 ~)

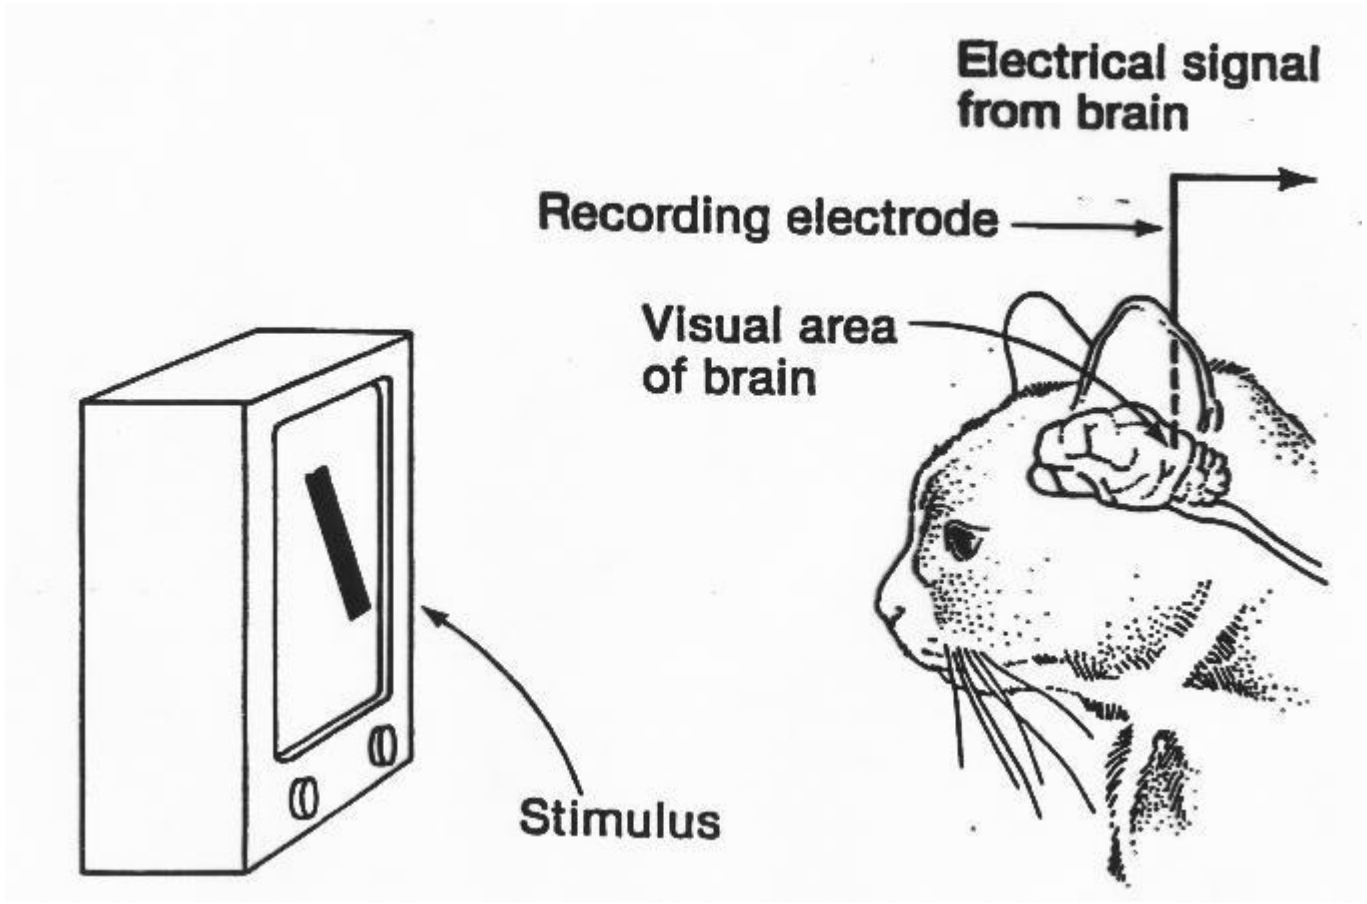

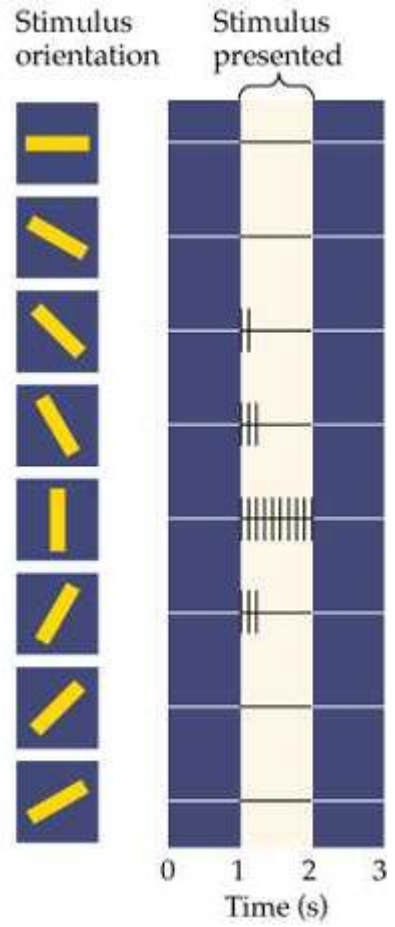

Neurons in the visual cortex respond selectively to oriented edges. Neurons in visual cortex typically respond vigorously to a bar of light oriented at a particular angle and weakly or not at all to other orientations.

Hubel & Wiesel (1959 ~)

Warning!! Visually Disturbing

https://www.youtube.com/watch?v=IOHayh06LJ4

Topographic Maps in Cortex

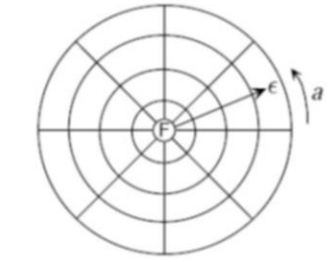

(Dayan and Abbott 2001)

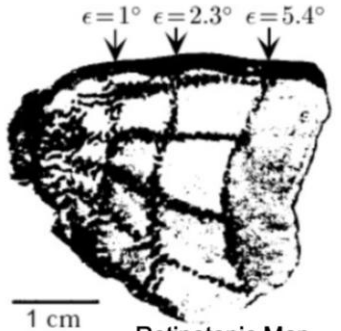

Retinotopic Map

- Each visual sensitive cell only responses to stimuli of a limited region (receptive field)
- Neighboring cells have partially overlapping receptive fields
- Neighboring points in a visual image evoke activity in neighboring regions of visual cortex
- In this manner, the visual system easily maintain the information of the spatial location of stimulus

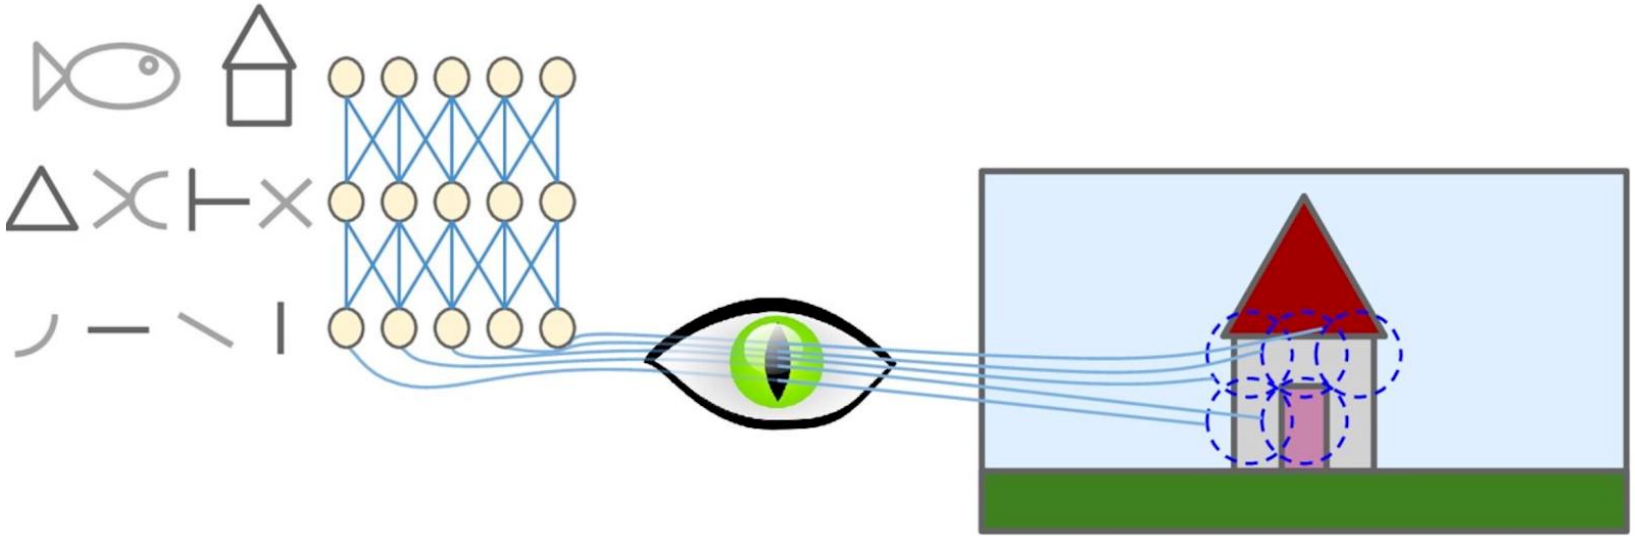

The human visual system

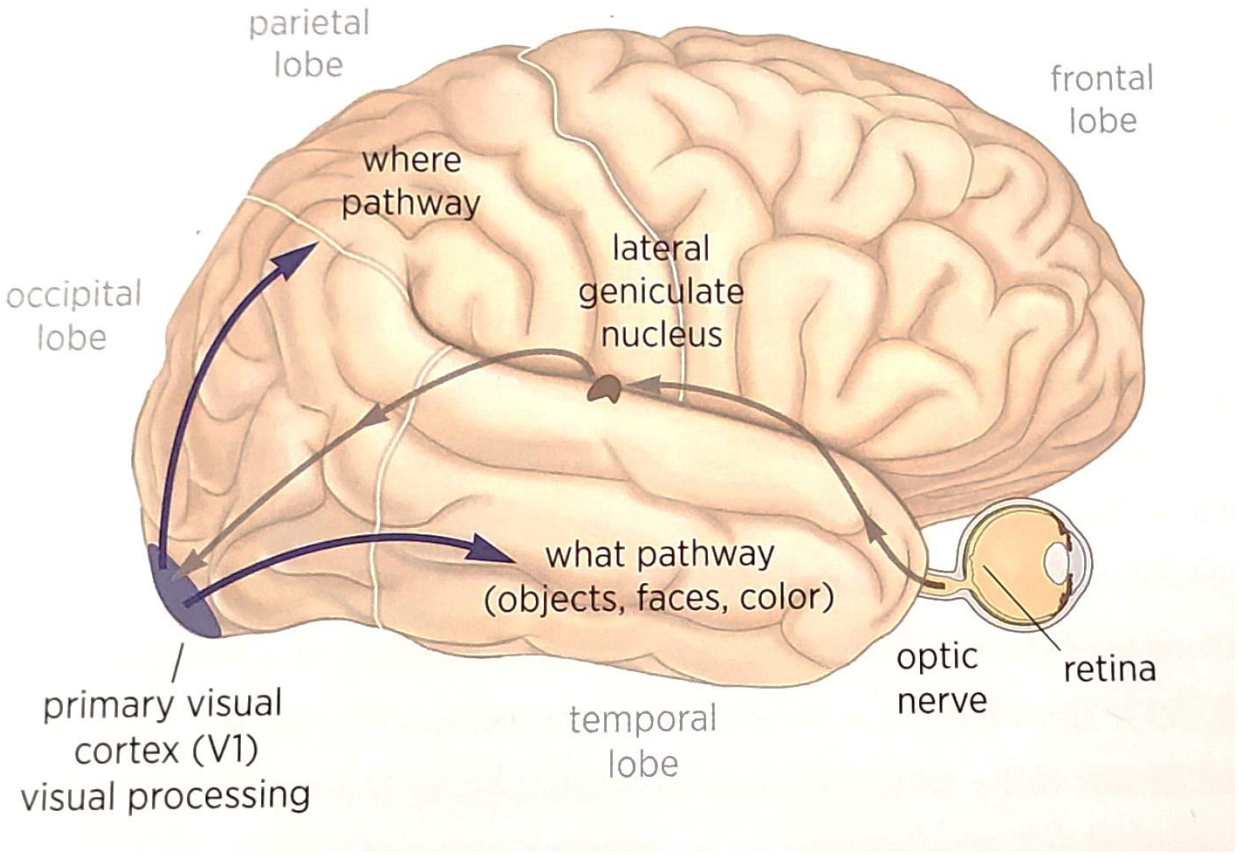

Image Courtesy: Kandel, "Reductionism in Art and Brain Science," 2016

- Retina: visual input
- Retina → Lateral Geniculate Nucleus
  - Visual information flows through the optic nerve
- Lateral Geniculate Nucleus (LGN):
  - A small, ovoid object at the end of the optic tract
  - One on each side of the brain
  - In humans, each LGN has six layers of neurons
  - Sends information to the primary visual cortex (V1)

e.g. Facial Recognition

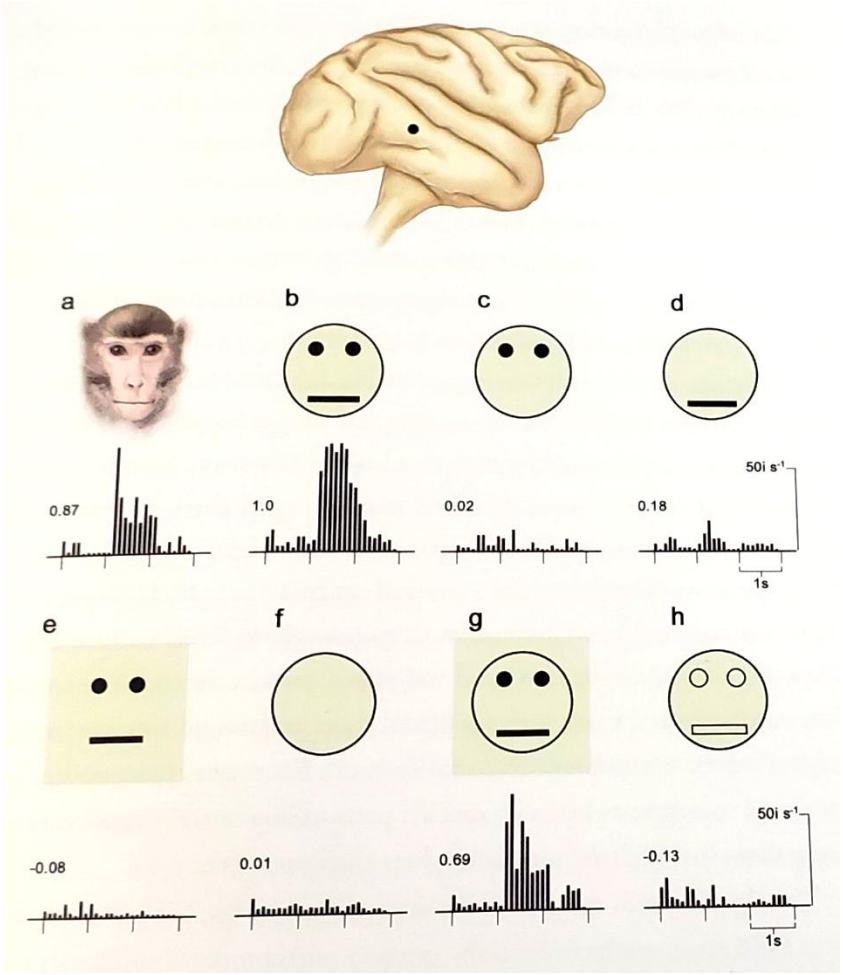

- Holistic face detection
  - "Face cell" in the inferior temporal cortex
  - Fires when there is a face-like object

Image Courtesy: Kandel, "Reductionism in Art and Brain Science," 2016

Convolutional Neural Networks

- Key idea:
  - As like how we (humans) understand a visual scene, if neural nets could see small pieces, understand patterns and textures, combine the pieces to see a bigger picture, computers should be able to recognize images.

- A bonus:
  - Typical neural networks are “fully connected”.
  - In an image domain, this means all the pixels are interconnected.
  - However, pixels far apart have no significant meaning…
  - By connecting only the nearing neighbors, computational load could be much lower.

Convolution Layer

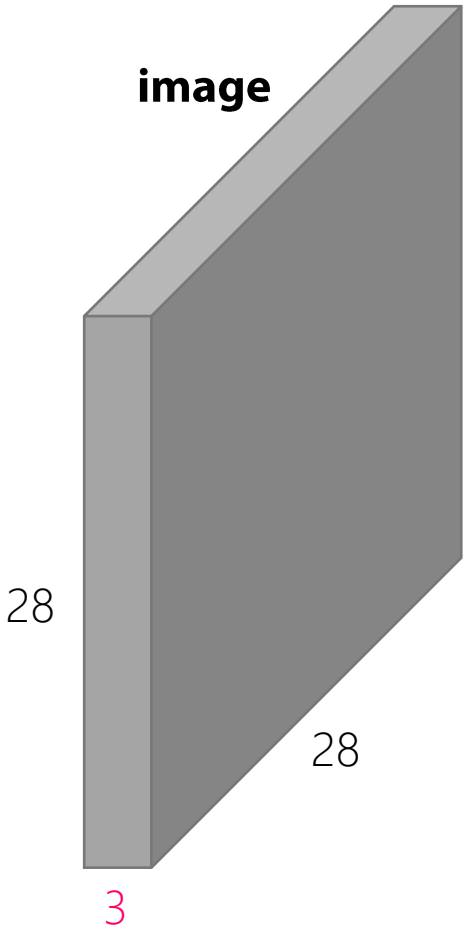

Convolution Layer

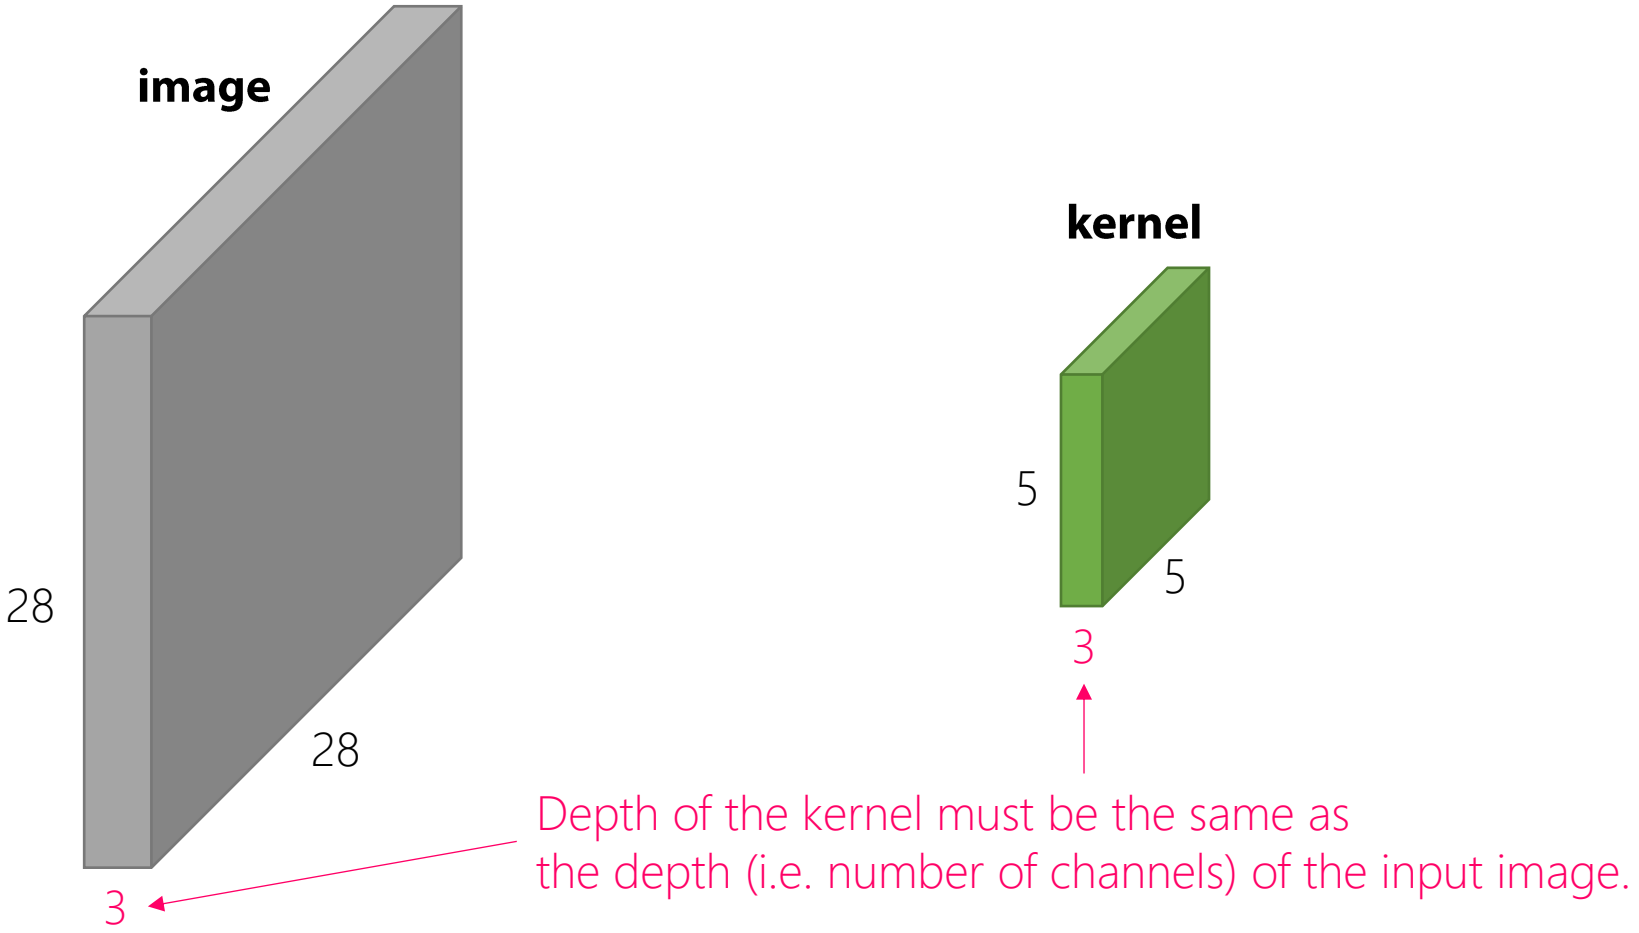

Convolution Layer

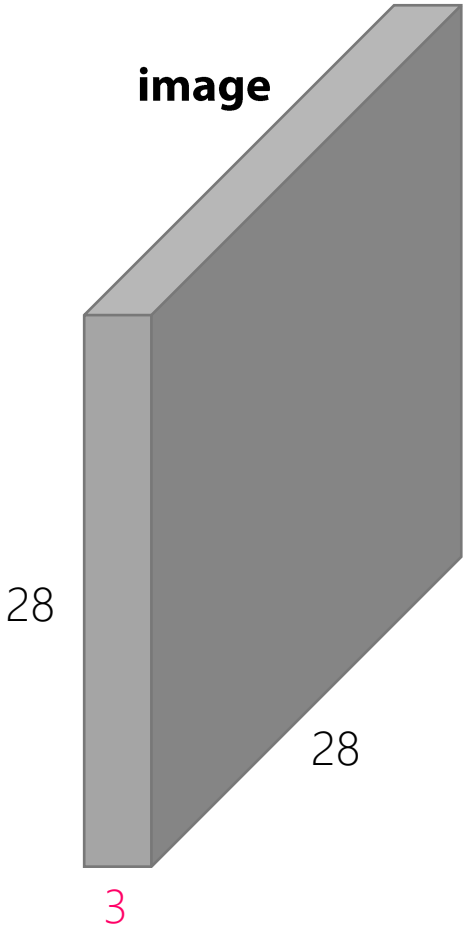

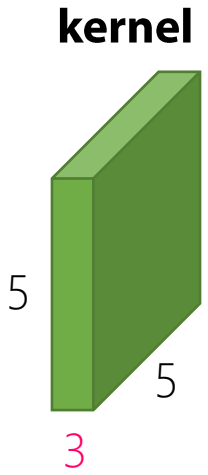

Convolve the kernel with the image!
In other words…
"slide the kernel over the image, compute dot products each time."

Convolution Layer

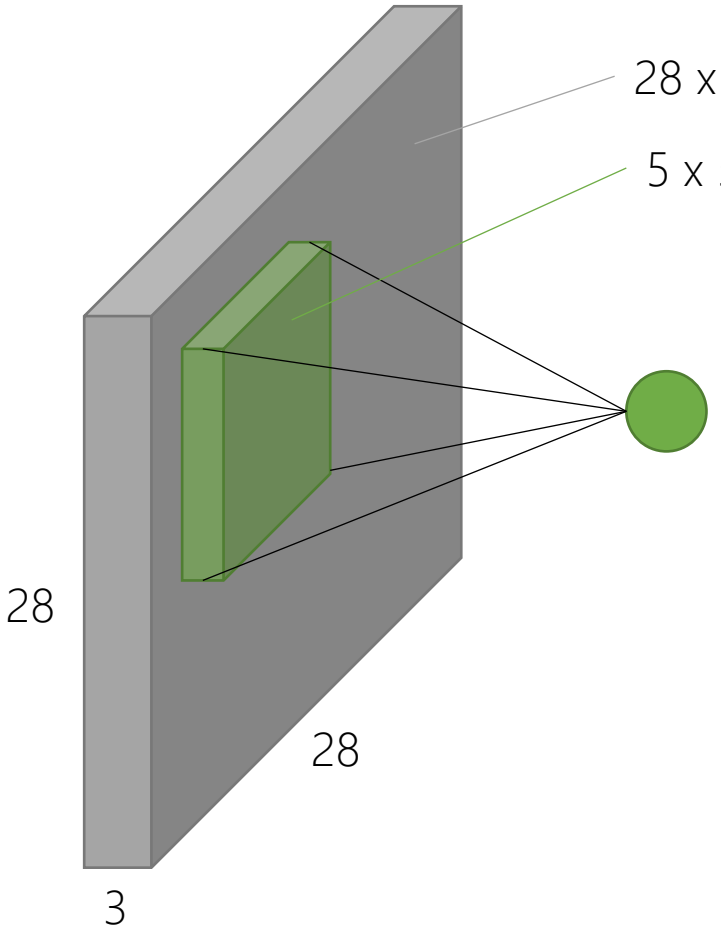

1 number as a result of taking dot product between the kernel and a small 5 x 5 x 3 chunk of the image, plus bias:

$$w^T x + b$$

Convolution Layer

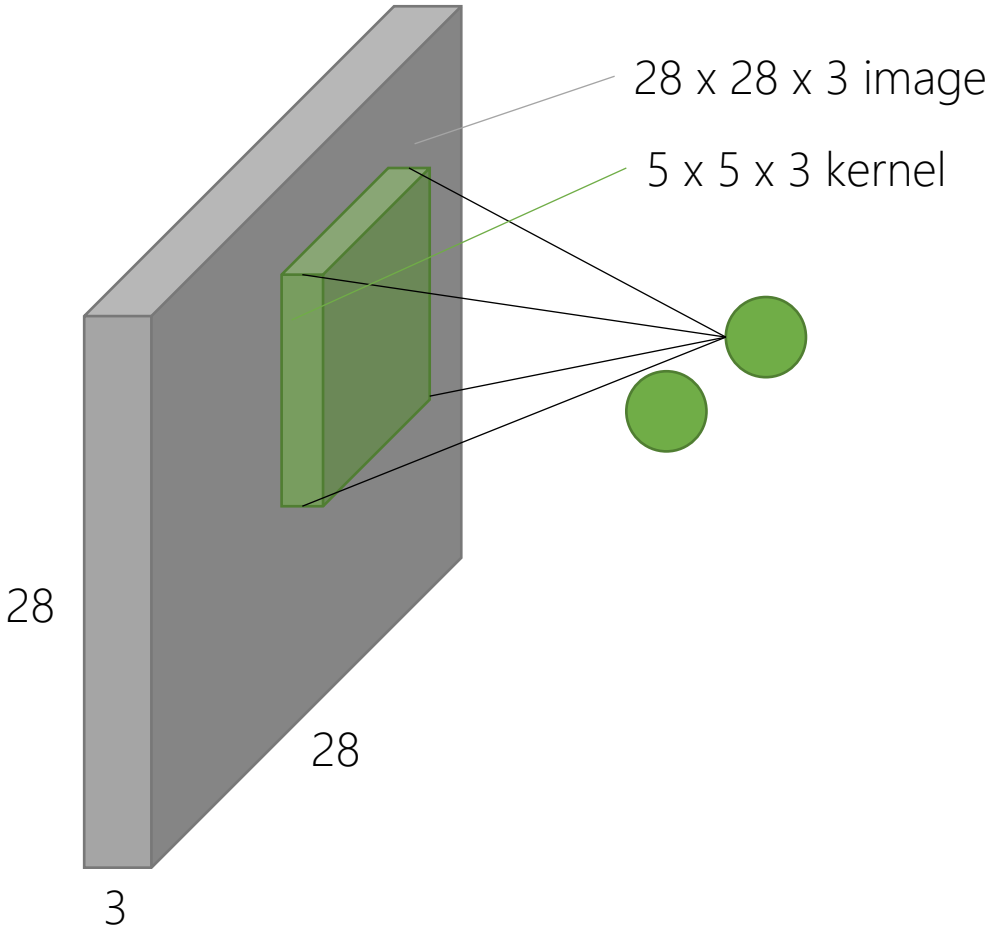

Convolution Layer

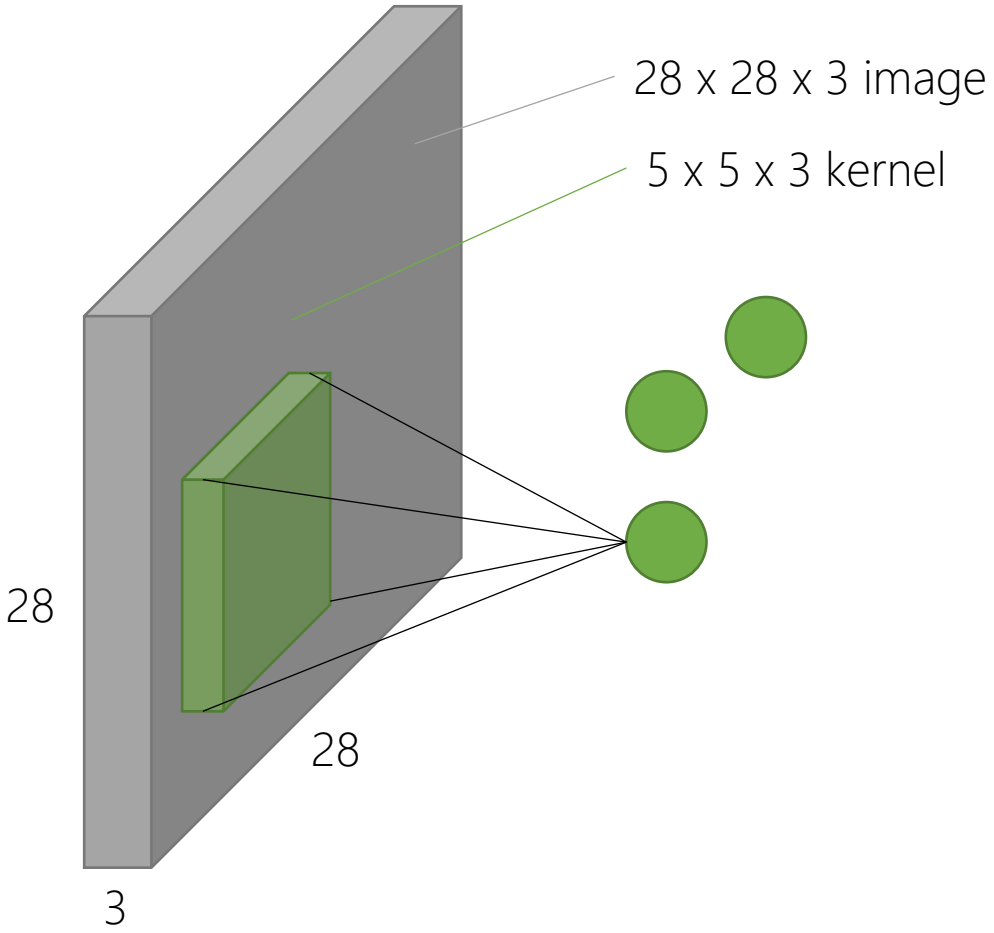

Convolution Layer

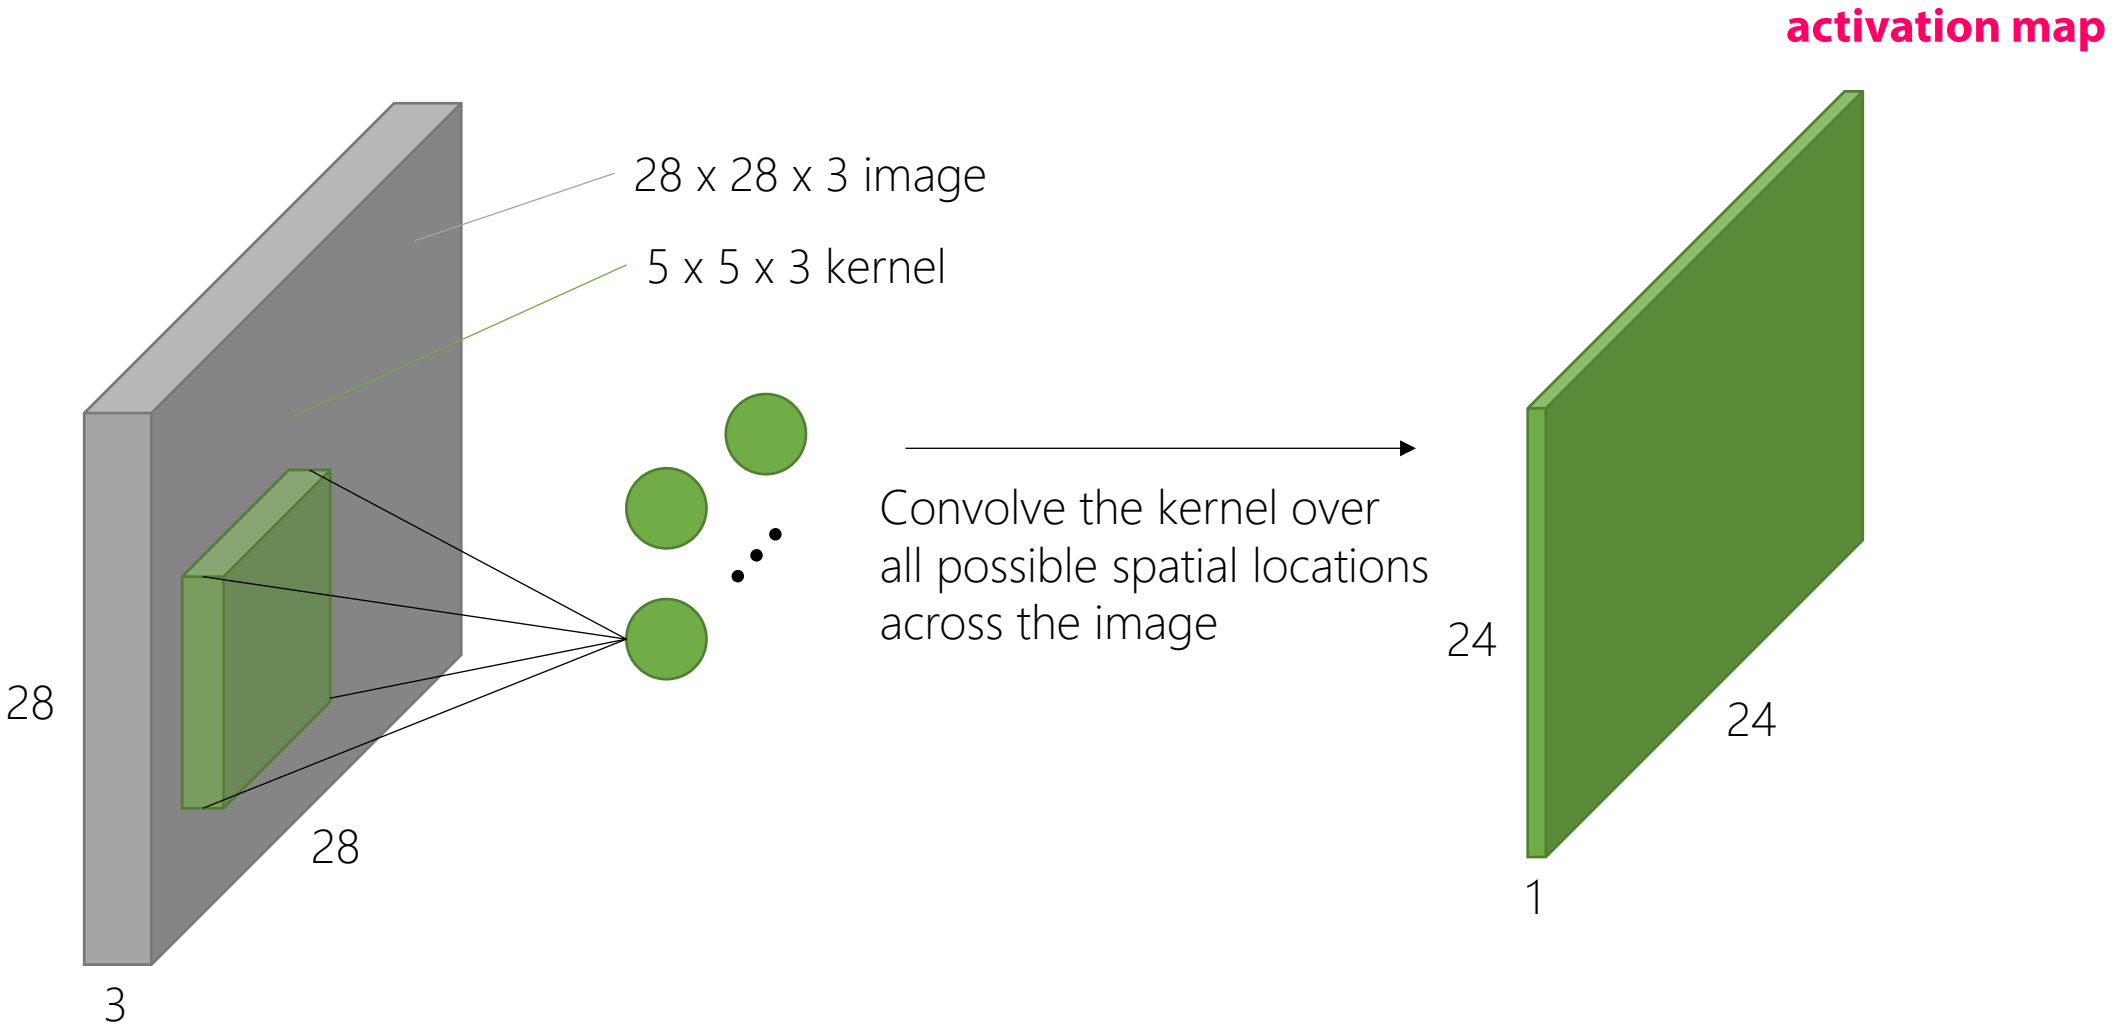

Convolution Layer

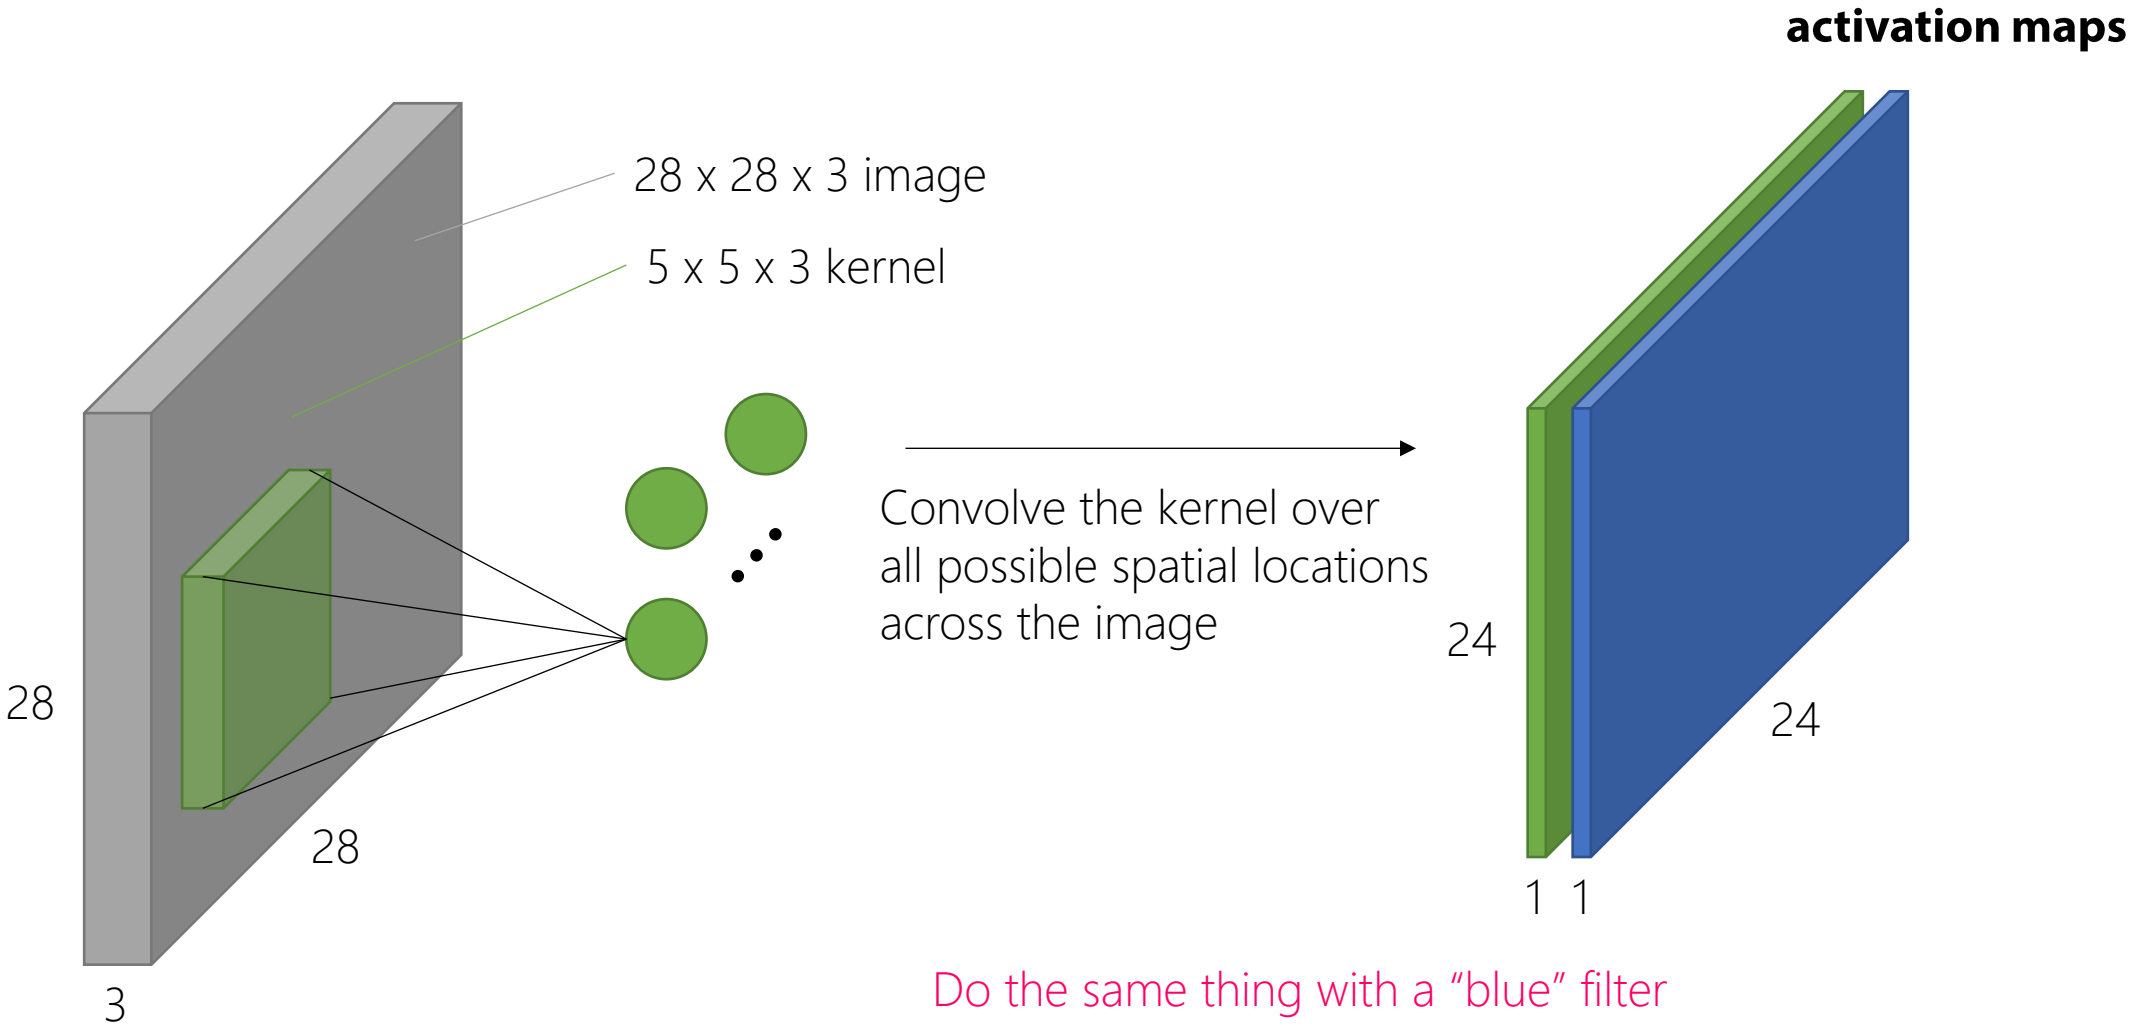

Do the same thing with a “blue” filter

Convolution Layer

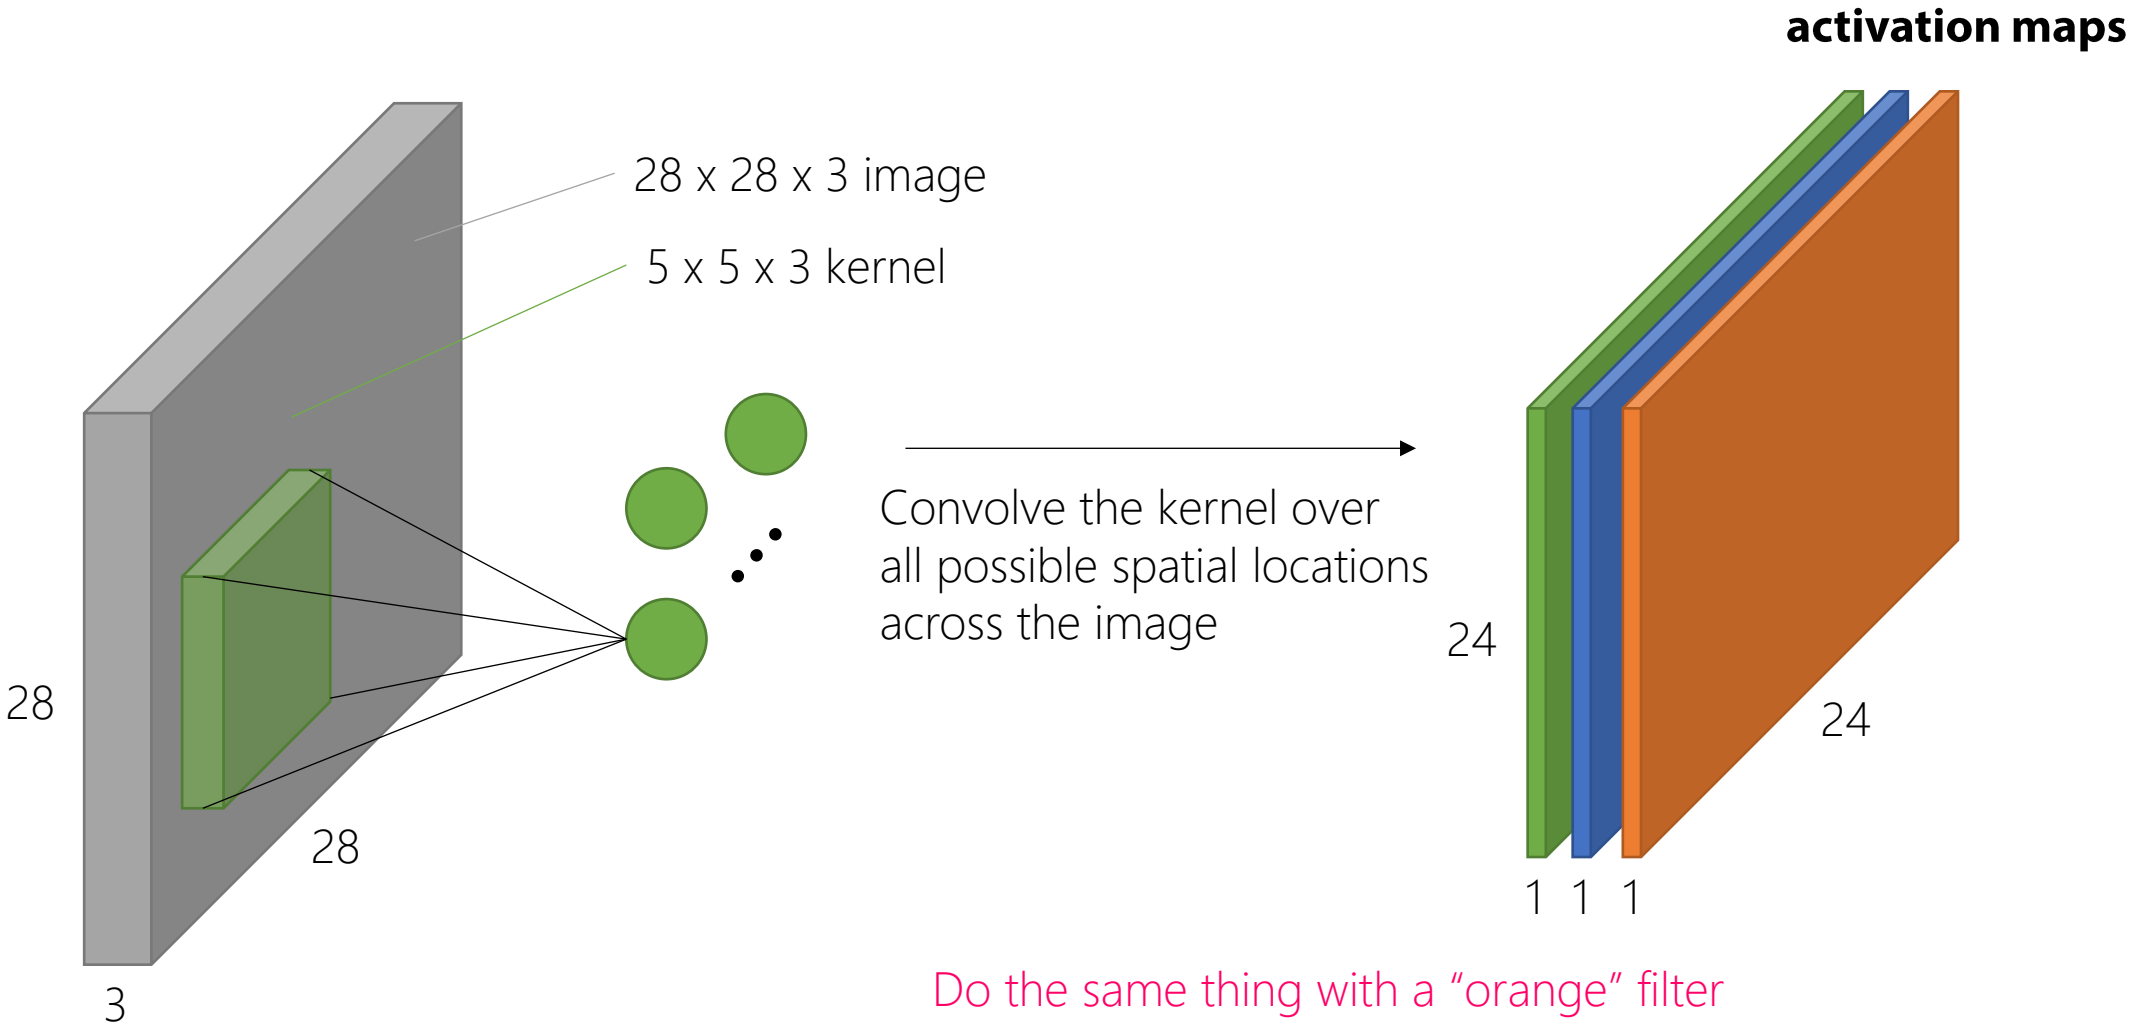

Do the same thing with a "orange" filter

Convolution Layer

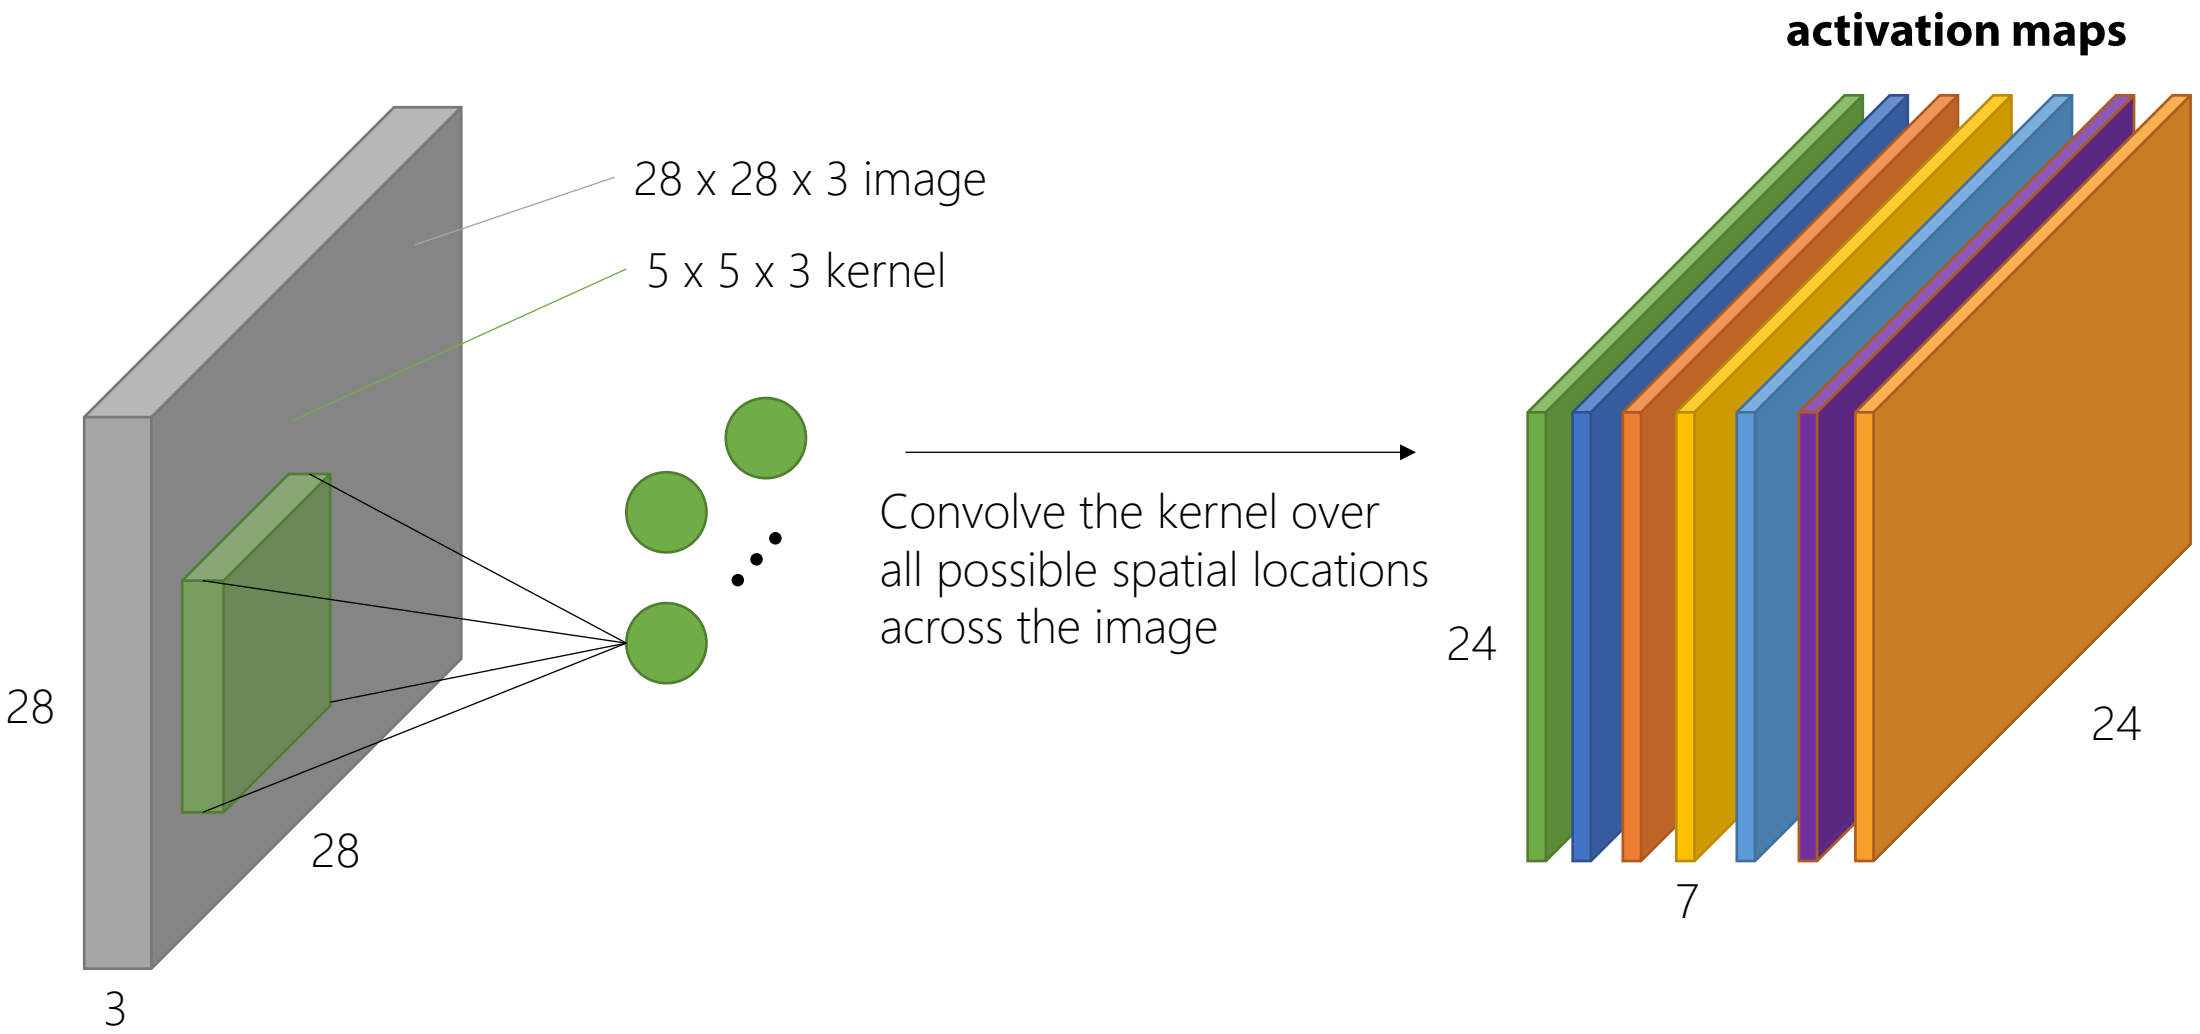

ConvNet is a sequence of convolution layers

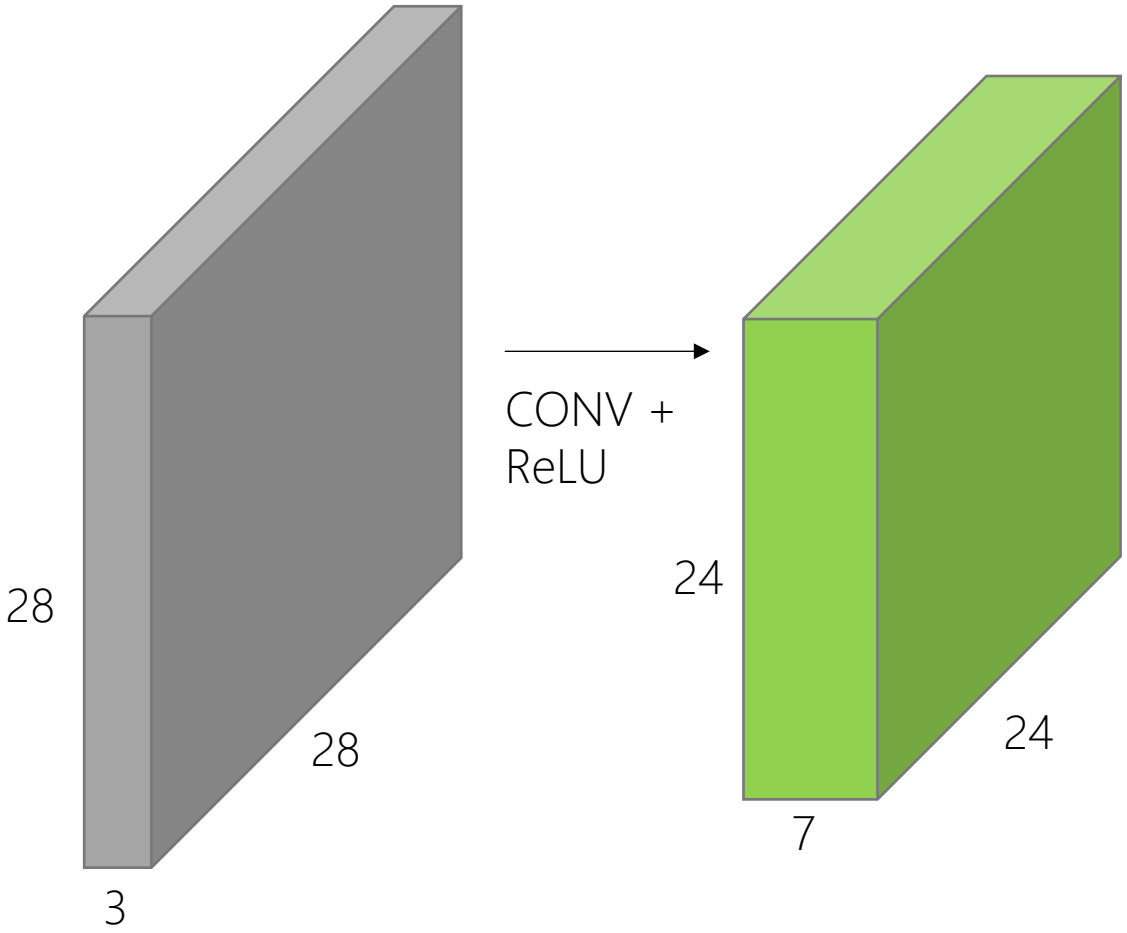

ConvNet is a sequence of convolution layers

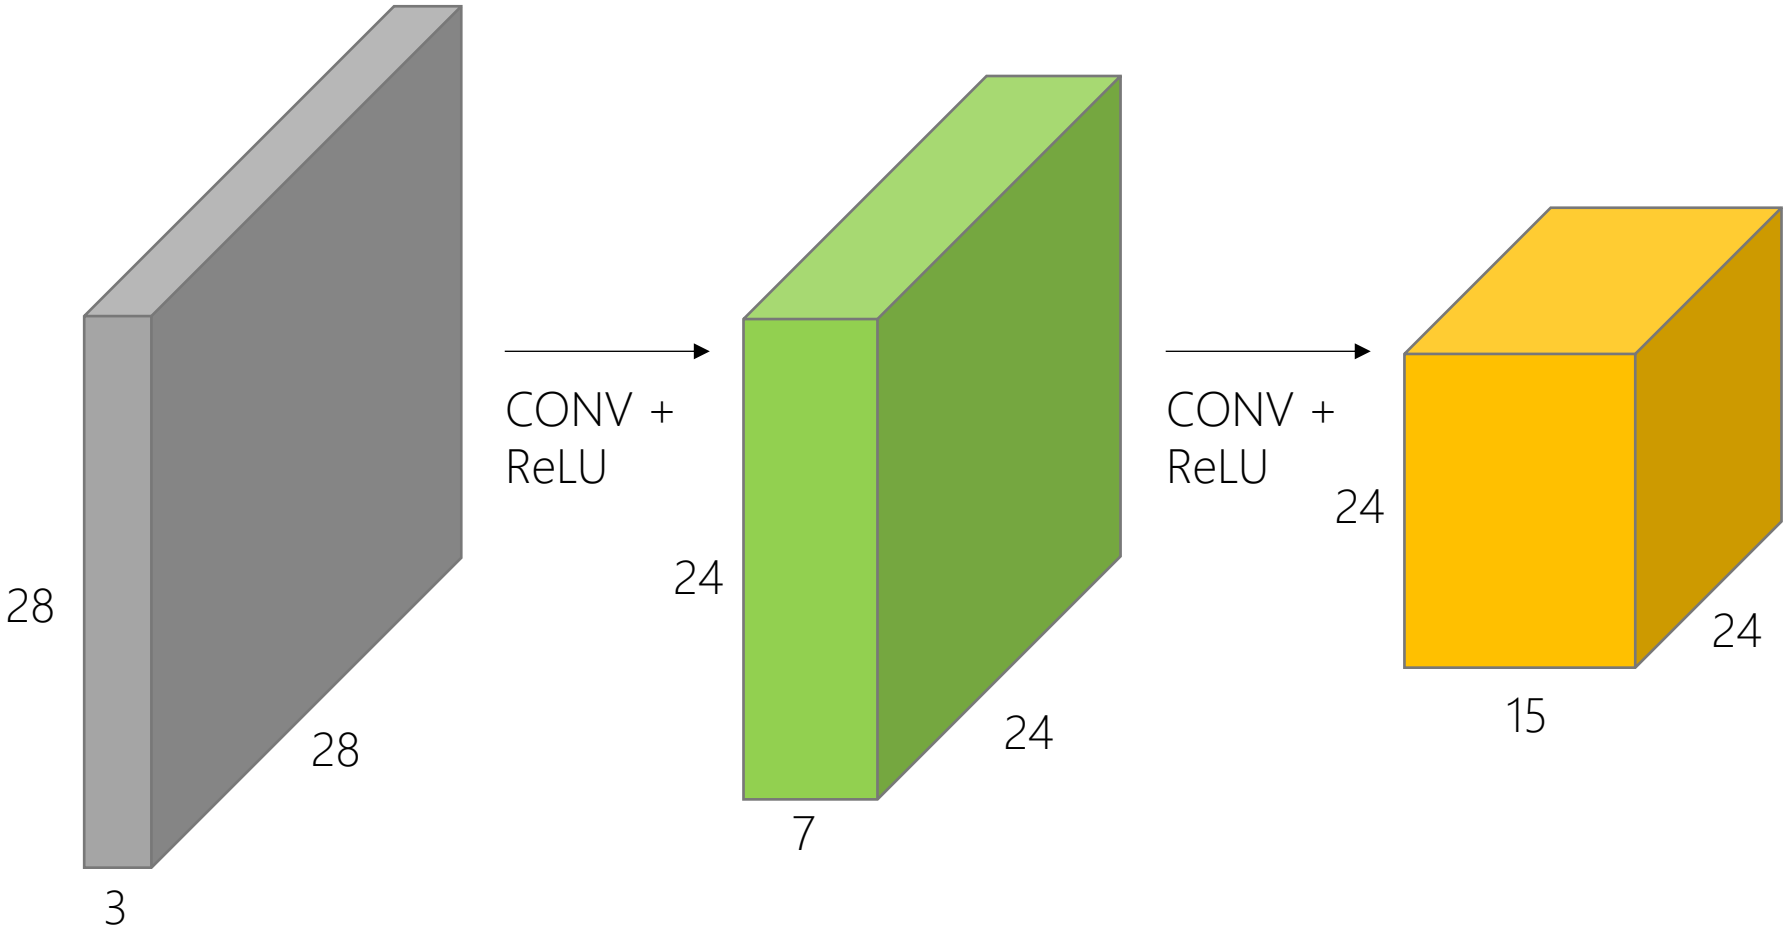

ConvNet is a sequence of convolution layers

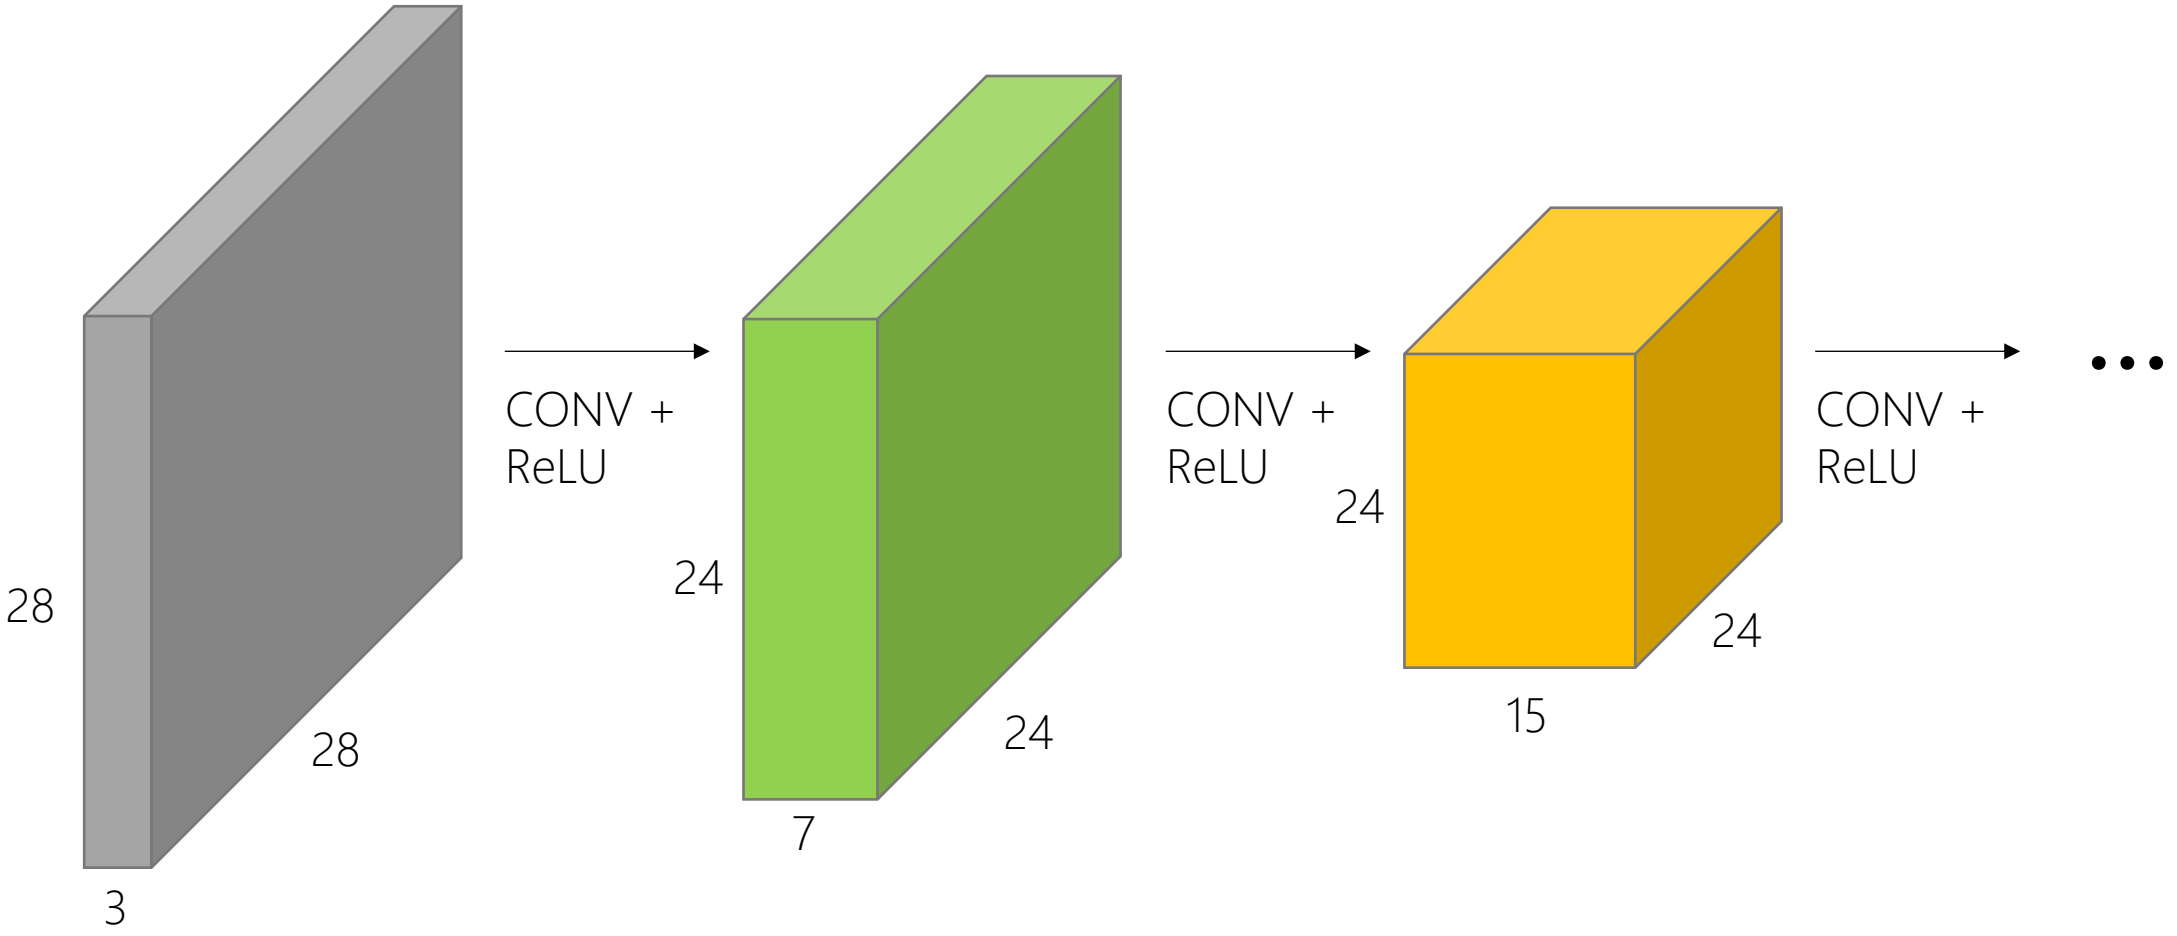

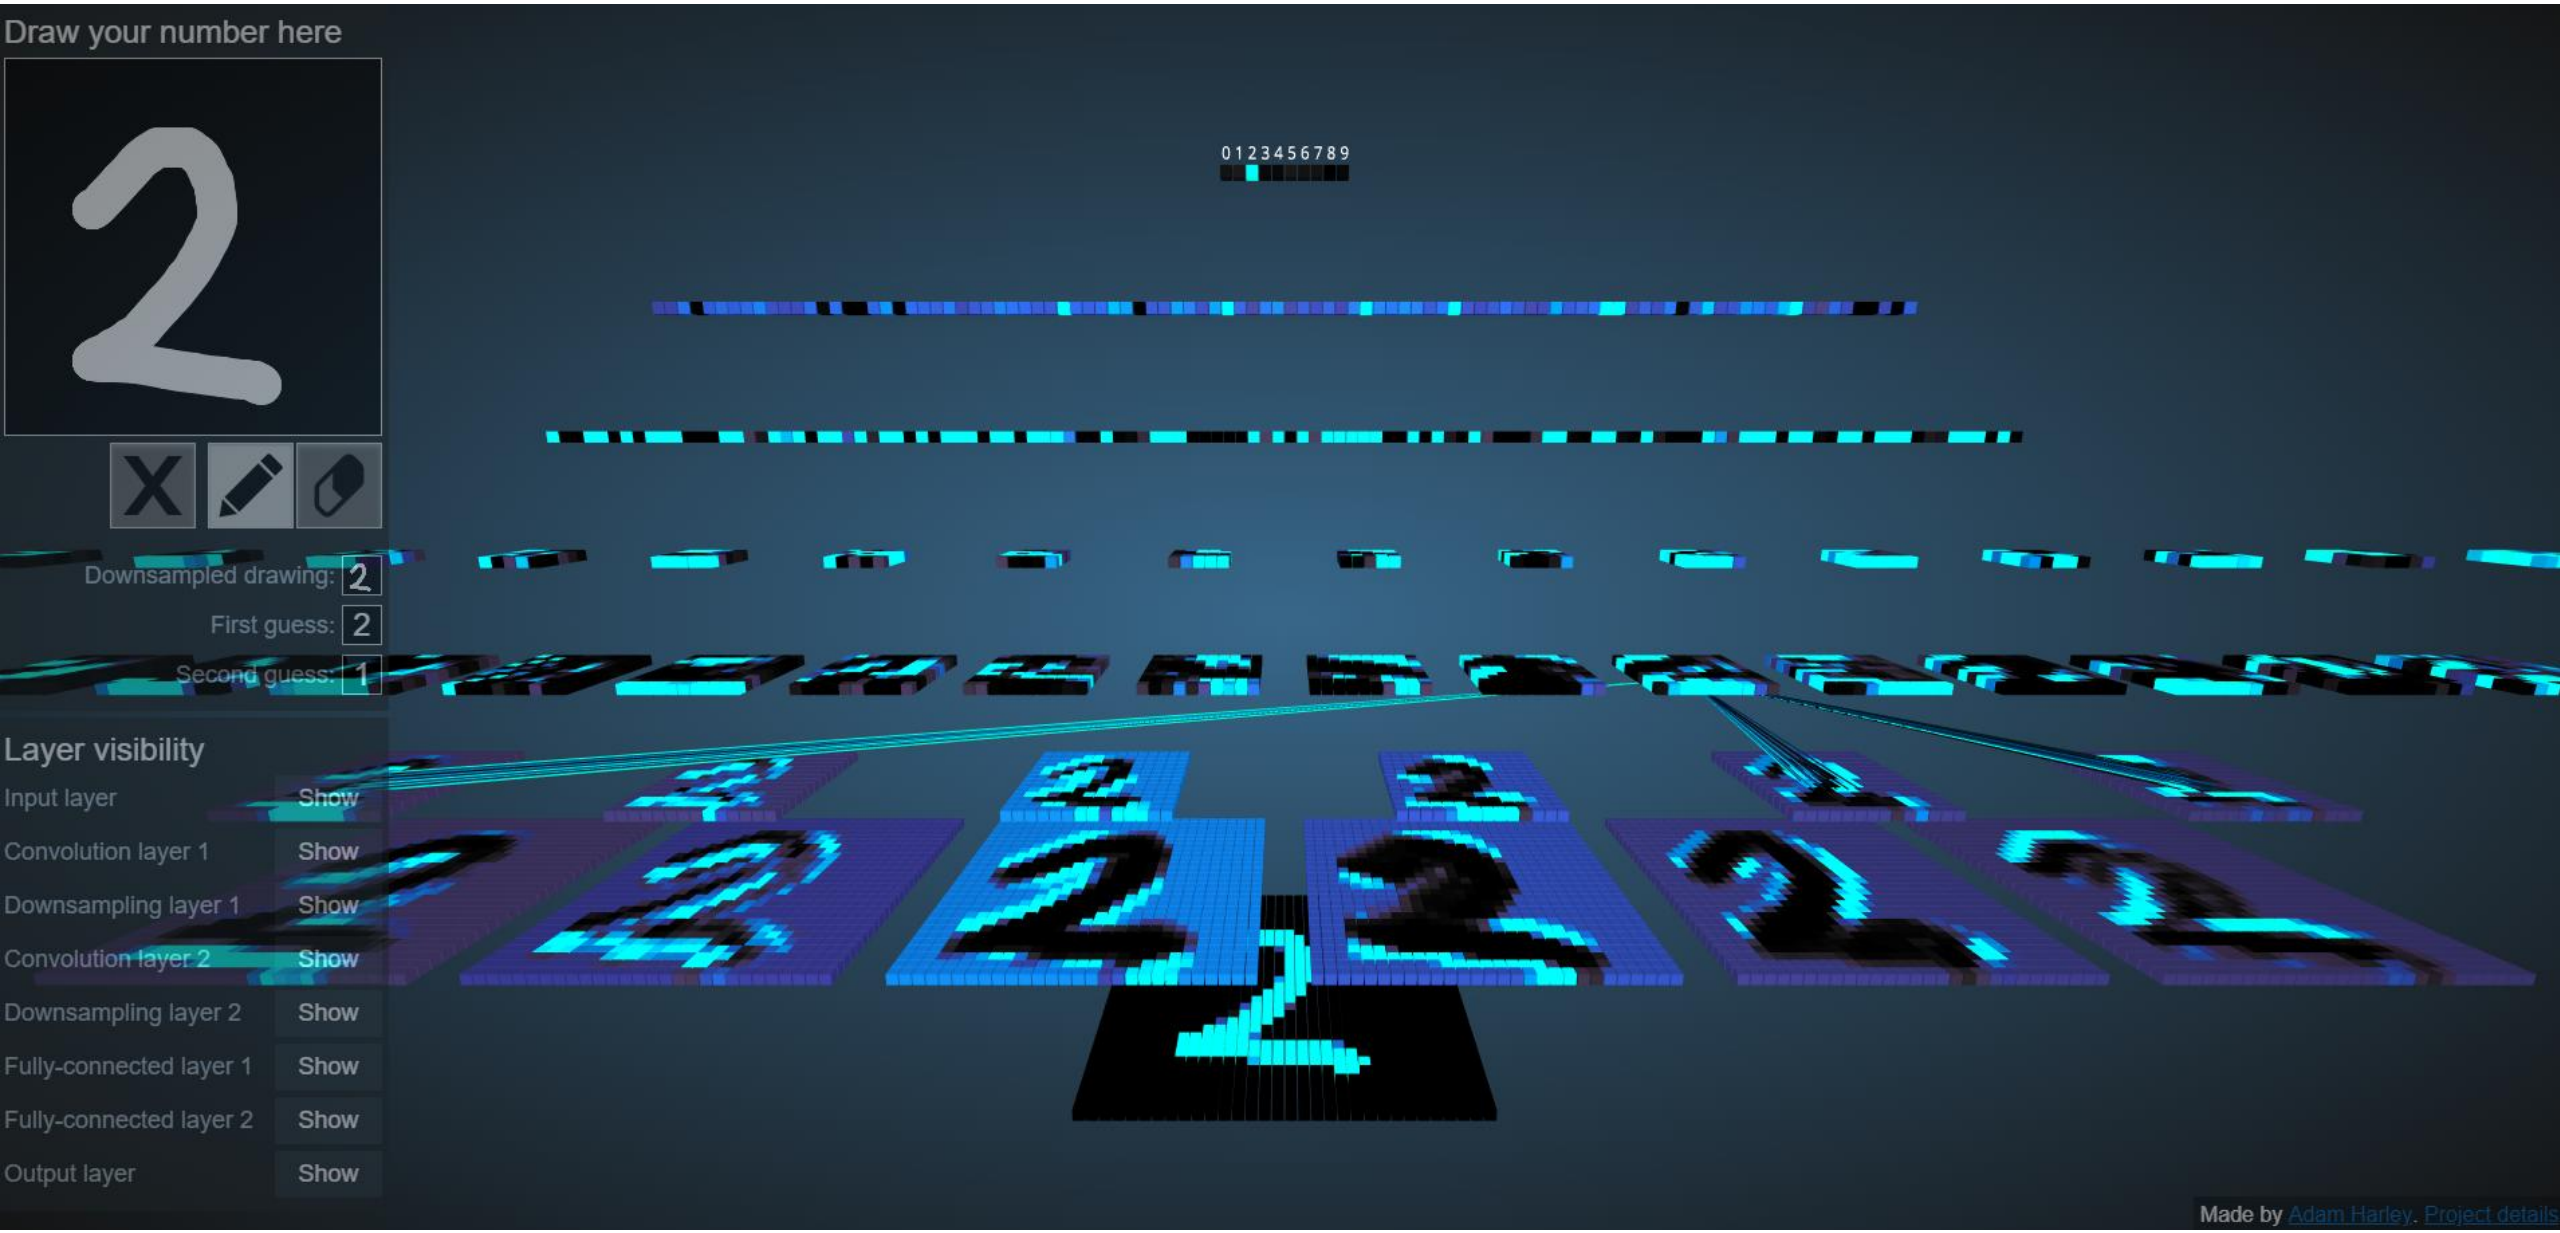

http://scs.ryerson.ca/~aharley/vis/conv/

ConvNet is a sequence of convolution layers

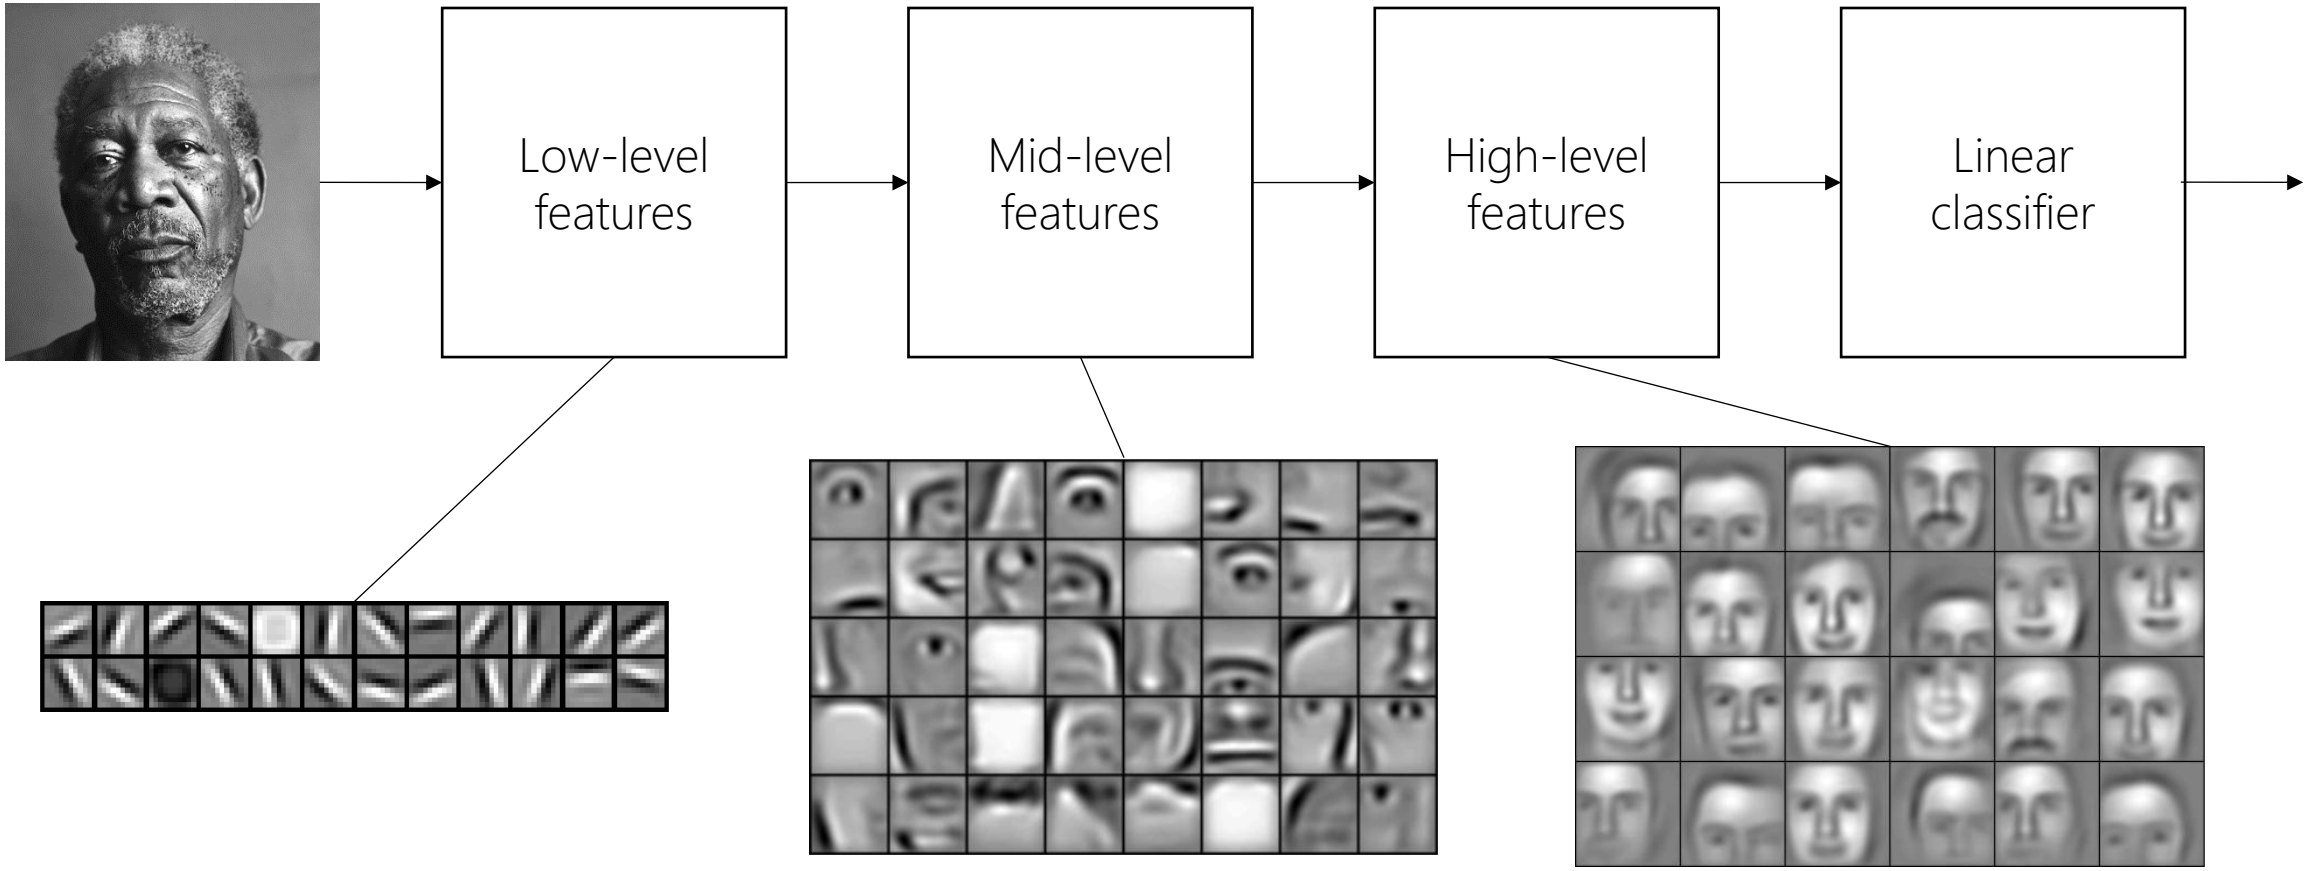

Lee *et al.* (2009)

ConvNet is a sequence of convolution layers

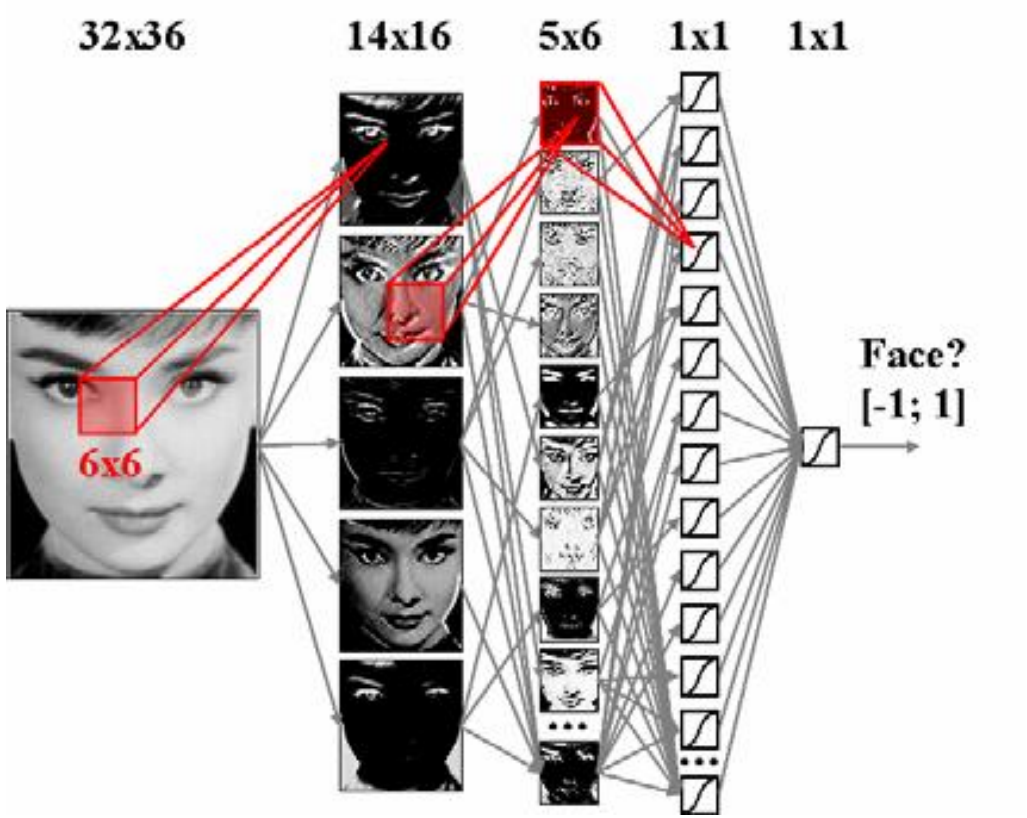

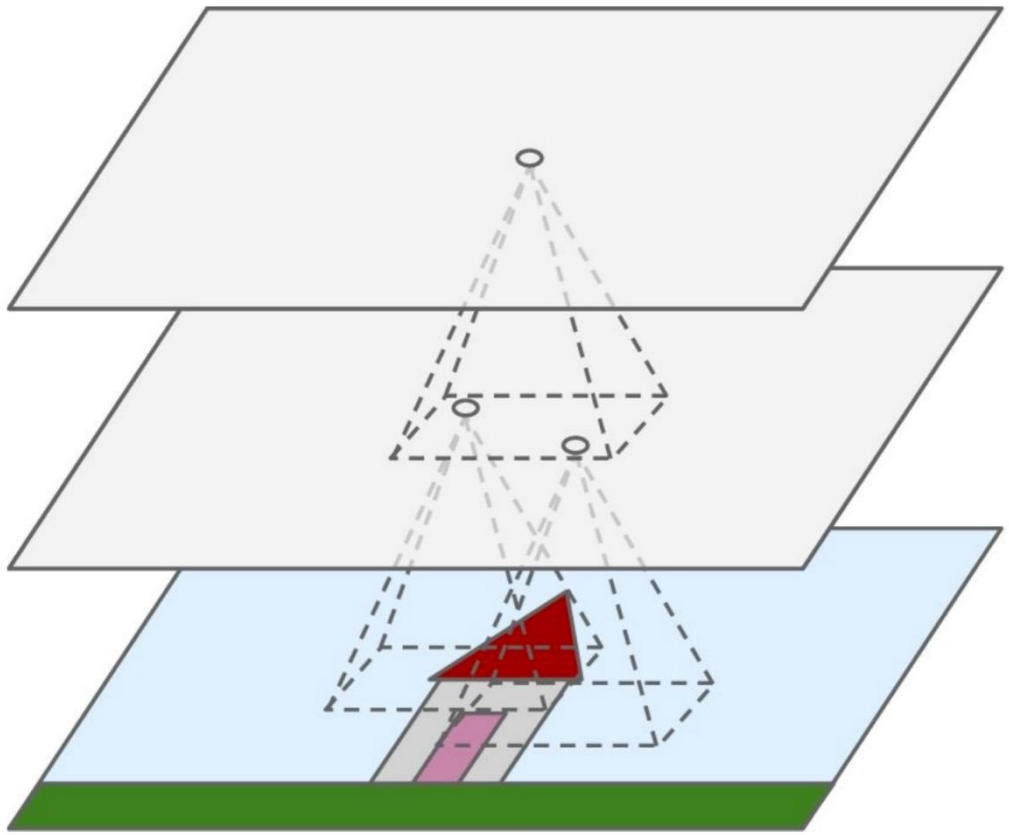

Conv kernels are trainable

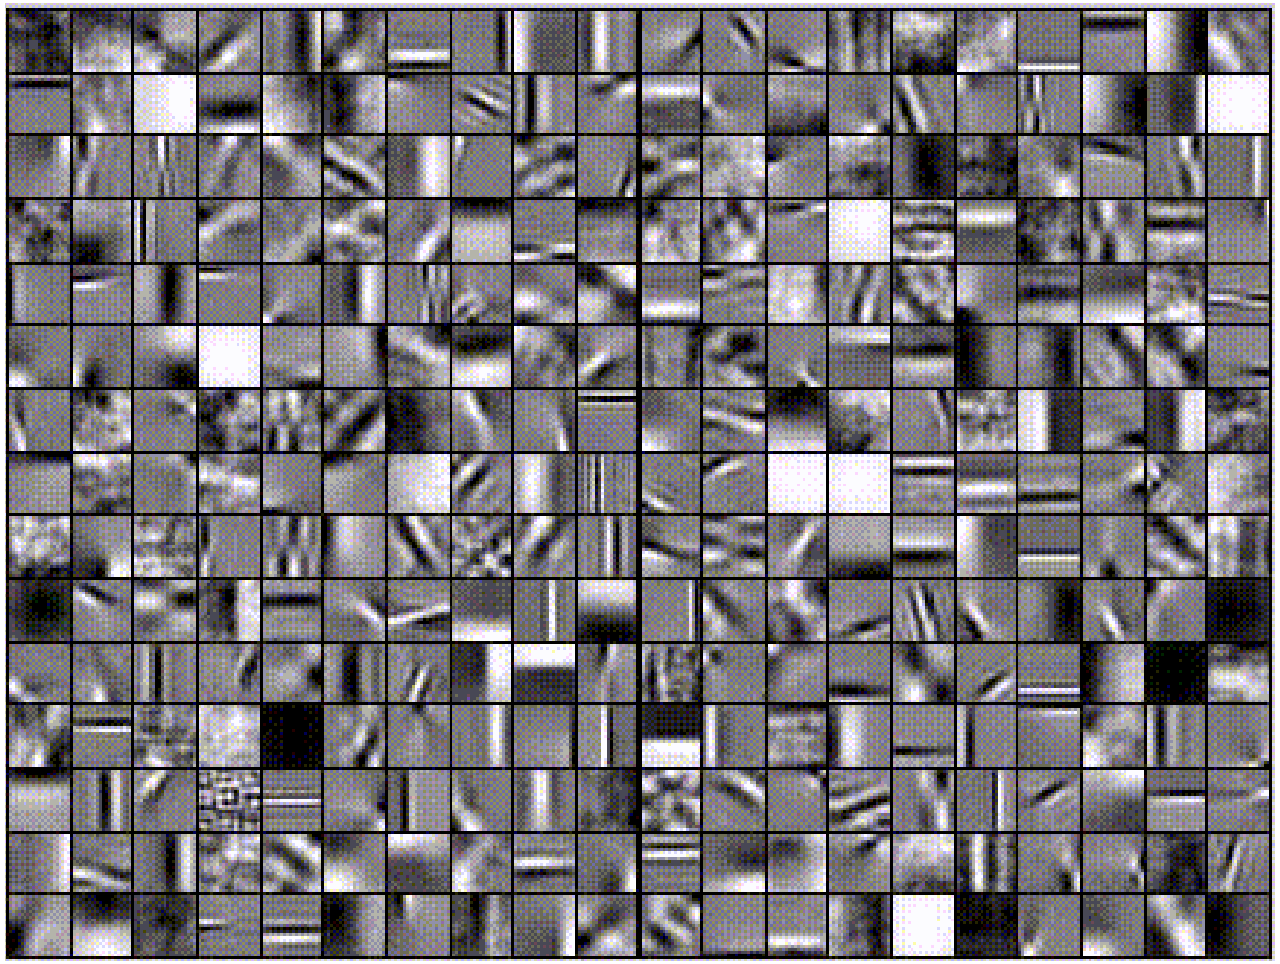

A closer look…

- Convolution

$$(f * g)(t) \coloneqq \int_{-\infty}^{\infty} f(\tau) g(t-\tau)\, d\tau$$

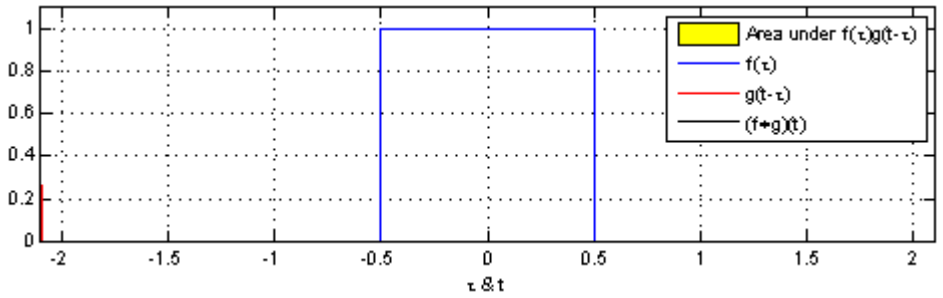

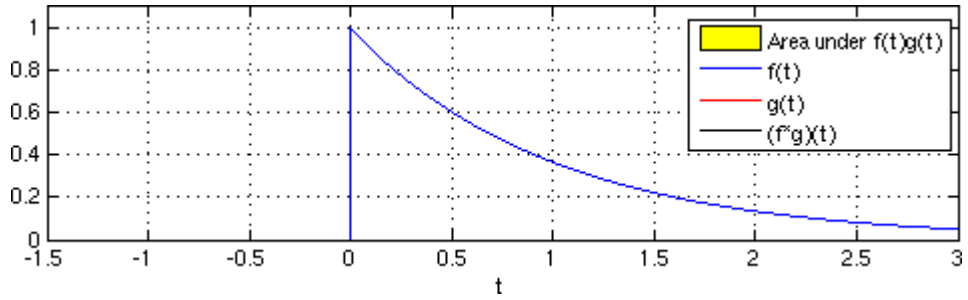

A closer look…

- 2D Discrete Convolution

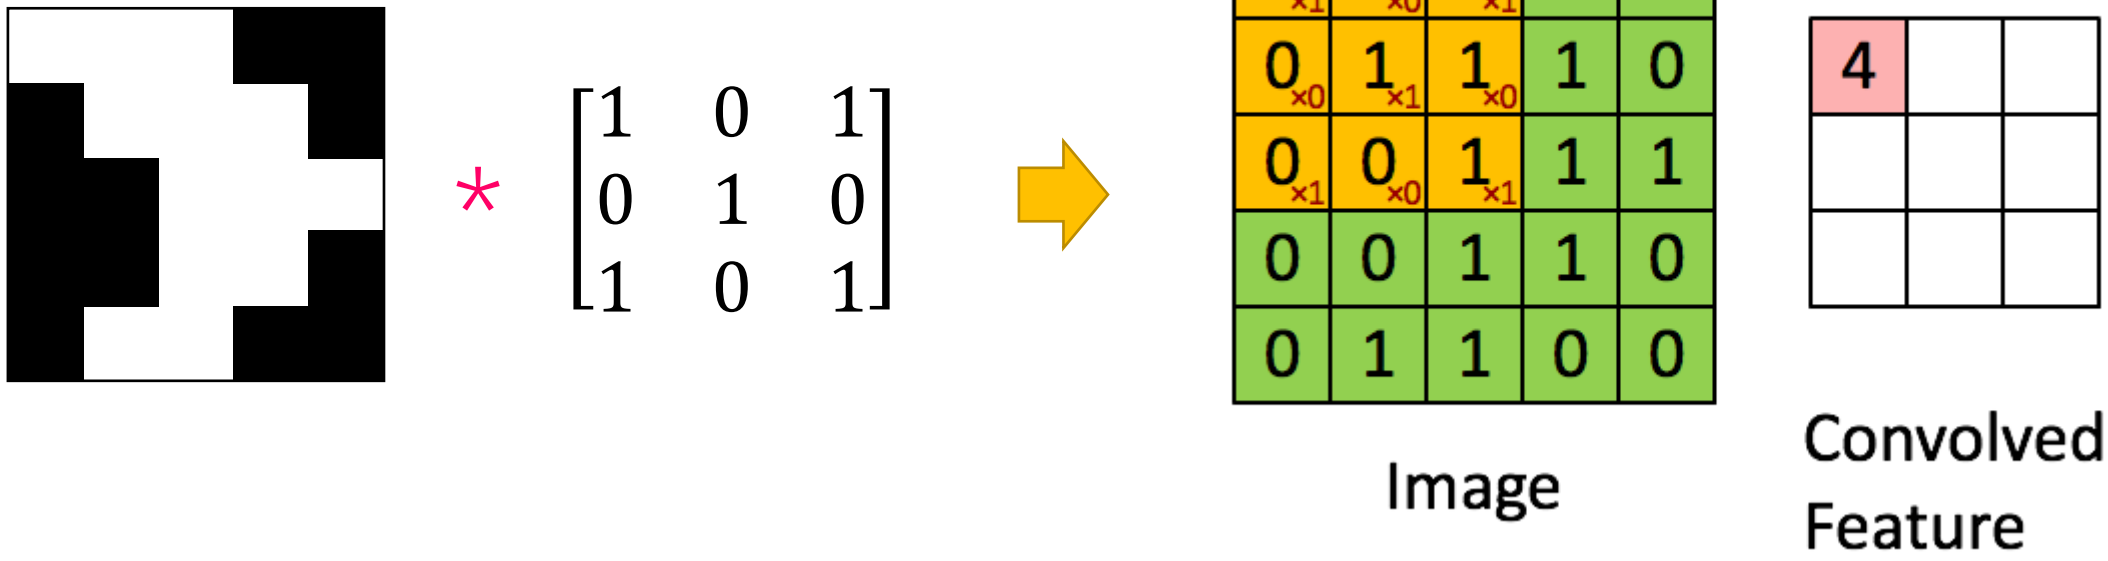

A closer look…

- 2D Discrete Convolution

$$* \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} =$$

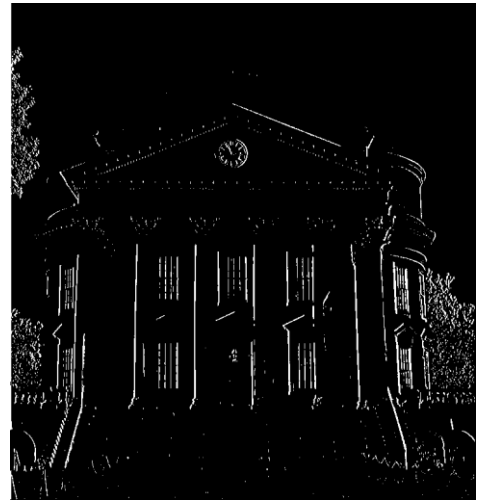

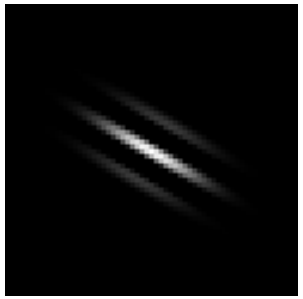

$$* \quad =$$

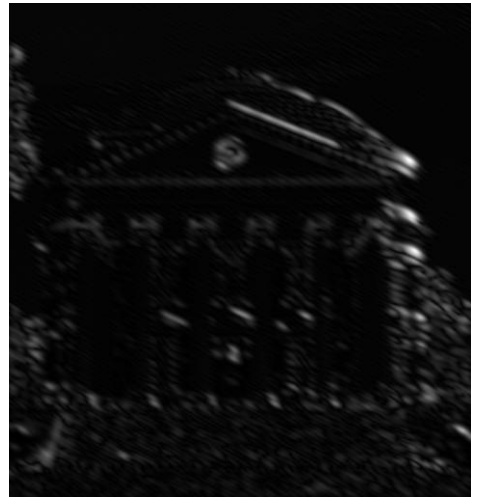

Stride & Padding

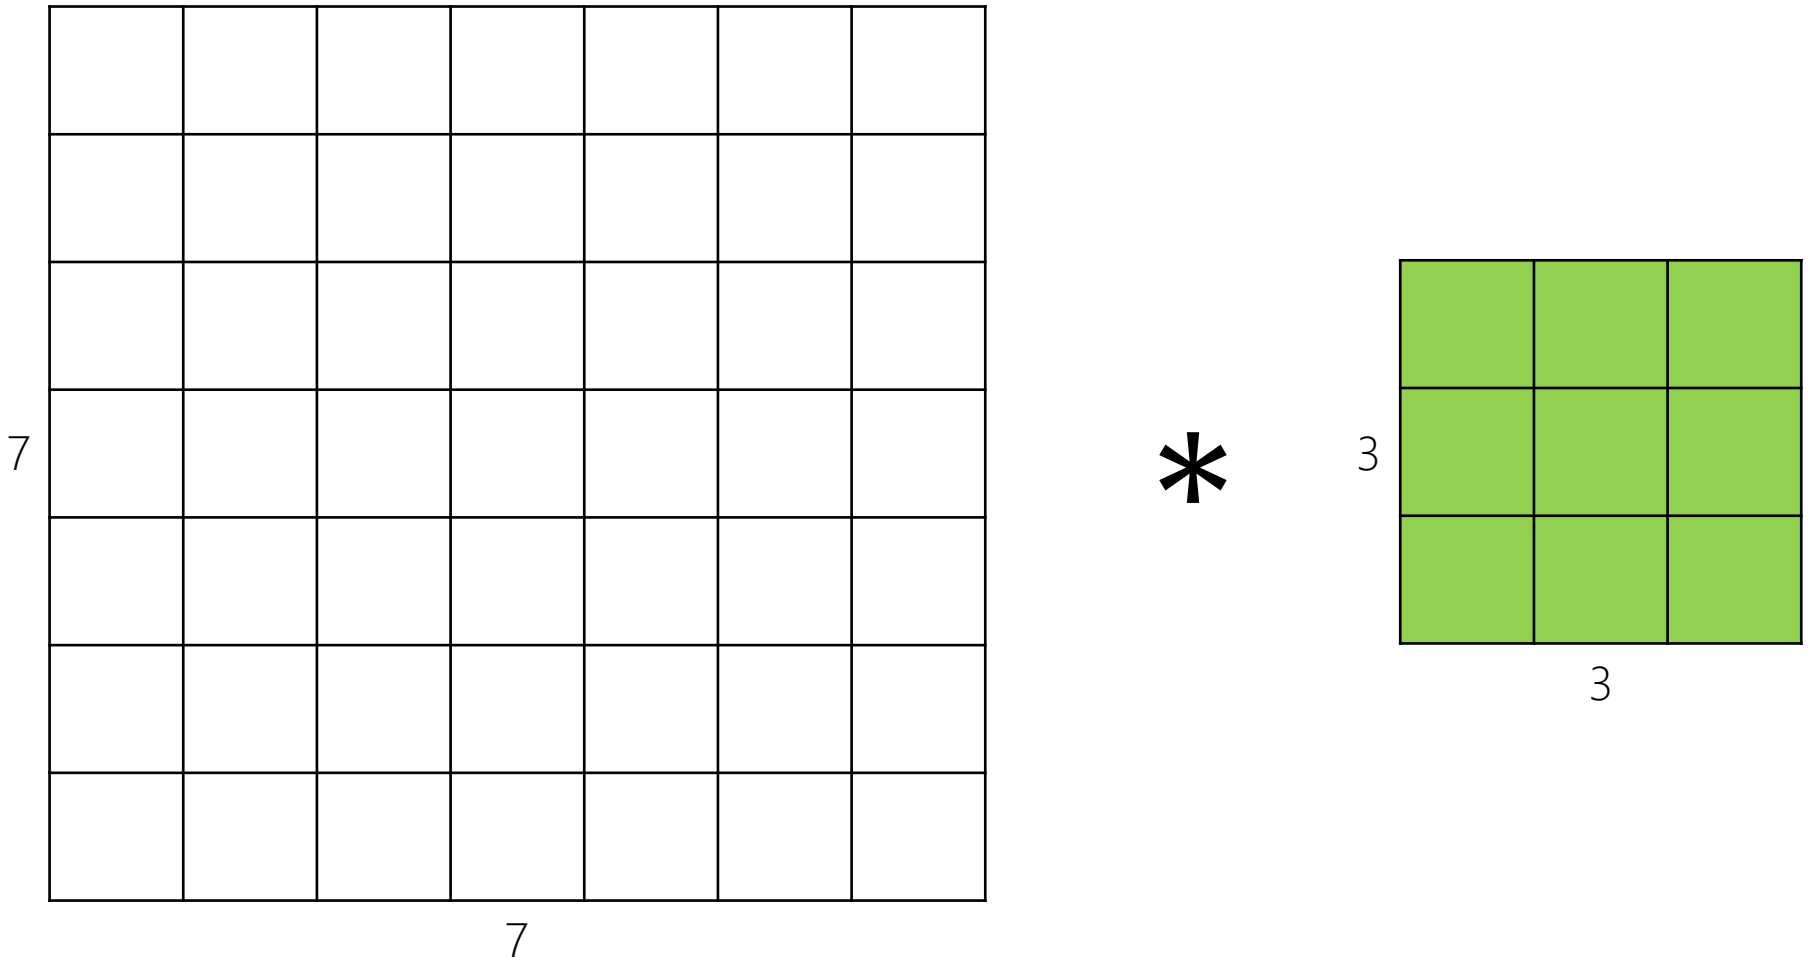

Stride & Padding

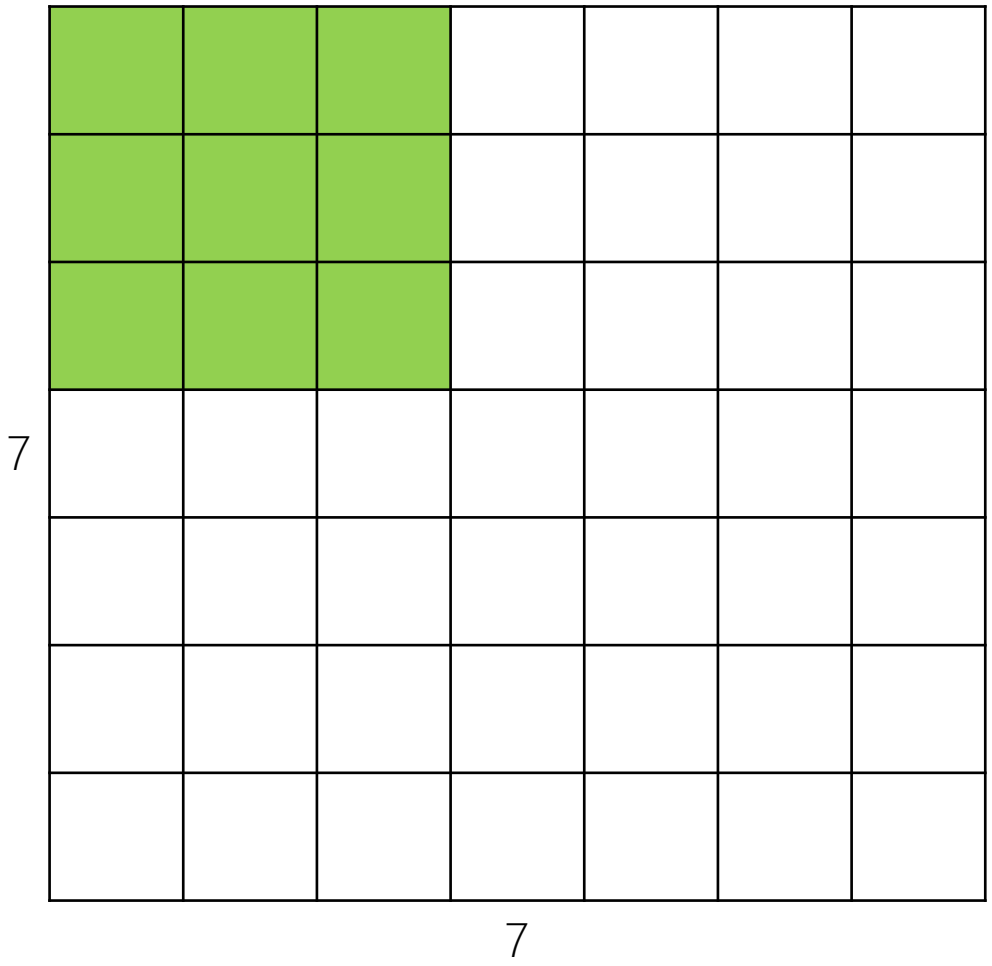

Stride & Padding

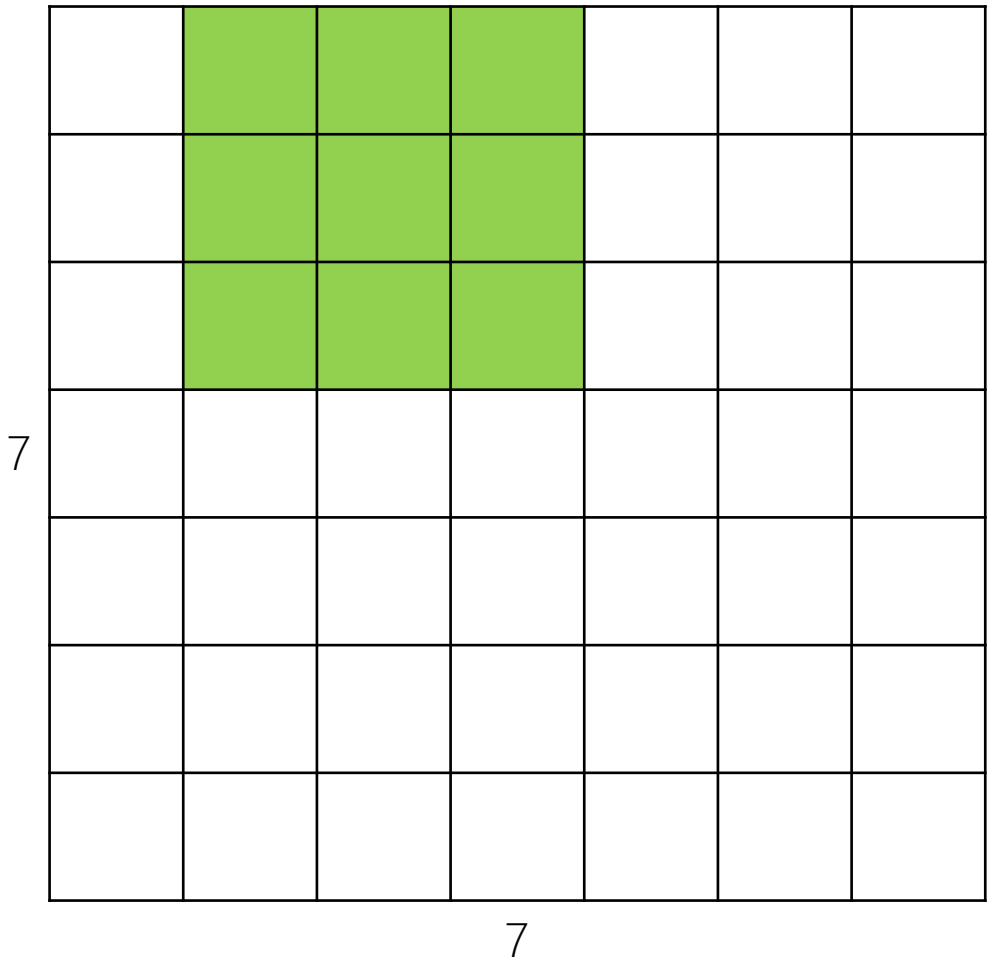

Stride & Padding

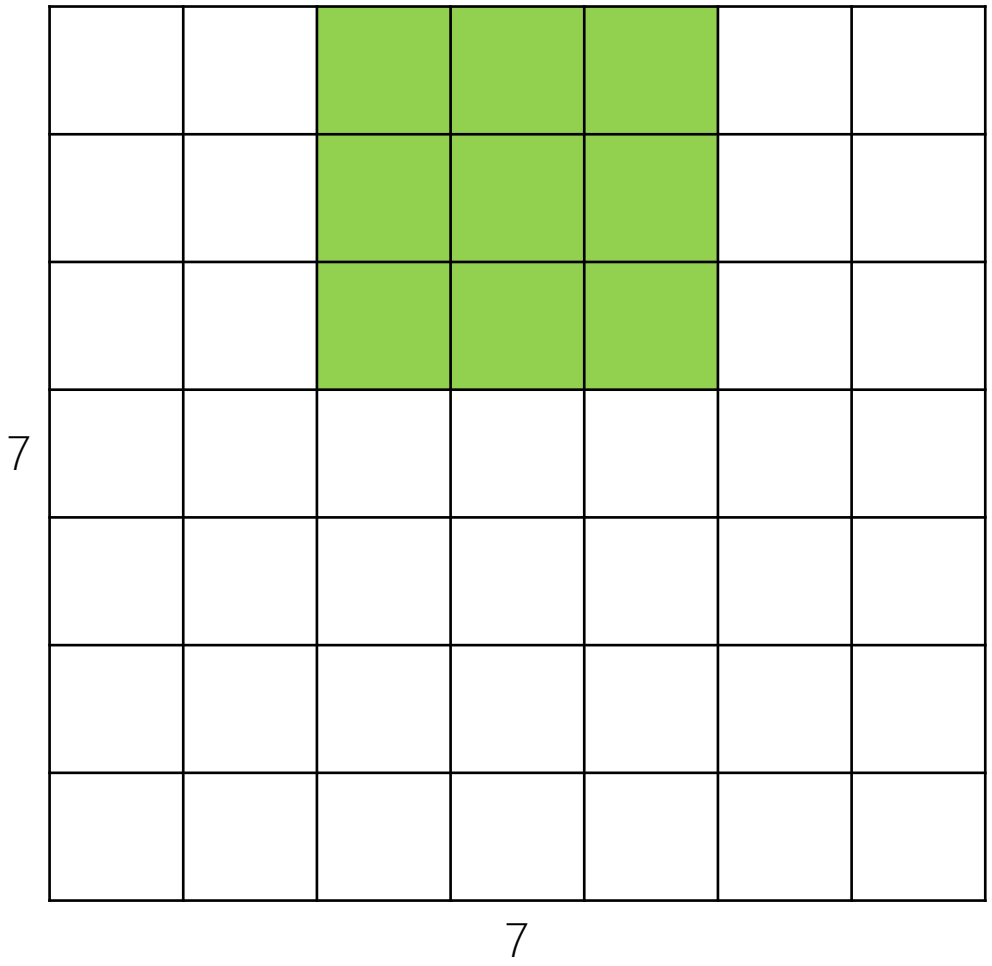

Stride & Padding

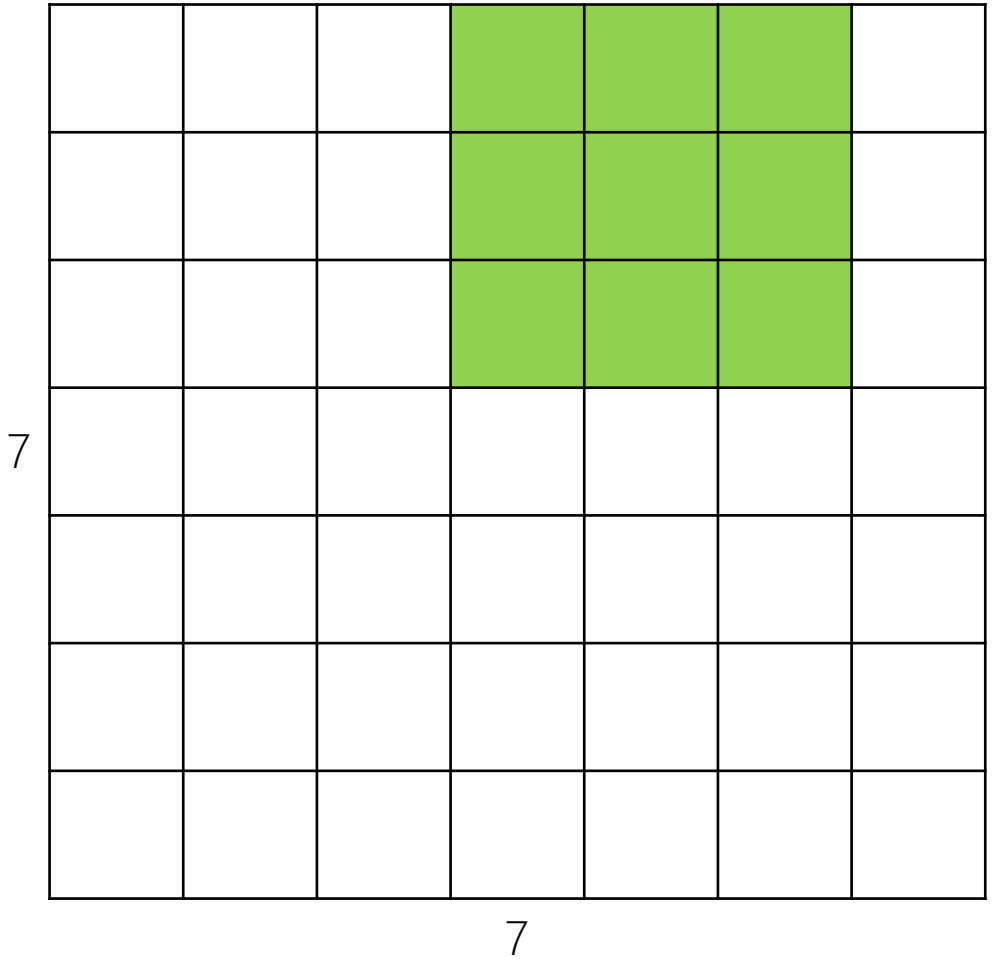

Stride & Padding

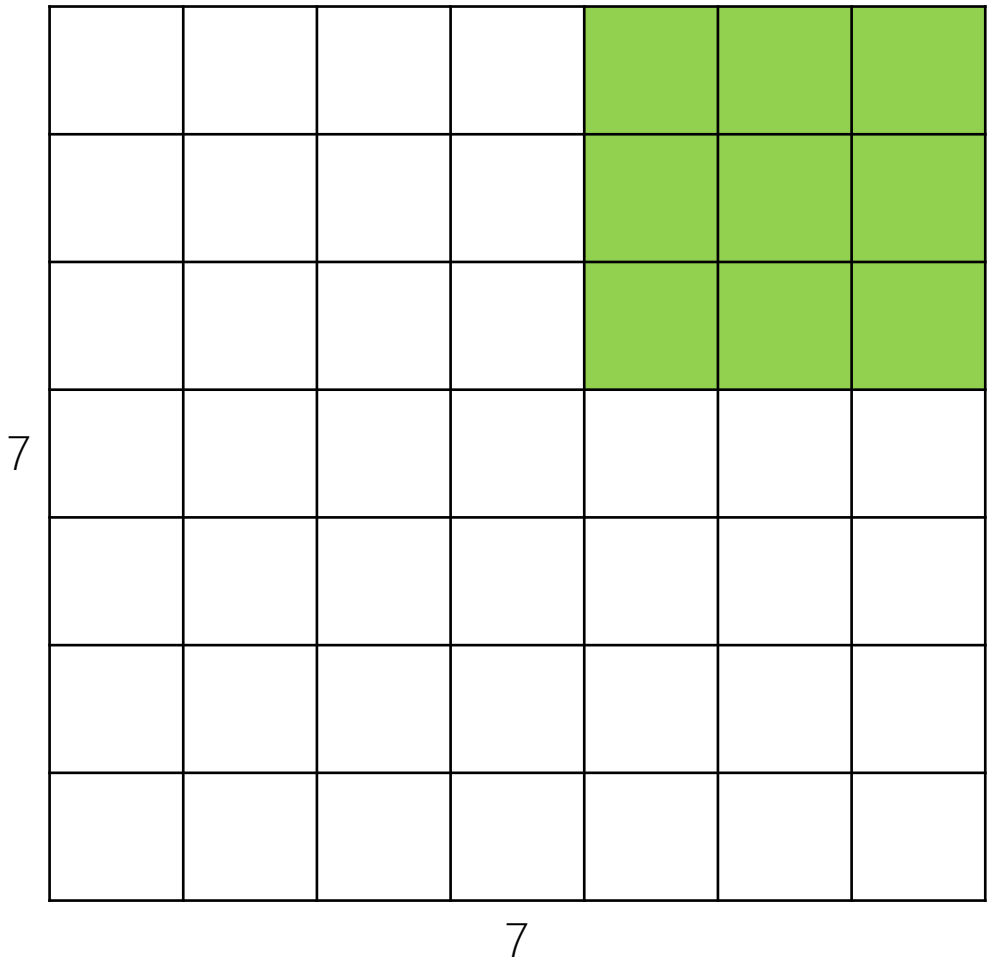

Stride & Padding

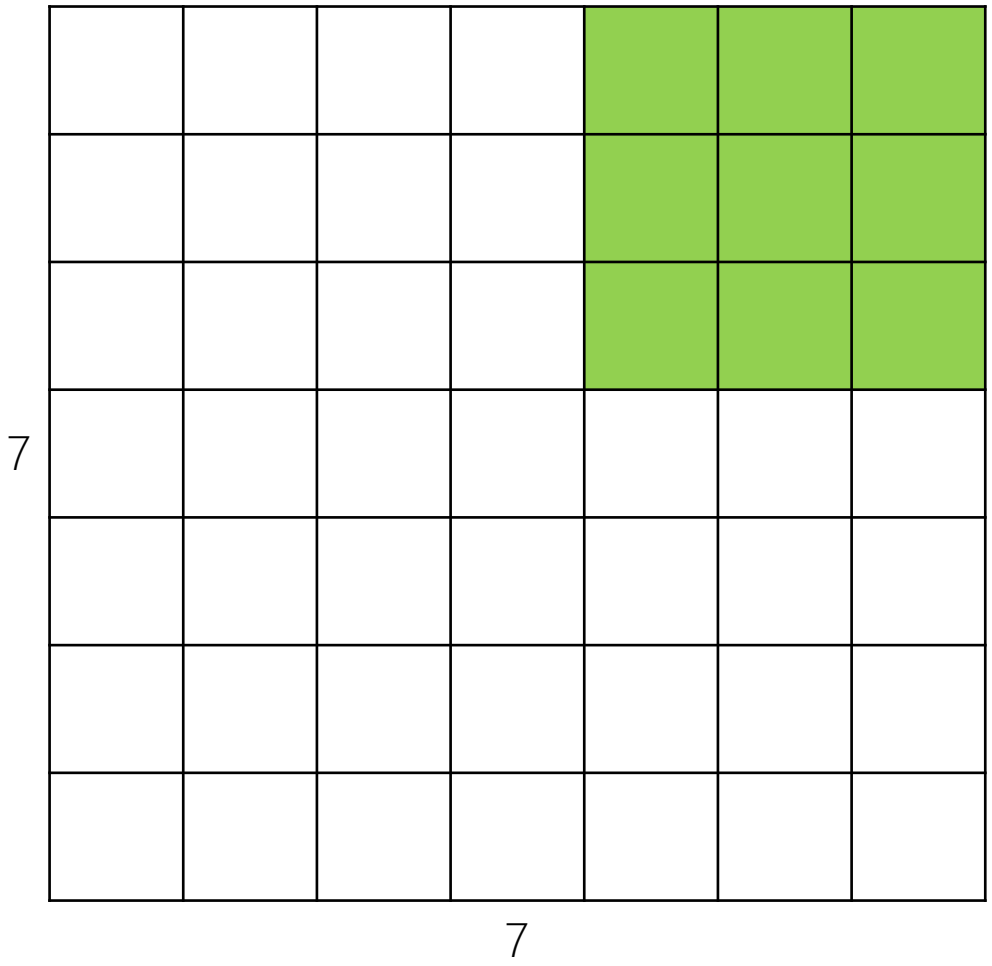

→ 5x5 output

Stride & Padding

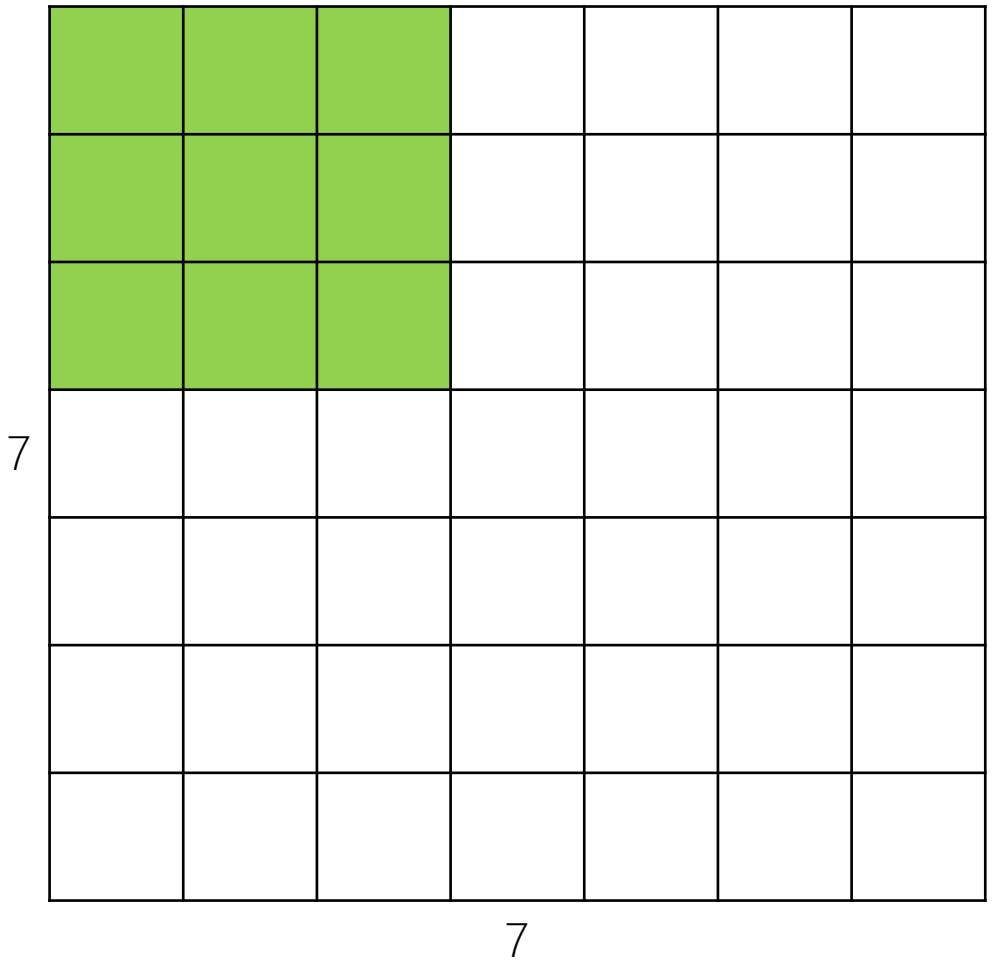

Now with stride 2

Stride & Padding

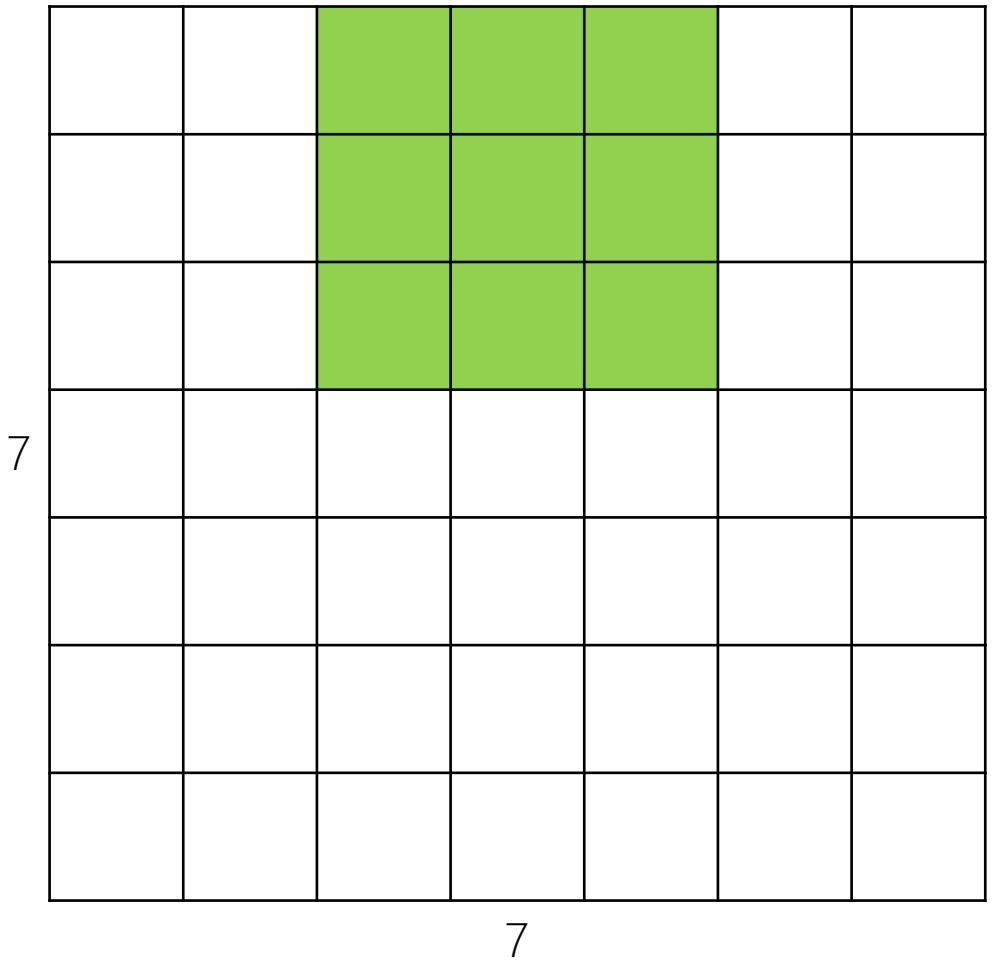

Now with stride 2

Stride & Padding

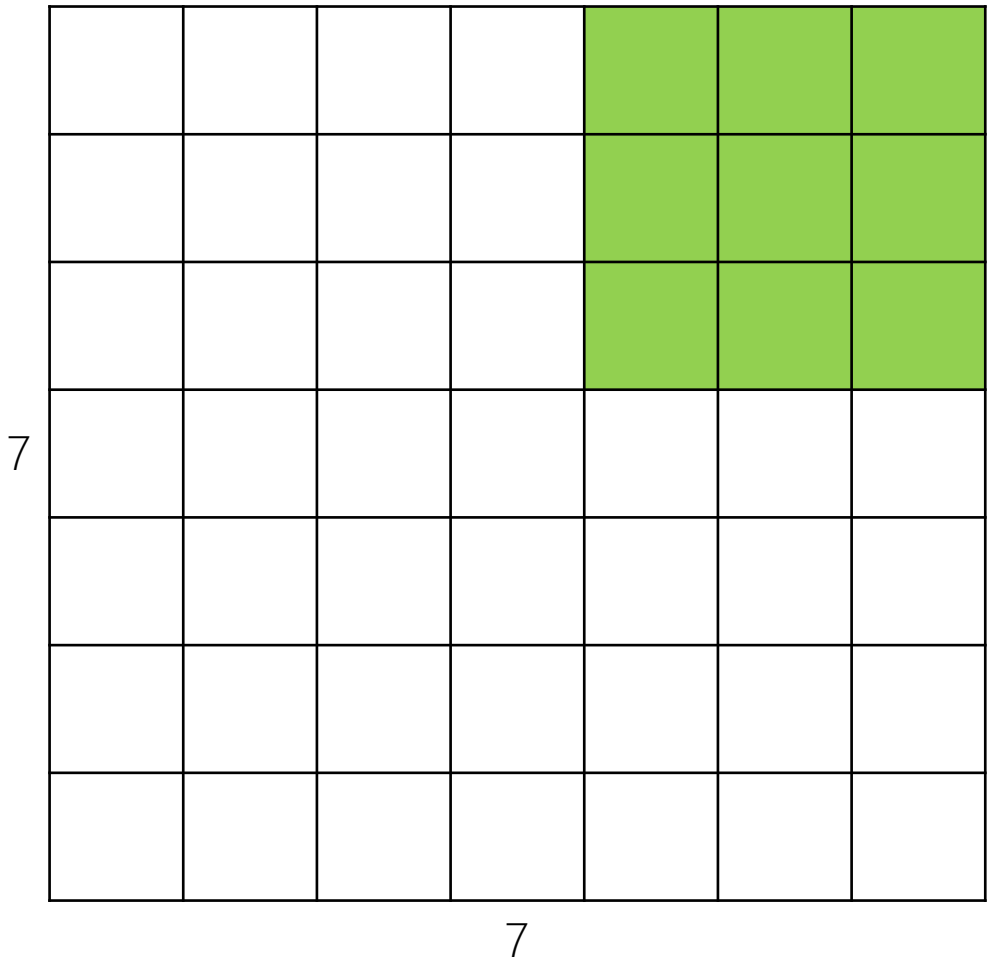

Now with stride 2

→ 3x3 output

Stride & Padding

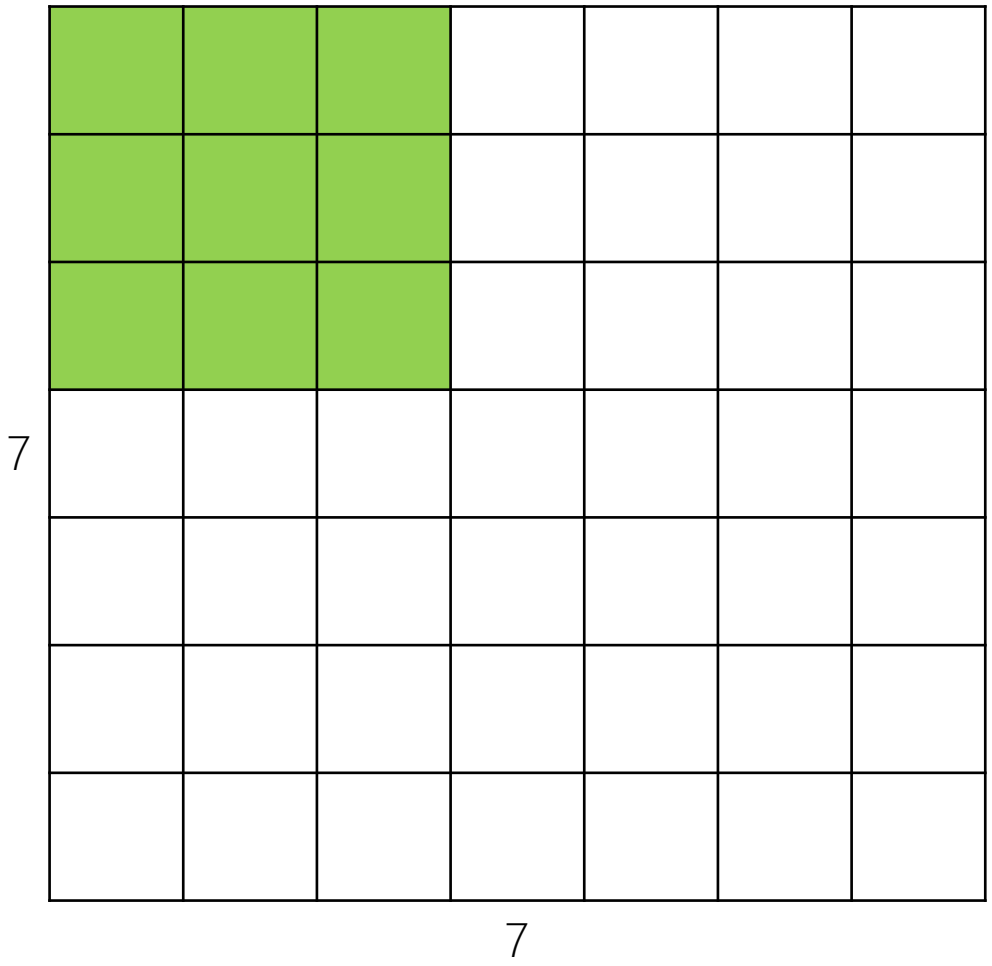

How about with stride 3?

Stride & Padding

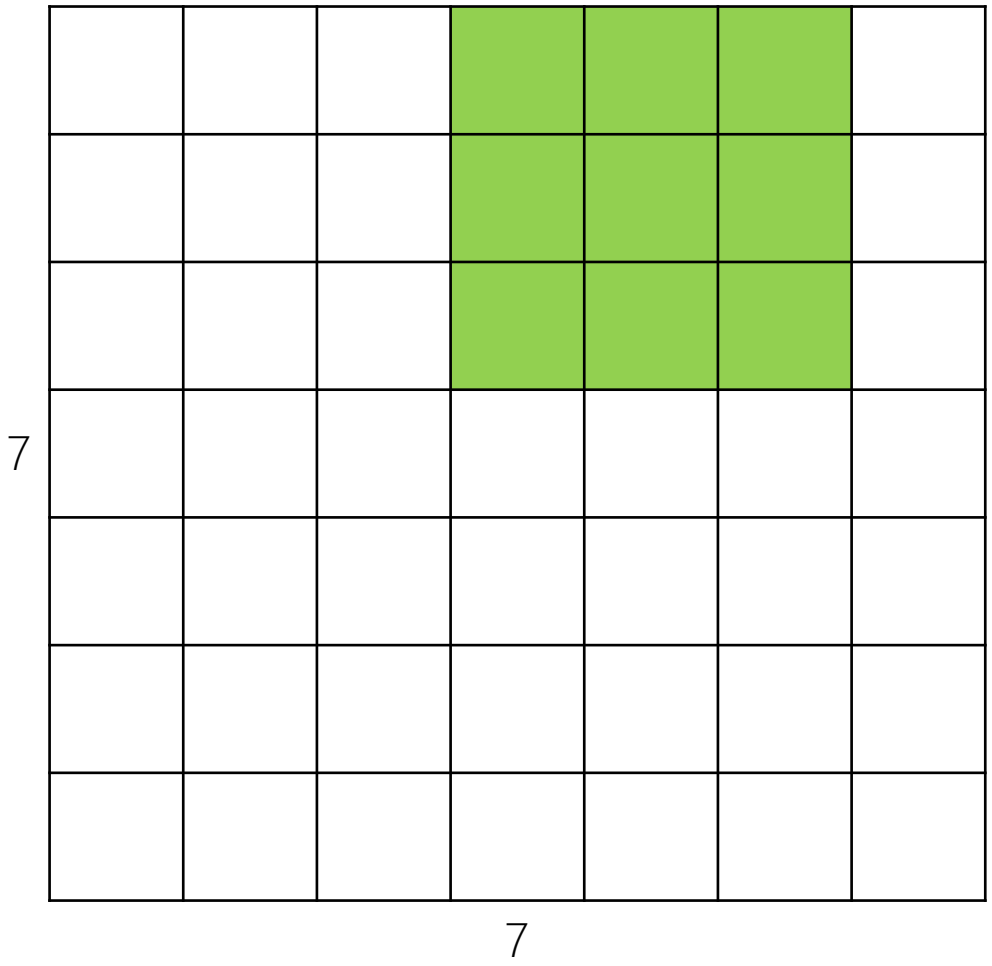

How about with stride 3?

Stride & Padding

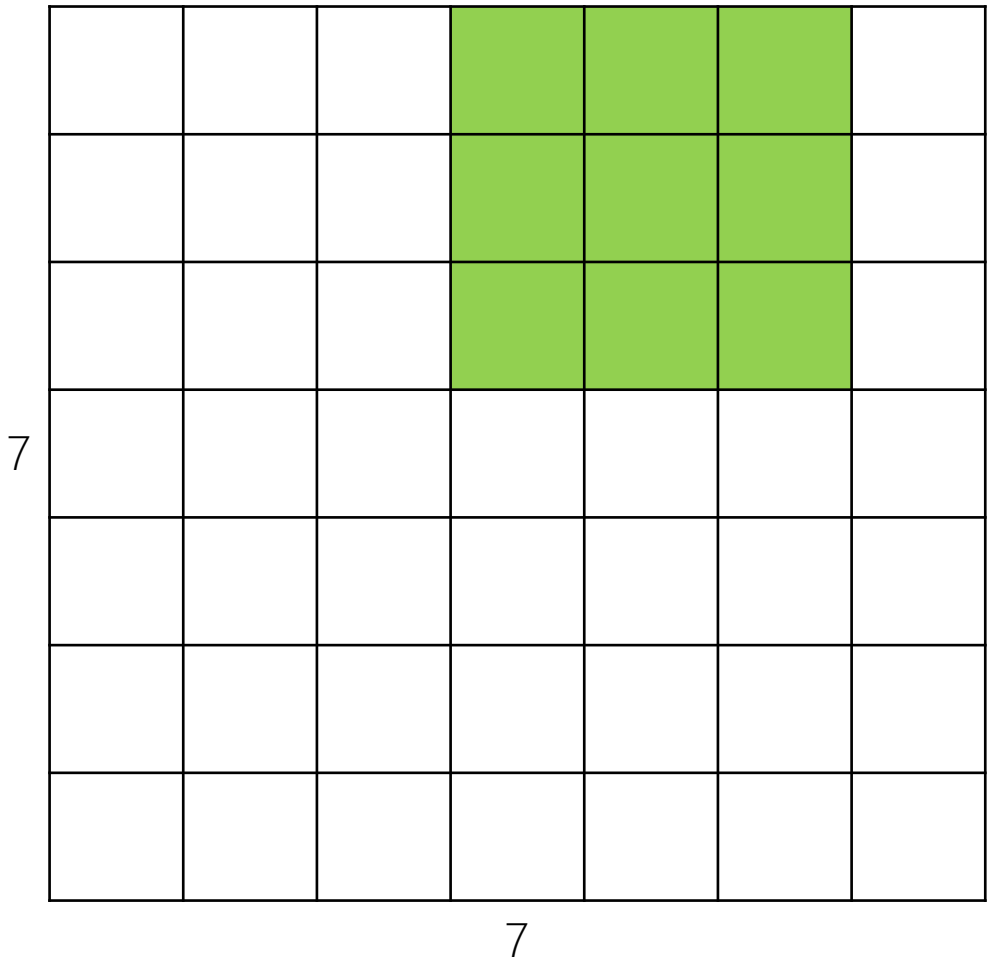

How about with stride 3?

→ …?

Stride & Padding

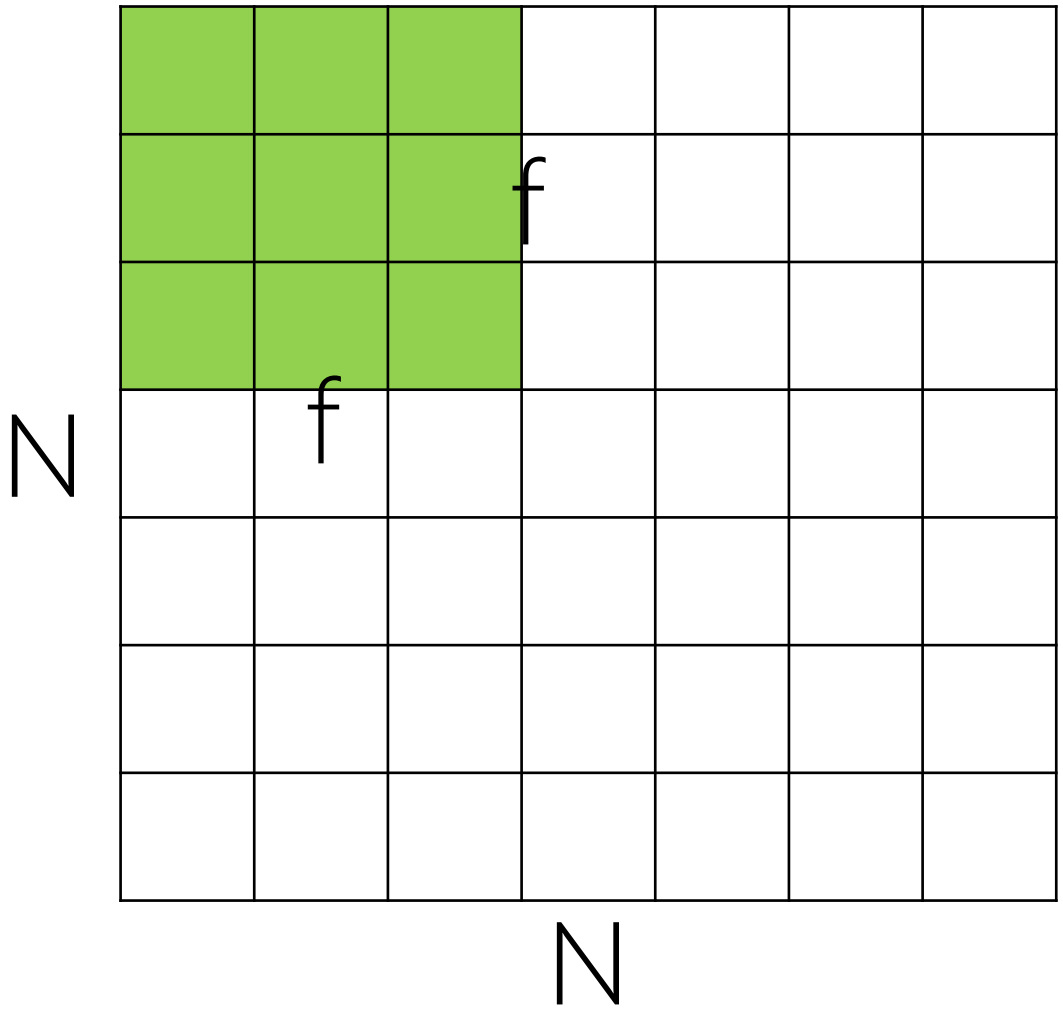

Output size:
(N-f)/stride + 1

Example, N=7, f=3:

Stride 1: (7-3)/1 + 1 = 5
Stride 2: (7-3)/2 + 1 = 3
Stride 3: (7-3)/3 + 1 = 2.3333

Stride & Padding

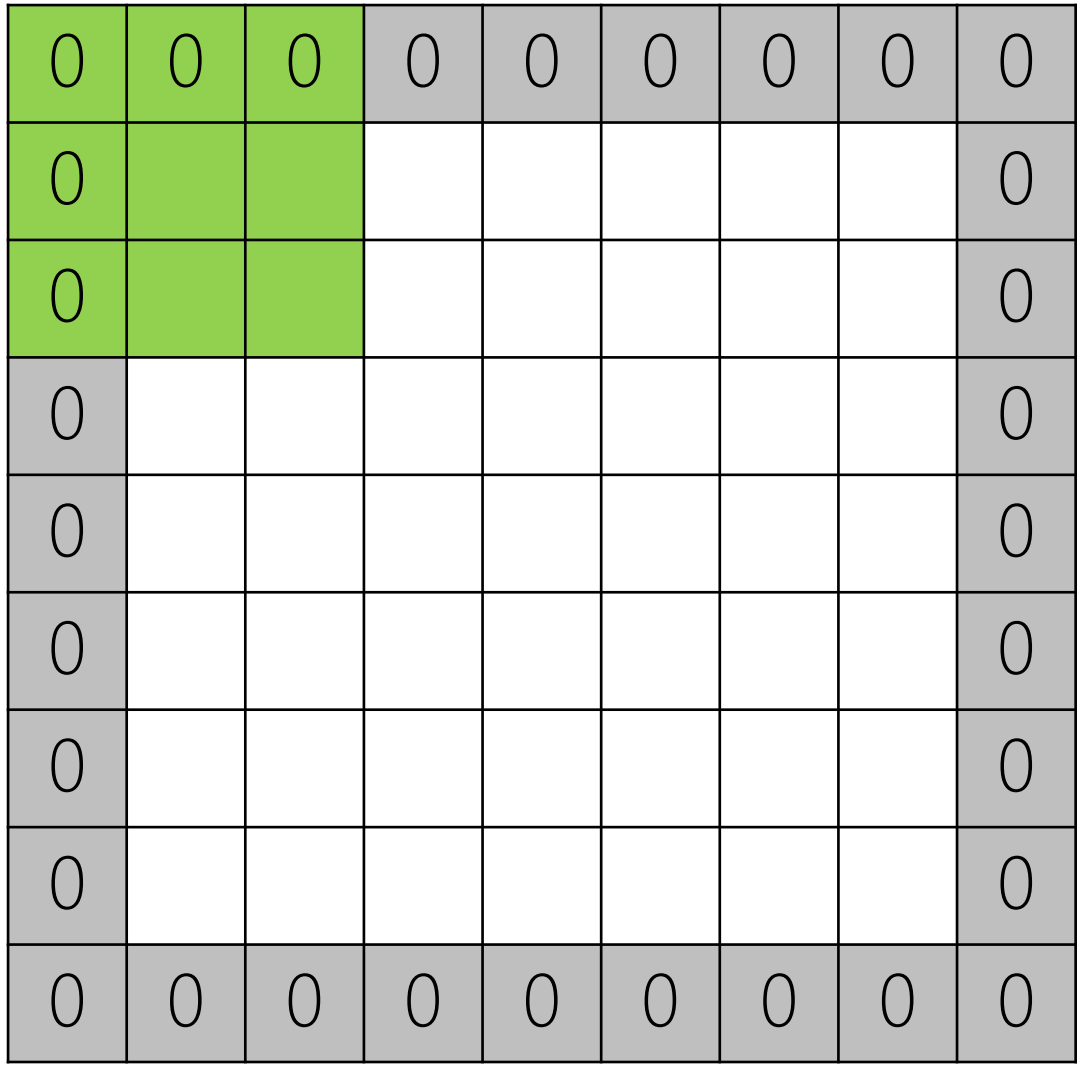

Output size:

(N-f)/stride + 1

Q. 7x7 input, 3x3 kernel, with stride 1, padded with 1 pixel. Output size?

Stride & Padding

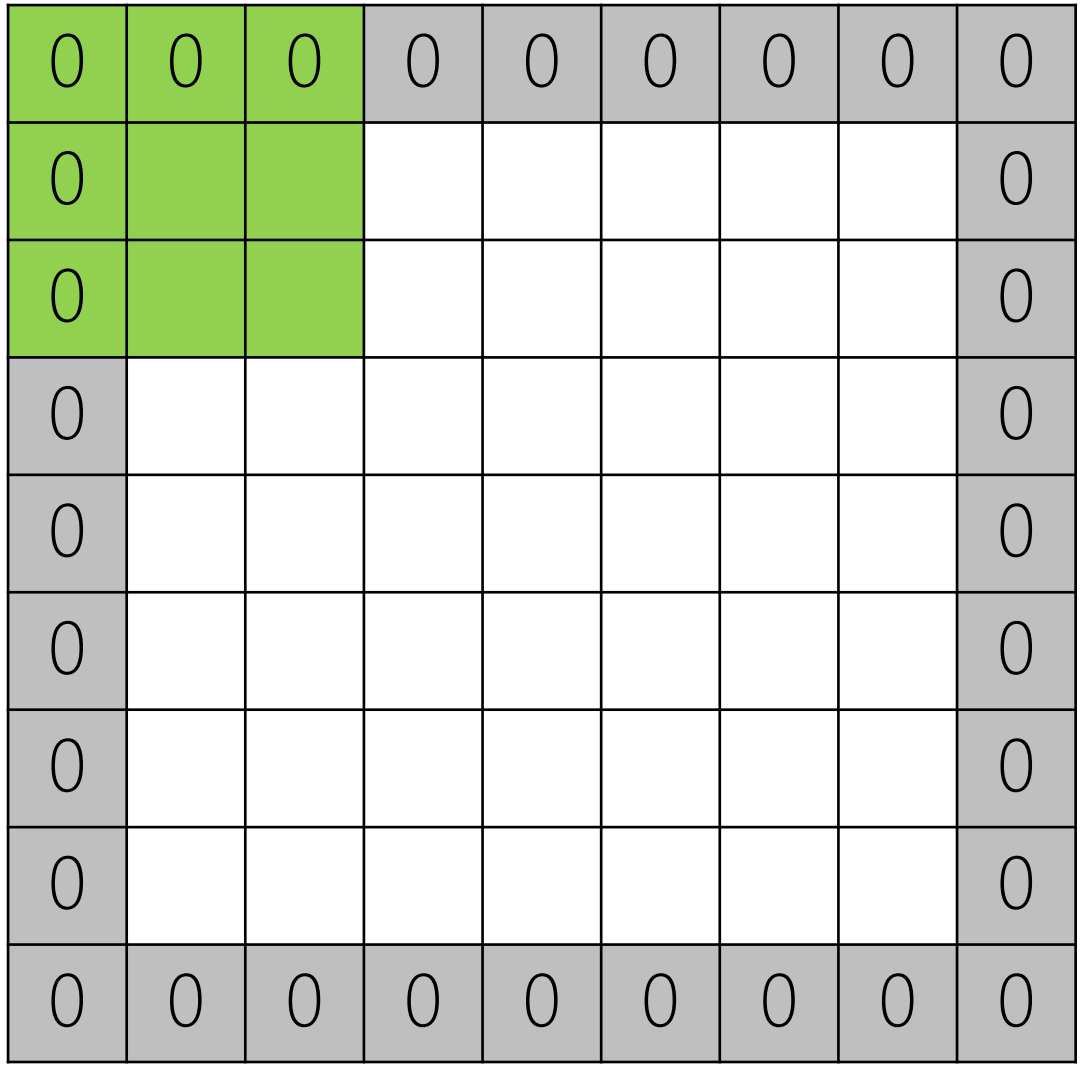

Output size:

(N-f)/stride + 1

Q. 7x7 input, 3x3 kernel, with stride 1, padded with 1 pixel. Output size?

→ 7x7

Stride & Padding

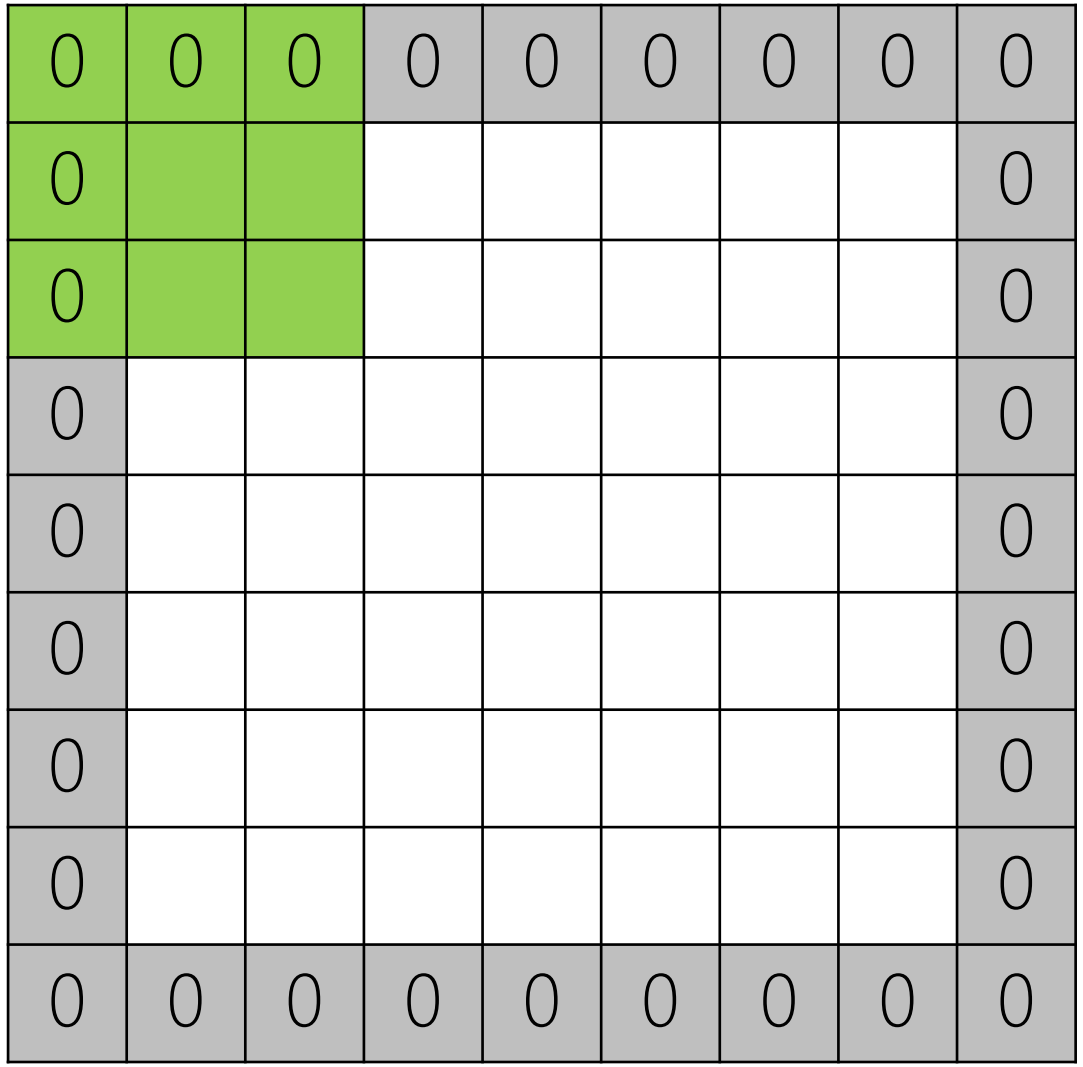

Output size:
(N-f)/stride + 1

Q. 7x7 input, 3x3 kernel, with stride 1, padded with 1 pixel. Output size?

→ 7x7

Q. 7x7 input, 3x3 kernel, with stride 3. You want to make 3x3 output, padding?

Stride & Padding

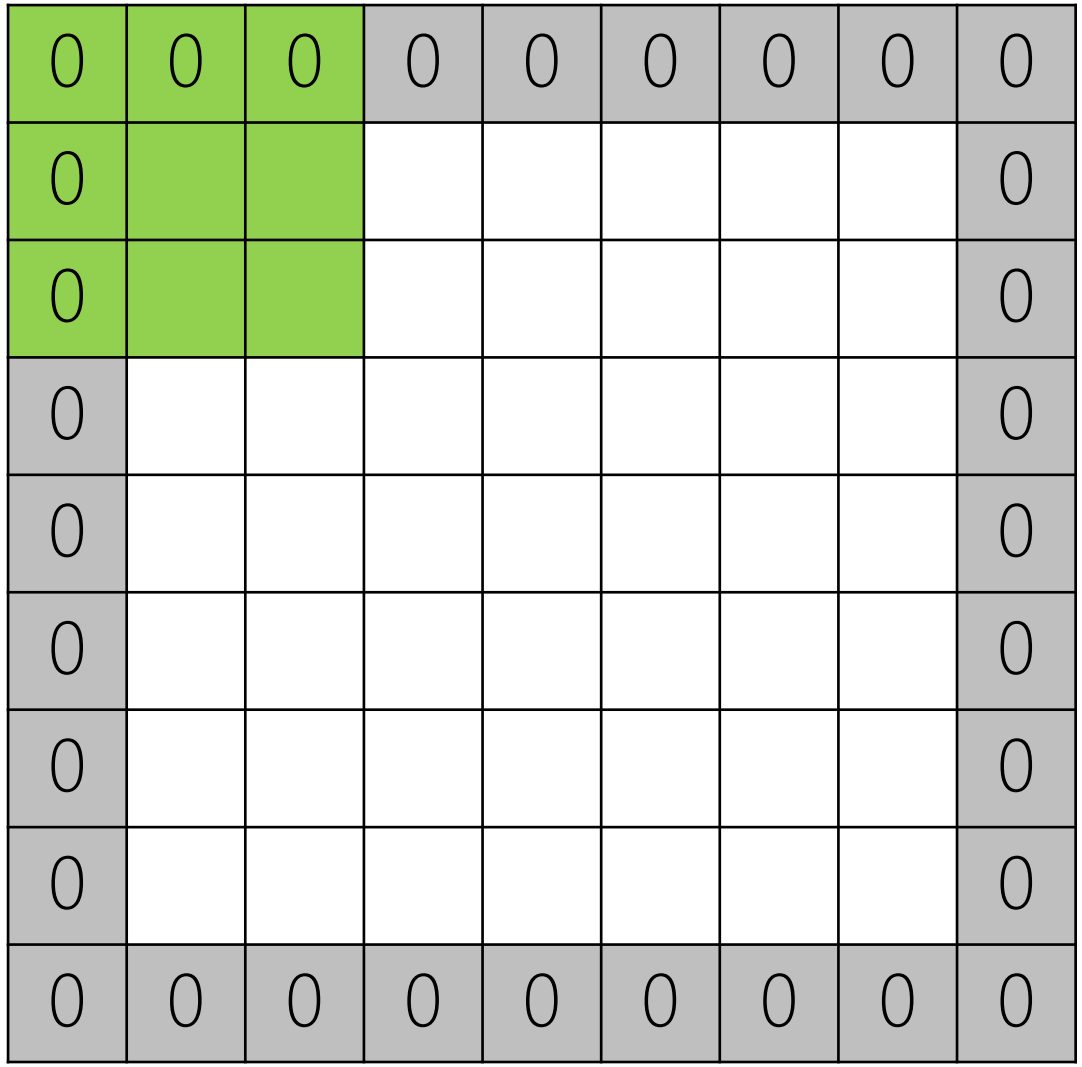

Output size:
(N-f)/stride + 1

Q. 7x7 input, 3x3 kernel, with stride 1, padded with 1 pixel. Output size?

→ 7x7

Q. 7x7 input, 3x3 kernel, with stride 3. You want to make 3x3 output, padding?

→ 1

Pooling Layers

- Pooling == ConvNet way of downsampling
- “Reduce the image size” to obtain smaller but more manageable features
- Receptive fields becomes relatively larger as the image size gets smaller
- Operates over each activation map independently

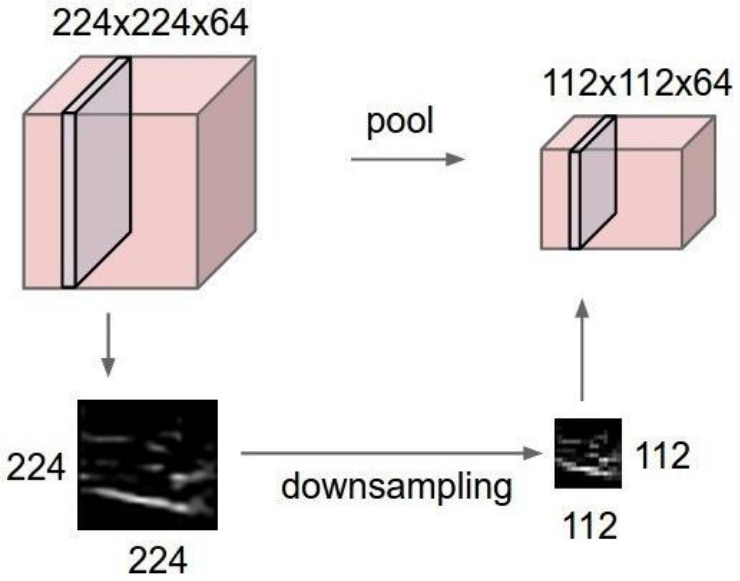

Image source: Stanford CS231n Convolutional Neural Networks for Visual Recognition

Max Pooling

- A non-linear down-sampling method
- An image is partitioned into a set of (non-overlapping) rectangles.
- The maximum of each such sub-region is sampled.

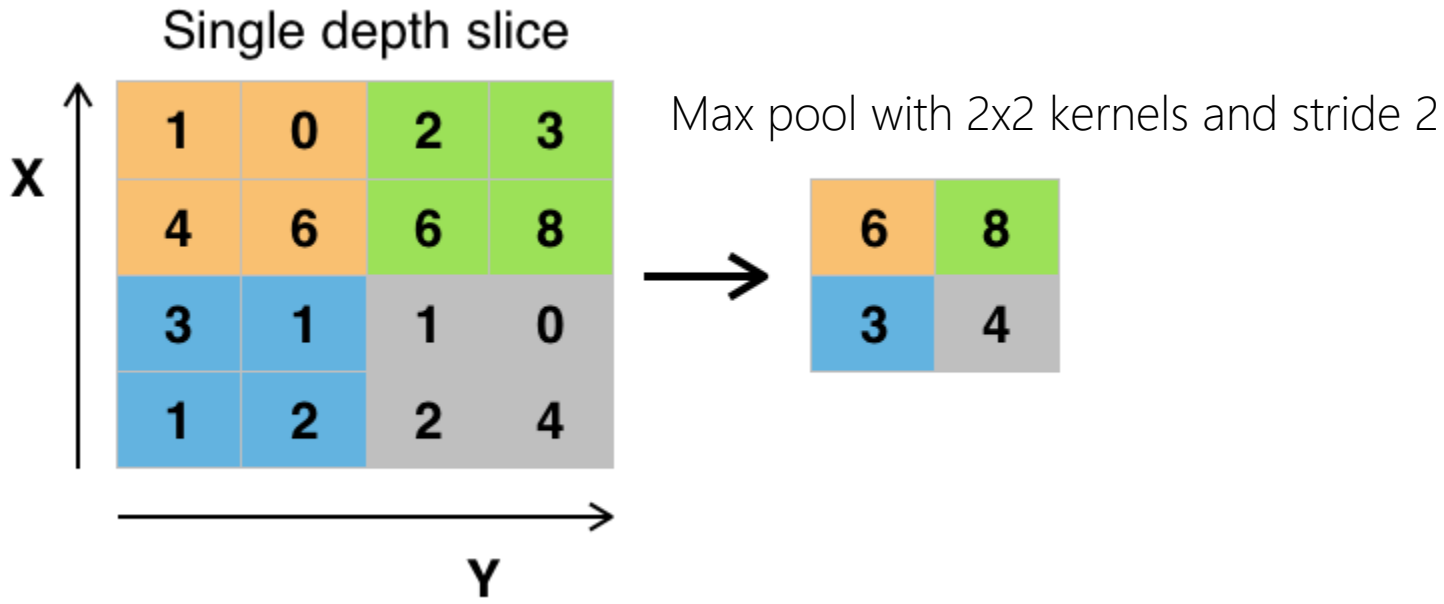

VGG Networks

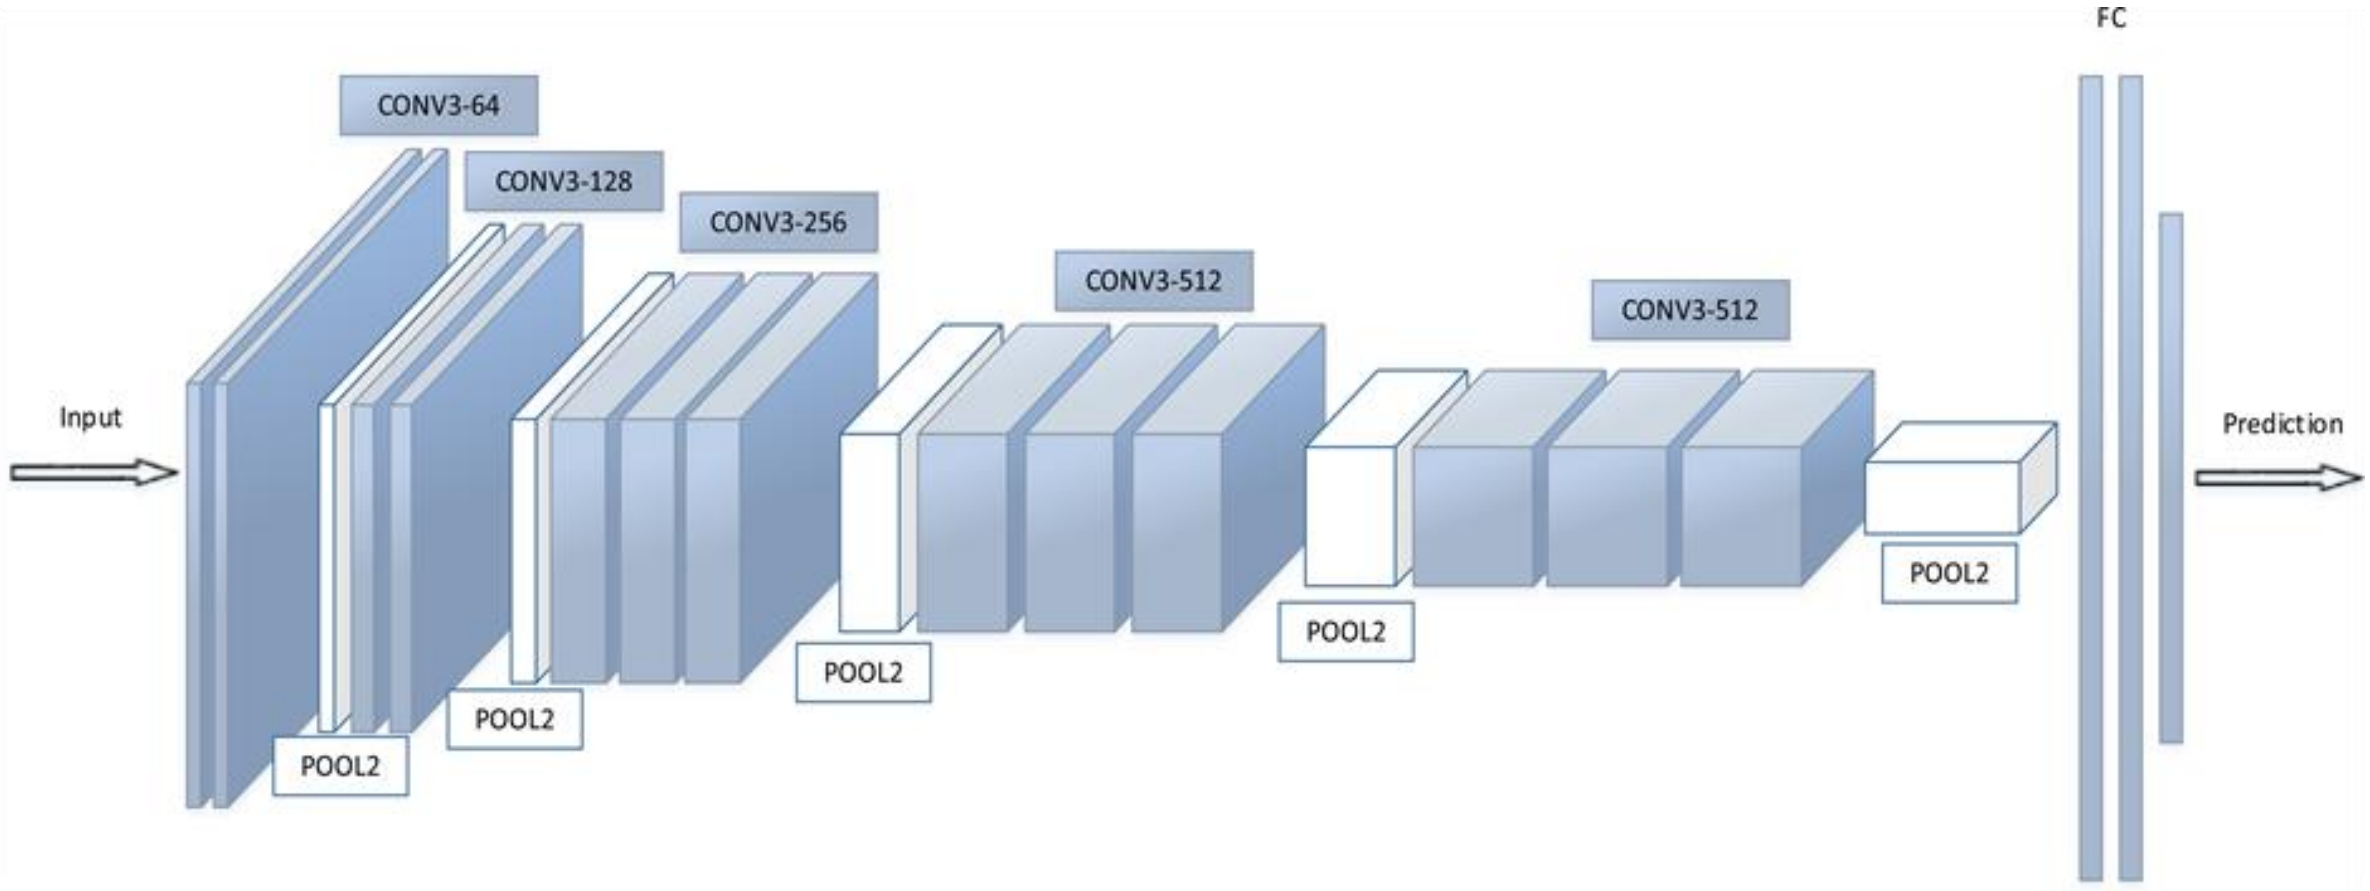

Abstraction of an Image

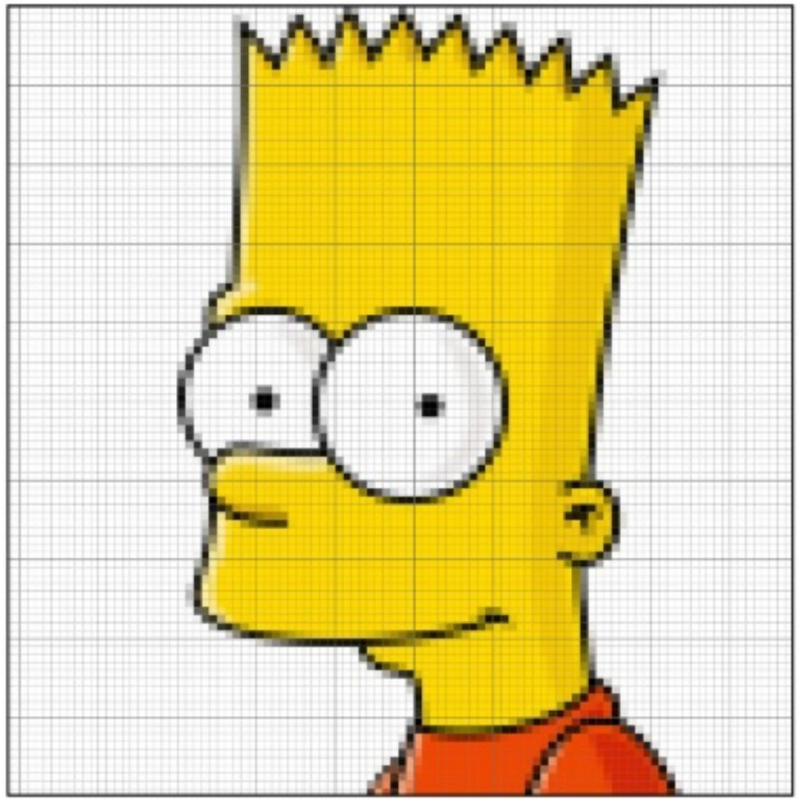

Conv2D + MaxPooling
+ Conv2D + MaxPooling
+ Conv2D + MaxPooling
+ ...

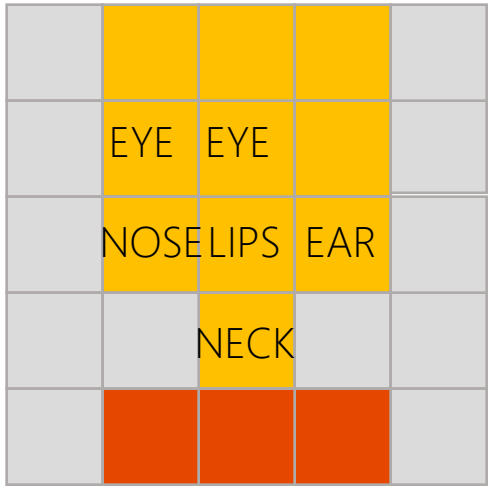

Abstraction of an Image

- Final decision is made by the fully-connected layers:

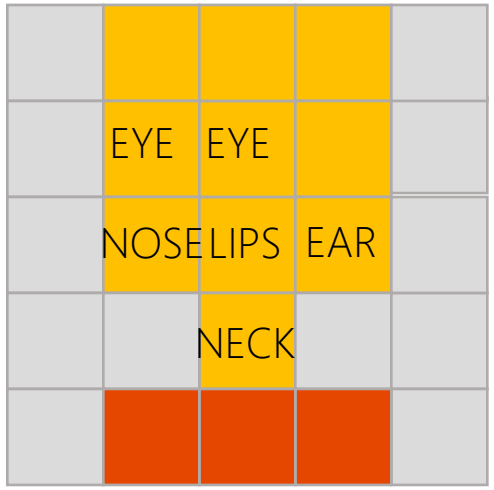

"Has two eyes, a nose, lips, an ear, and a neck and wears t-shirts"

Cat	X
Dog	X
Person	O
Orange	X