




Neural Operators

Stephen Baek

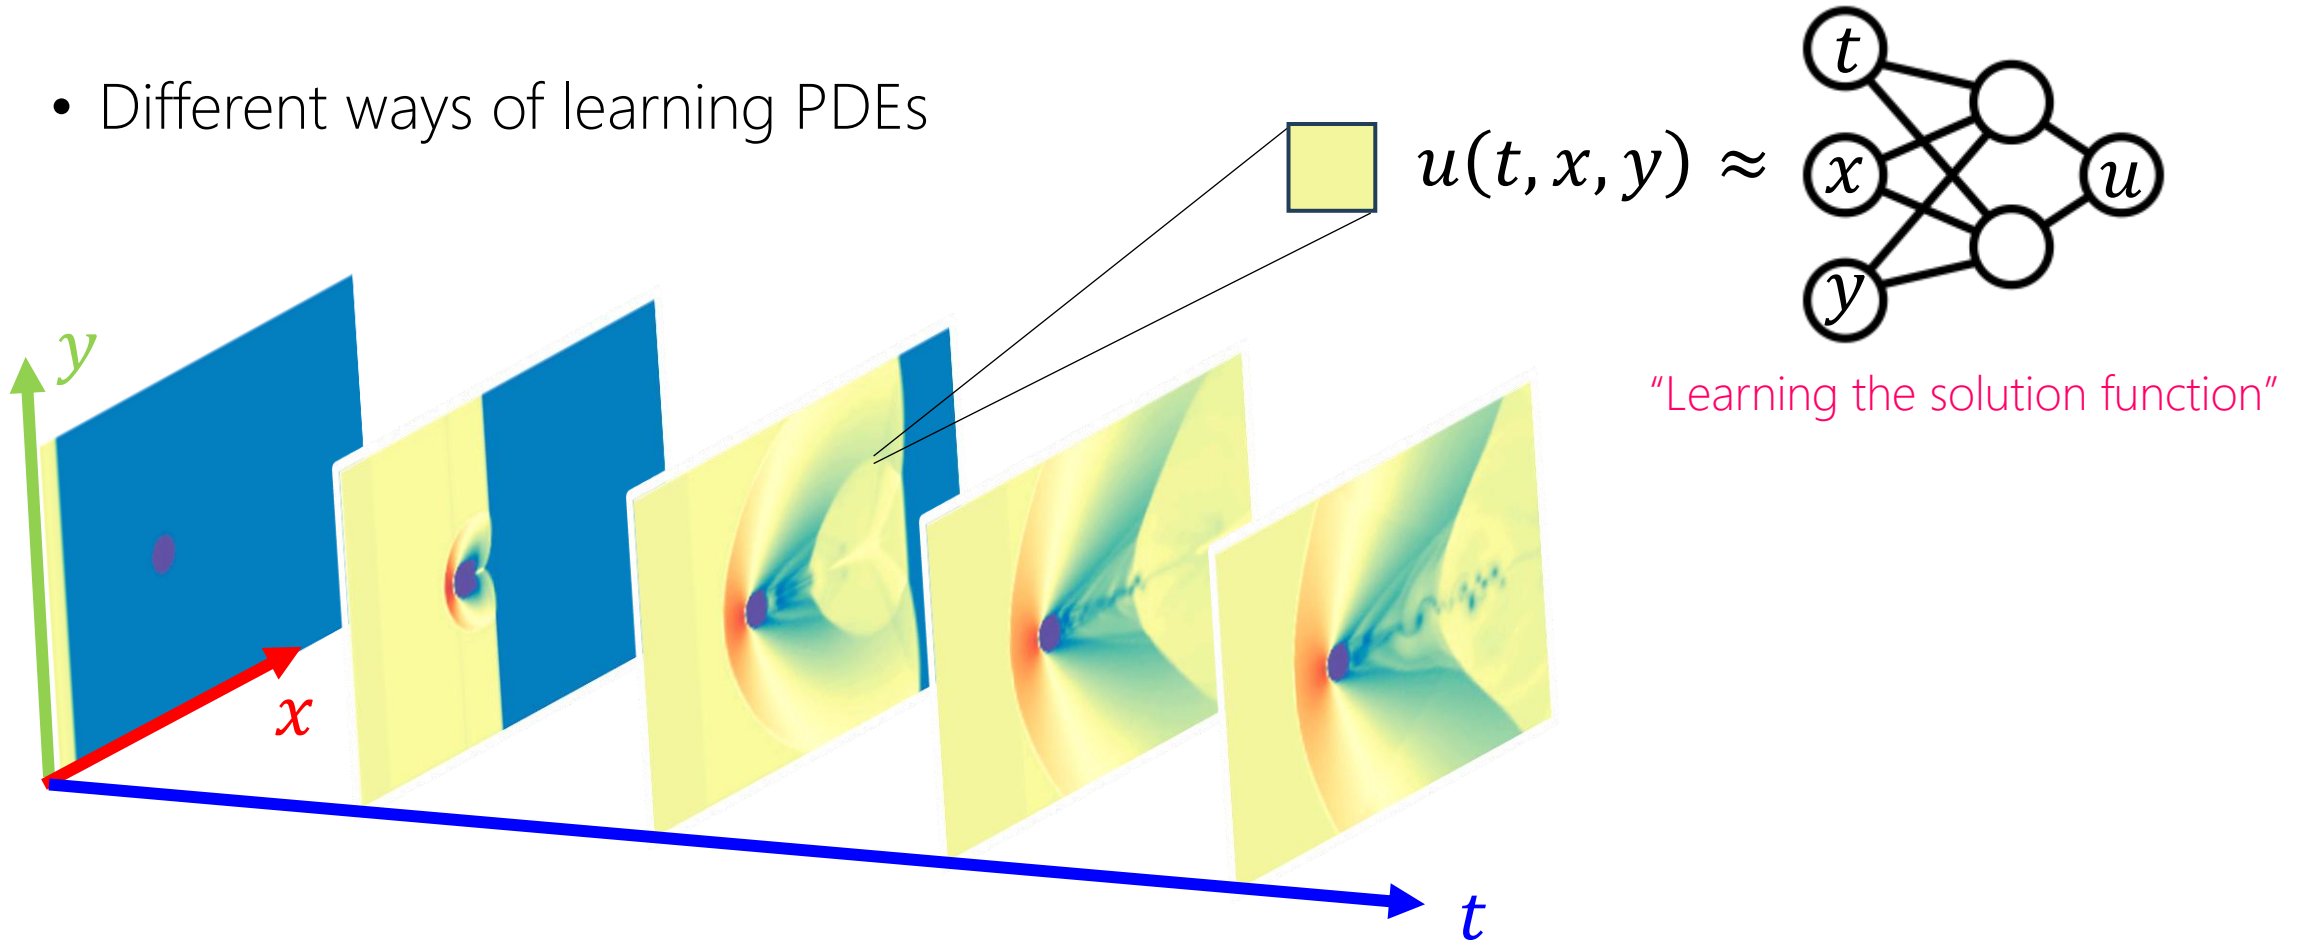
Machine Learning on Function Spaces

- Different ways of learning PDEs




Machine Learning on Function Spaces

- Different ways of learning PDEs



Machine Learning on Function Spaces

- Different ways of learning PDEs



Recap on Terminology

- Function
 - Mapping from value to value.
 - “The usual ones”
 - For instance,
 - $f(x) = \sin(x)$ maps a scalar to a scalar ($\sin: \mathbb{R} \rightarrow \mathbb{R}$)
 - Linear function $\mathbf{y} = A\mathbf{x} + \mathbf{b}$ maps a vector $\mathbf{x} \in \mathbb{R}^n$ to a vector $\mathbf{y} \in \mathbb{R}^m$
 - The solution function $u(t, x, y)$ of a PDE maps a given time t and location x, y to a physical quantity u .

Recap on Terminology

- Function
 - Mapping from **value to value**
 - “The usual ones”
 - For instance,
 - $f(x) = \sin(x)$ maps a scalar to a scalar ($\sin: \mathbb{R} \rightarrow \mathbb{R}$)
 - Linear function $\mathbf{y} = A\mathbf{x} + \mathbf{b}$ maps a vector $\mathbf{x} \in \mathbb{R}^n$ to a vector $\mathbf{y} \in \mathbb{R}^m$
 - The solution function $u(t, \mathbf{x}, \mathbf{y})$ of a PDE maps a given time t and location \mathbf{x}, \mathbf{y} to a physical quantity u .
- Operator
 - Mapping from **function to function**.
 - May not be familiar with the term, but you may already know it.
 - For instance,
 - Differential operator d/dx maps the function $f(x) = 3x^2 + 2$ to the function $f'(x) = 6x$.
 - Integral operator \int maps the function $\sin(x)$ to the function $-\cos(x) + C$, $C \in \mathbb{R}$.
 - Convolution operator T maps a function $f(x)$ to $(Tf)(x) = \int_{-\infty}^{\infty} f(y)g(x-y)dy$.

Ex: Solution Operator of ODE

- Second order systems, e.g., homogeneous spring-mass-damper system:
$$m\ddot{x}(t) + c\dot{x}(t) + kx(t) = 0$$

Ex: Solution Operator of ODE

- Second order systems, e.g., homogeneous spring-mass-damper system:

$$m\ddot{x}(t) + c\dot{x}(t) + kx(t) = 0$$

- Convert to first-order system:

$$\frac{d}{dt}\mathbf{u}(t) = \mathbf{A}\mathbf{u}(t)$$

- ... where: $\mathbf{u}(t) = \begin{bmatrix} x(t) \\ \dot{x}(t) \end{bmatrix}$, $\mathbf{A} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{c}{m} \end{bmatrix}$

Ex: Solution Operator of ODE

- Second order systems, e.g., homogeneous spring-mass-damper system:

$$m\ddot{x}(t) + c\dot{x}(t) + kx(t) = 0$$

- Convert to first-order system:

$$\frac{d}{dt}\mathbf{u}(t) = \mathbf{A}\mathbf{u}(t)$$

- ... where: $\mathbf{u}(t) = \begin{bmatrix} x(t) \\ \dot{x}(t) \end{bmatrix}$, $\mathbf{A} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{c}{m} \end{bmatrix}$

- Solution to IVP: $\mathbf{u}(t) = e^{t\mathbf{A}}\mathbf{u}(0)$

Ex: Solution Operator of ODE

- Second order systems, e.g., homogeneous spring-mass-damper system:

$$m\ddot{x}(t) + c\dot{x}(t) + kx(t) = 0$$

- Convert to first-order system:

$$\frac{d}{dt}\mathbf{u}(t) = \mathbf{A}\mathbf{u}(t)$$

- ... where: $\mathbf{u}(t) = \begin{bmatrix} x(t) \\ \dot{x}(t) \end{bmatrix}$, $\mathbf{A} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{c}{m} \end{bmatrix}$

- Solution to IVP: $\mathbf{u}(t) = e^{t\mathbf{A}}\mathbf{u}(0)$

Solution operator

Formal Definition

- Generic family of PDEs:

$$\begin{aligned} (L_a u)(x) &= f(x), & x \in D, \\ u(x) &= 0, & x \in \partial D \end{aligned}$$

- $a \in \mathcal{A}$ where \mathcal{A} is the Banach space* of physics parameters
- $u: D \rightarrow \mathbb{R}$: solution function that lives in the Banach space \mathcal{U}
- $f: D \rightarrow \mathbb{R}$: source term in the dual \mathcal{U}^* of \mathcal{U} .
- $D \subset \mathbb{R}^d$: a bounded physics domain
- $L_a: \mathcal{A} \rightarrow \mathcal{L}(\mathcal{U}; \mathcal{U}^*)$: mapping from the parameter Banach space \mathcal{A} to the space of linear operators that map \mathcal{U} to its dual \mathcal{U}^*

*Banach Space

= complete normed vector space

We have a way to measure the "length" and "distance" of arrows.

If you have a sequence of arrows v_1, v_2, \dots that get closer and closer to some limit $v = \lim v_i$, the limit v is also an arrow in the same vector space.

Formal Definition

- Generic family of PDEs:

$$\begin{aligned} (L_a u)(x) &= f(x), & x \in D, \\ u(x) &= 0, & x \in \partial D \end{aligned}$$

- $a \in \mathcal{A}$ where \mathcal{A} is the Banach space* of physics parameters
- $u: D \rightarrow \mathbb{R}$: solution function that lives in the Banach space \mathcal{U}
- $f: D \rightarrow \mathbb{R}$: source term in the dual \mathcal{U}^* of \mathcal{U} .
- $D \subset \mathbb{R}^d$: a bounded physics domain
- $L_a: \mathcal{A} \rightarrow \mathcal{L}(\mathcal{U}; \mathcal{U}^*)$: mapping from the parameter Banach space \mathcal{A} to the space of linear operators that map \mathcal{U} to its dual \mathcal{U}^*
- Solution operator to this PDE $\mathcal{G}: \mathcal{A} \times \mathcal{U}^* \rightarrow \mathcal{U}$

*Banach Space

= complete normed vector space

We have a way to measure the "length" and "distance" of arrows.

If you have a sequence of arrows v_1, v_2, \dots that get closer and closer to some limit $v = \lim v_i$, the limit v is also an arrow in the same vector space.

Formal Definition

- Generic family of PDEs:

$$\begin{aligned} (L_a u)(x) &= f(x), & x \in D, \\ u(x) &= 0, & x \in \partial D \end{aligned}$$

- $a \in \mathcal{A}$ where \mathcal{A} is the Banach space* of physics parameters
- $u: D \rightarrow \mathbb{R}$: solution function that lives in the Banach space \mathcal{U}
- $f: D \rightarrow \mathbb{R}$: source term in the dual \mathcal{U}^* of \mathcal{U} .
- $D \subset \mathbb{R}^d$: a bounded physics domain
- $L_a: \mathcal{A} \rightarrow \mathcal{L}(\mathcal{U}; \mathcal{U}^*)$: mapping from the parameter Banach space \mathcal{A} to the space of linear operators that map \mathcal{U} to its dual \mathcal{U}^*
- Solution operator to this PDE $\mathcal{G}: \mathcal{A} \times \mathcal{U}^* \rightarrow \mathcal{U}$
- Parametrize \mathcal{G} with a neural network! i.e. $\mathcal{G} \approx \mathcal{G}_\theta$
(universal approximation theorem of operators says you can do it**)

*Banach Space


= complete normed vector space

We have a way to measure the "length" and "distance" of arrows.

If you have a sequence of arrows v_1, v_2, \dots that get closer and closer to some limit $v = \lim v_i$, the limit v is also an arrow in the same vector space.

**See Kovachki et al. (2021) for more details.

Learning Solution Operators



Universal Approximation Theorem for Operators

- (Chen & Chen, 1995) Suppose that σ is a continuous non-polynomial function, X is a Banach space, $K_1 \subset X$, $K_2 \subset \mathbb{R}^d$ are two compact sets in X and \mathbb{R}^d , respectively, V is a compact set in $C(K_1)$, G is a continuous operator, which maps V into $C(K_2)$.

Then for any $\epsilon > 0$, there are positive integers n, p, m , constants $c_i^k, \xi_{ij}^k, \theta_i^k, \zeta_k \in \mathbb{R}$, $w_k \in \mathbb{R}^d$, $x_j \in K_1$, $i = 1, \dots, n$, $k = 1, \dots, p$, $j = 1, \dots, m$, such that

$$\left| G(u)(y) - \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k u(x_j) + \theta_i^k \right) \sigma(w_k \cdot y + \zeta_k) \right| < \epsilon$$

holds for all $u \in V$ and $y \in K_2$.

Universal Approximation Theorem for Operators

- (Chen & Chen, 1995) Suppose that σ is an activation function, X is a Banach space, $K_1 \subset X$, $K_2 \subset \mathbb{R}^d$ are two compact sets in X and \mathbb{R}^d , respectively, V is a compact set in $C(K_1)$, \mathcal{G} is a continuous operator, which maps V into $C(K_2)$.

Then for any $\epsilon > 0$, there are positive integers n, p, m , constants $c_i^k, \xi_{ij}^k, \theta_i^k, \zeta_k \in \mathbb{R}$, $w_k \in \mathbb{R}^d$, $x_j \in K_1$, $i = 1, \dots, n$, $k = 1, \dots, p$, $j = 1, \dots, m$, such that

$$\left| \mathcal{G}(u)(y) - \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k u(x_j) + \theta_i^k \right) \sigma(w_k \cdot y + \zeta_k) \right| < \epsilon$$

holds for all $u \in V$ and $y \in K_2$.

Universal Approximation Theorem for Operators

- (Chen & Chen, 1995) Suppose that σ is an activation function, K_1 is an input physical domain (where the input physical quantities u are sampled), K_2 is an output physical domain (where the output physical quantities are sampled), V is a compact set in $C(K_1)$, \mathcal{G} is a continuous operator, which maps V into $C(K_2)$.

Then for any $\epsilon > 0$, there are positive integers n, p, m , constants $c_i^k, \xi_{ij}^k, \theta_i^k, \zeta_k \in \mathbb{R}$, $w_k \in \mathbb{R}^d$, $x_j \in K_1$, $i = 1, \dots, n$, $k = 1, \dots, p$, $j = 1, \dots, m$, such that

$$\left| \mathcal{G}(u)(y) - \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k u(x_j) + \theta_i^k \right) \sigma(w_k \cdot y + \zeta_k) \right| < \epsilon$$

holds for all $u \in V$ and $y \in K_2$.

Universal Approximation Theorem for Operators

- (Chen & Chen, 1995) Suppose that σ is an activation function, K_1 is an input physical domain (where the input physical quantities u are sampled), K_2 is an output physical domain (where the output physical quantities are sampled), V is a set of continuous functions (physical quantities u) defined on K_1 , \mathcal{G} is a continuous operator, which maps V into $\mathcal{C}(K_2)$.

Then for any $\epsilon > 0$, there are positive integers n, p, m , constants $c_i^k, \xi_{ij}^k, \theta_i^k, \zeta_k \in \mathbb{R}$, $w_k \in \mathbb{R}^d$, $x_j \in K_1$, $i = 1, \dots, n$, $k = 1, \dots, p$, $j = 1, \dots, m$, such that

$$\left| \mathcal{G}(u)(y) - \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k u(x_j) + \theta_i^k \right) \sigma(w_k \cdot y + \zeta_k) \right| < \epsilon$$

holds for all $u \in V$ and $y \in K_2$.

Universal Approximation Theorem for Operators

- (Chen & Chen, 1995) Suppose that σ is an activation function, K_1 is an input physical domain (where the input physical quantities u are sampled), K_2 is an output physical domain (where the output physical quantities are sampled), V is a set of continuous functions (physical quantities u) defined on K_1 , \mathcal{G} is a **solution operator**, which **maps input physical quantities u to output physical quantities evaluated at y** .

Then for any $\epsilon > 0$, there are positive integers n, p, m , constants $c_i^k, \xi_{ij}^k, \theta_i^k, \zeta_k \in \mathbb{R}$, $w_k \in \mathbb{R}^d$, $x_j \in K_1$, $i = 1, \dots, n$, $k = 1, \dots, p$, $j = 1, \dots, m$, such that

$$\left| \mathcal{G}(u)(y) - \sum_{k=1}^p \sum_{i=1}^n c_i^k \sigma \left(\sum_{j=1}^m \xi_{ij}^k u(x_j) + \theta_i^k \right) \sigma(w_k \cdot y + \zeta_k) \right| < \epsilon$$

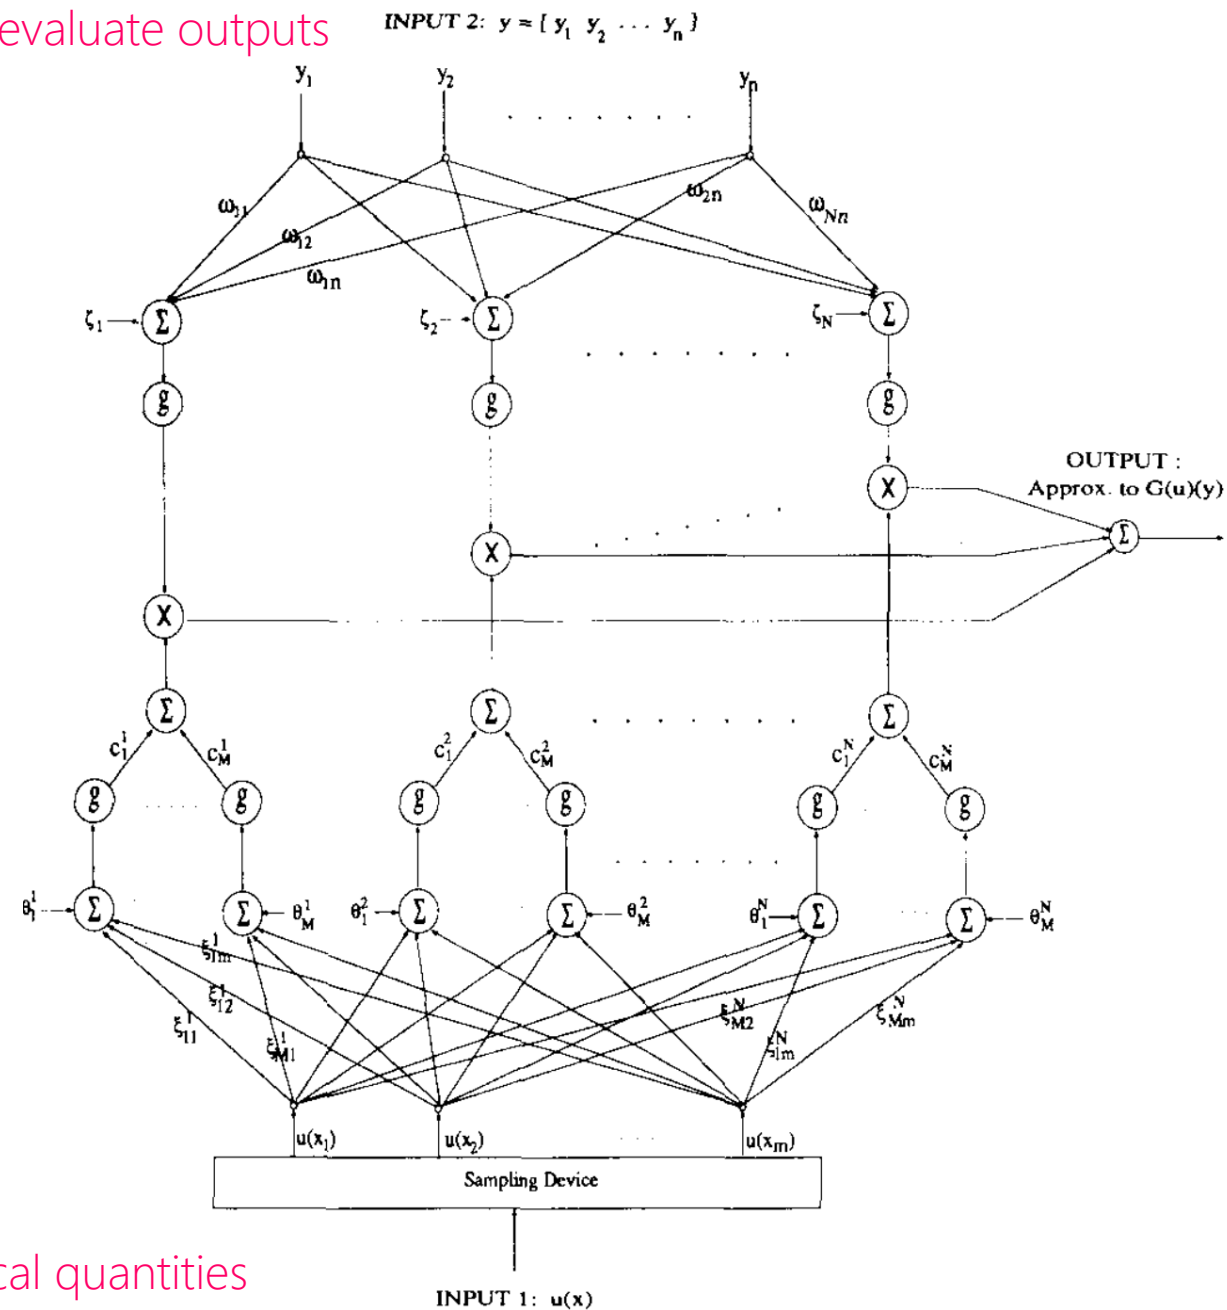
holds for all $u \in V$ and $y \in K_2$.

Universal Approximation Theorem for Operators

- (Chen & Chen, 1995) Suppose that σ is an activation function, K_1 is an input physical domain (where the input physical quantities \mathbf{u} are sampled), K_2 is an output physical domain (where the output physical quantities are sampled), V is a set of continuous functions (physical quantities \mathbf{u}) defined on K_1 , \mathcal{G} is a solution operator, which maps input physical quantities \mathbf{u} to output physical quantities evaluated at \mathbf{y} .

Then, there exists a neural network that looks like the one in the next slide that can approximate \mathcal{G} with an error bound ϵ .

Locations where you want to evaluate outputs




Input physical quantities

A modern twist: Deep-O-Net

Lu et al., *Nature Mach Intell*, 2021

$$\mathcal{G}(u)(y) \approx \sum_{k=1}^p \underbrace{b_k(u)}_{\text{branch}} \underbrace{t_k(y)}_{\text{trunk}}$$

D Unstacked DeepONet



In Simple Terms: Image-to-Image Problem

- “Operator learning can be taken as an image-to-image problem.”
- Zongyi Li website.
- ...where an “Image” could be:
 - A snapshot of the solution function at a given time t .
 - Initial condition/configuration of a physical system. (snapshot at $t=0$)
 - Boundary geometry.
 - Distribution of system coefficients.
 - Etc.

Image-to-Image Problem

- For instance, a solution operator:
 - Consider the Darcy equation describing the flow of a fluid through a porous medium:
$$-\operatorname{div}(a\nabla u) = f$$
 - ..., where
 - $u \in \mathcal{U}$: temperature or pressure
 - $a \in \mathcal{A}$: conductance or permeability
 - $f \in \mathcal{U}^*$: source term


$f = 0$
Assumed in
this example

×



a : Permeability coefficients

$\mathcal{G}: \mathcal{A} \times \mathcal{U}^* \rightarrow \mathcal{U}$
solution operator



u : Pressure distribution

Image-to-Image Problem


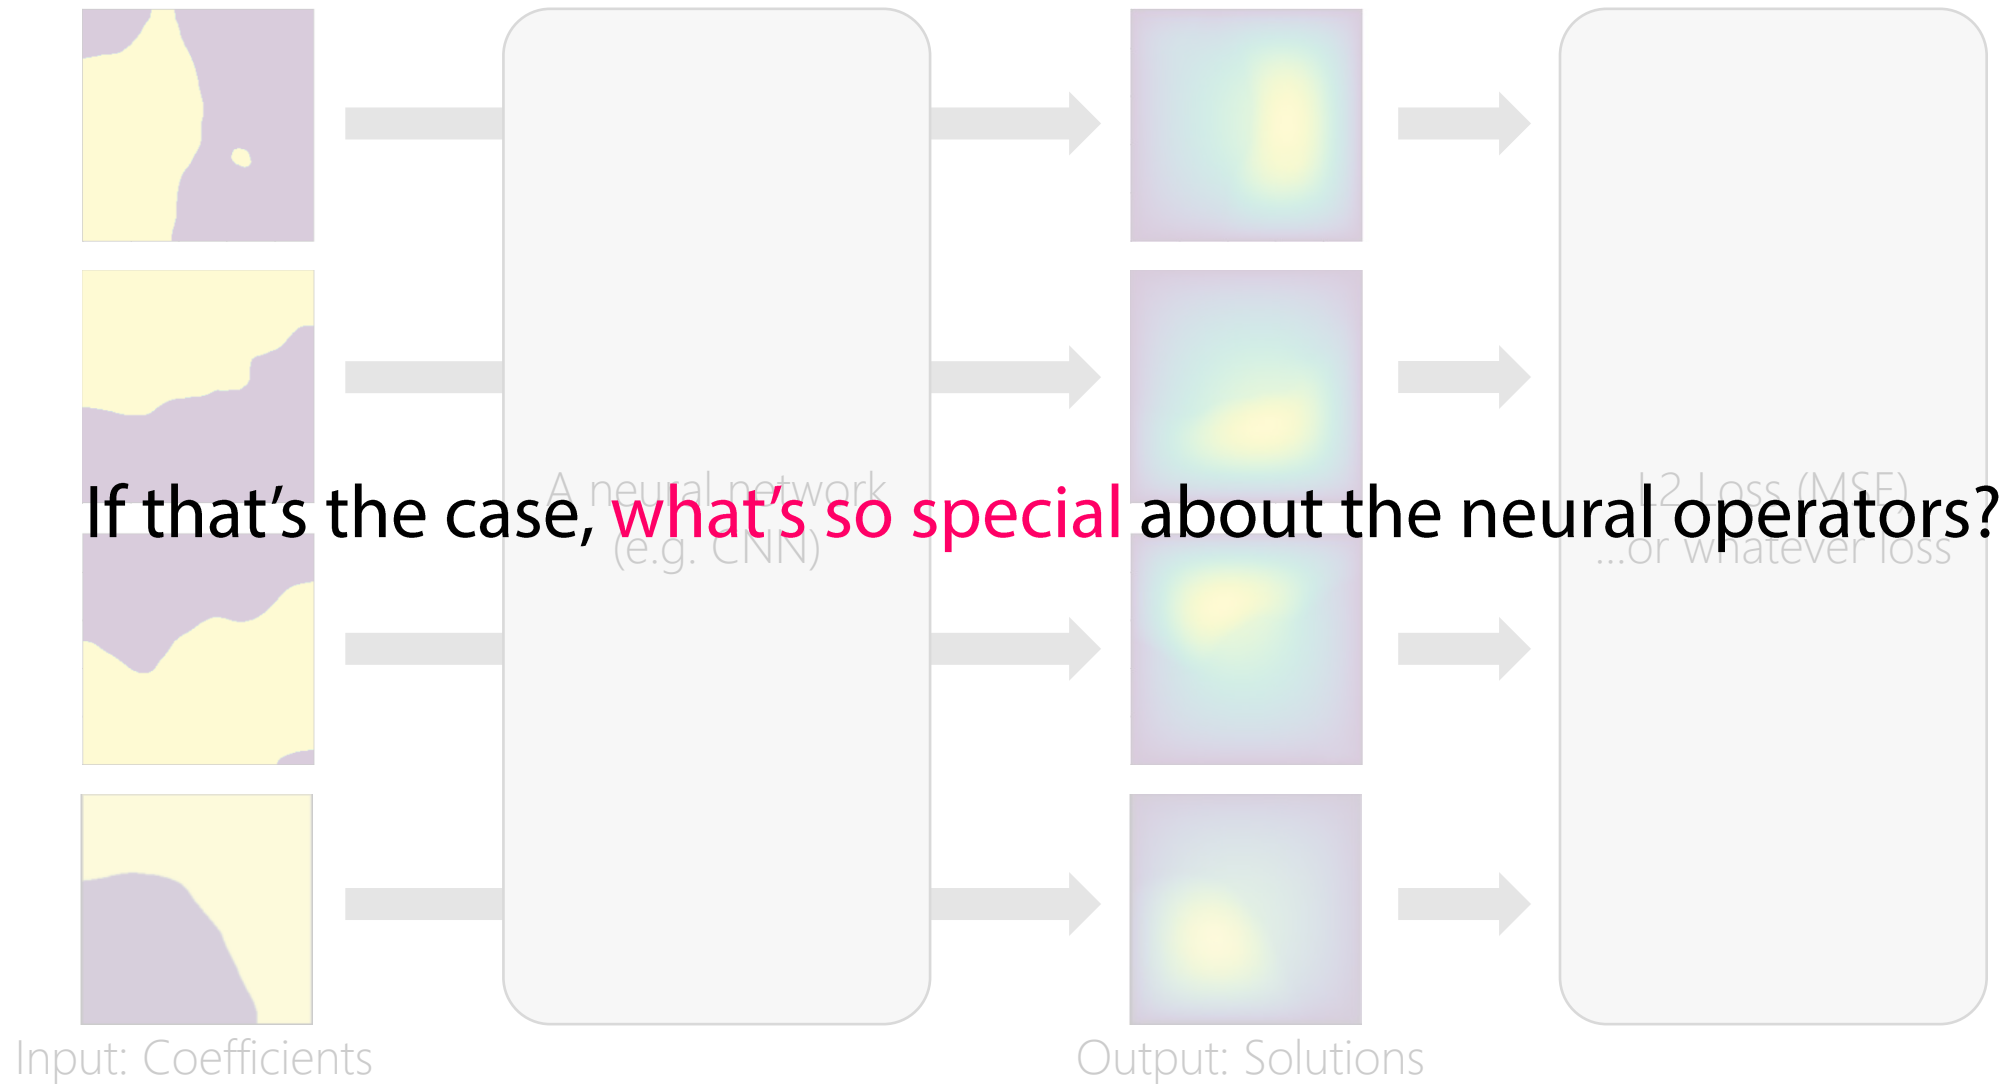



Image-to-Image Problem




Machine Learning on Functions




Our Dataset

...



...

Machine Learning on Functions




Discretization:

$$a_i := a(x_i)$$

$$u_i := u(y_i)$$


Machine Learning on Functions





Discretization:

$$a_i := a(x_i)$$

$$u_i := u(y_i)$$



Machine Learning on Functions




Discretization:

$$a_i := a(x_i)$$

$$u_i := u(y_i)$$

a_0	a_1	a_2	a_3	a_4	\dots	a_n
u_0	u_1	u_2	\dots	u_m		


Machine Learning on Functions



Discretization:


$$a_i := a(x_i)$$

$$u_i := u(y_i)$$




Linear neural network layer

$$u_j = \sigma \left(\sum_{i=0}^n K_{ji} a_i \right)$$



Machine Learning on Functions


What if the discretization changes?



K_{j0} K_{j1} K_{j2} K_{j3} K_{j4} \dots K_{jn}

→ Are these still useful?

Machine Learning on Functions




Instead, consider “indexing” the weight K with x_i and y_j

$$u_j = \sigma \left(\sum_{i=0}^n K_{ji} a_i \right)$$

➔

$$u(y_j) = \sigma \left(\sum_{i=0}^n K(y_j, x_i) a(x_i) \right)$$

Machine Learning on Functions



Instead, consider “indexing” the weight K with x_i and y_j


$$u_j = \sigma \left(\sum_{i=0}^n K_{ji} a_i \right)$$

➔

$$u(y_j) = \sigma \left(\sum_{i=0}^n K(y_j, x_i) a(x_i) \right)$$

$$= \sigma \left(\sum_{i=0}^n \underbrace{\kappa(y_j, x_i) a(x_i) \Delta x_i}_{\text{Riemann Sum}} \right)$$

Machine Learning on Functions



Instead, consider “indexing” the weight K with x_i and y_j

$$u_j = \sigma \left(\sum_{i=0}^n K_{ji} a_i \right)$$




$$u(y_j) = \sigma \left(\sum_{i=0}^n K(y_j, x_i) a(x_i) \right)$$

$$= \sigma \left(\sum_{i=0}^n \underbrace{\kappa(y_j, x_i) a(x_i) \Delta x_i}_{\text{Riemann Sum}} \right)$$



$$u(y) = \sigma \left(\int \kappa(y, x) a(x) dx \right)$$

Machine Learning on Functions



Instead, consider “indexing” the weight K with x_i and y_j

$$u_j = \sigma \left(\sum_{i=0}^n K_{ji} a_i \right)$$

Linear matrix operator acting on **fixed** discretization

$$u(y_j) = \sigma \left(\sum_{i=0}^n K(y_j, x_i) a(x_i) \right)$$



Linear integral operator acting on **continuous** functions

Riemann Sum



$$u(y) = \sigma \left(\int \kappa(y, x) a(x) dx \right)$$

Kernel function



Assuming we learned the kernel function...





$$u(y) = \sigma \left(\int \kappa(y, x) a(x) dx \right)$$




Assuming we learned the kernel function...




$$u(y) = \sigma \left(\int \kappa(y, x) a(x) dx \right)$$




Assuming we learned the kernel function...



Assuming we learned the kernel function...



Assuming we learned the kernel function...



Works great in theory, but how?

$$u(y) = \sigma \left(\int \kappa(y, x) a(x) dx \right)$$

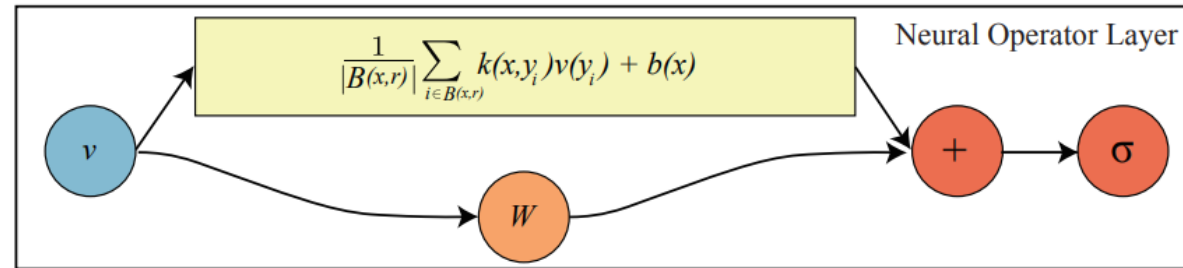
- Practically, we need to compute (in each neural network layer):
 - Continuous kernel function $\kappa(y, x)$
 - Continuous integral $\int \cdot dx$
- ...which are **computationally expensive**!



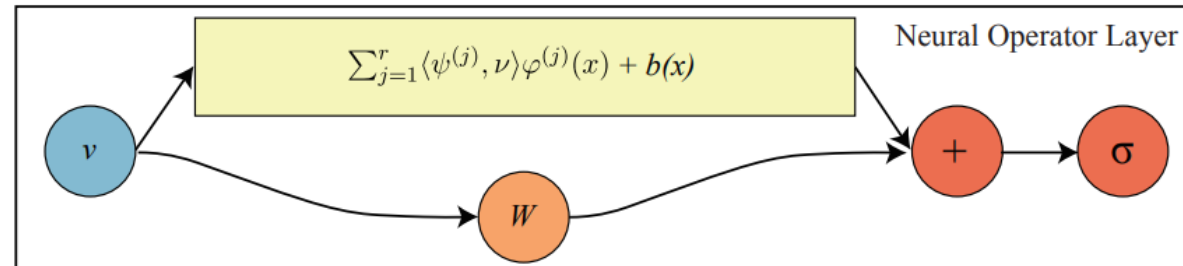
Works great in theory, but how?

- Kovachki, Li, et al. (2022) – Different versions of the linear integral operator

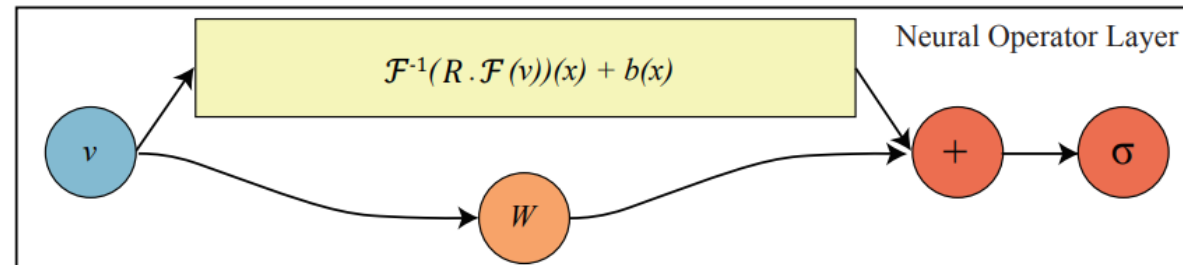
Graph Neural Operator



Low-rank Neural Operator



Fourier Neural Operator



Fourier Neural Operators

- Convolution Theorem:
 - The convolution $f * g$ of two functions f and g is equivalent to a simple element-wise (Hadamard) product \odot in the Fourier (spectral) domain*:
$$\mathcal{F}(f * g) = \mathcal{F}(f) \odot \mathcal{F}(g)$$

Fourier Neural Operators

- Convolution Theorem:
 - The convolution $f * g$ of two functions f and g is equivalent to a simple element-wise (Hadamard) product \odot in the Fourier (spectral) domain*:

$$\mathcal{F}(f * g) = \mathcal{F}(f) \odot \mathcal{F}(g)$$

- If the integral operator was the convolution operator, i.e.:

$$u(y) = \sigma \left(\int \kappa(y, x) a(x) dx \right) \rightarrow u(y) := \sigma((\kappa * a)(y)) = \sigma \left(\int \kappa(y - x) a(x) dx \right)$$

- Then by the convolution theorem:

$$(\kappa * a)(y) = \mathcal{F}^{-1}(R \odot \mathcal{F}(a))(y)$$


- $\mathcal{F}(a)$: Fourier transform of the input function $a(x)$
 - $R := \mathcal{F}(\kappa)$: Kernel function learned as Fourier coefficients (no need to evaluate κ directly)

Fourier Neural Operators

- FNO Layer

1. Fourier transform $\mathcal{F}(a)$
2. Linear transform R
3. Inverse Fourier transform \mathcal{F}^{-1}

$$\mathcal{F}^{-1}(R \odot \mathcal{F}(a))(y)$$




For a uniform grid, fast Fourier transform is available. Otherwise, discrete Fourier transform (slow).

Fourier Neural Operators

- The Fourier Filters



Filters in CNN




Fourier Filters


Fourier Neural Operators

- The full architecture of FNO

(a)




(b)



W : Local (point-wise) linear operator (skip connection)
to "keep the track of the location information and non-periodic boundary"*

Graph Neural Operators

- Message Passing Neural Networks (Gilmer et al. 2017)
 - A graph G with node features $\mathbf{v}_i \in V$ and edge features $\mathbf{e}_{i,j} \in E$.





Graph Neural Operators

- Message Passing Neural Networks (Gilmer et al. 2017)
 - A graph G with node features $\mathbf{v}_i \in V$ and edge features $\mathbf{e}_{i,j} \in E$.
 - Message Passing Phase—Runs for K time steps. Update hidden states $\mathbf{h}_i^{(k)}$ at each node based on messages $\mathbf{m}_i^{(k+1)}$ according to:

$$\mathbf{m}_i^{(k)} = \sum_{j \in N(i)} M^{(k)}(\mathbf{h}_i^{(k-1)}, \mathbf{h}_j^{(k-1)}, \mathbf{e}_{i,j})$$

$$\mathbf{h}_i^{(k)} = U^{(k)}(\mathbf{h}_i^{(k-1)}, \mathbf{m}_i^{(k)})$$

..., where $M^{(k)}$ is a message function (typically a neural net layer), $U^{(k)}$ is vertex update function (again, a neural network layer), and $N(i)$ is the neighbors of \mathbf{v}_i in G .



Graph Neural Operator

- Graph Neural Operator using Message Passing Graph Networks

$$v_{t+1}(x) = \sigma \left(W v_t(x) + \frac{1}{|N(x)|} \sum_{y \in N(x)} \kappa_{\phi}(e(x, y)) v_t(y) \right)$$

- $\kappa_{\phi}(e(x, y))$: A message passing neural network taking edge features as input.
- Edge features $e(x, y)$: For example, the edge weight can be defined as

$$e(x, y) = (x, y, a(x), a(y))$$


Positions

Coefficients/physics parameters at those positions


Examples & Use Cases

- Supervised learning of dynamics


Initial Vorticity




t=15




t=20







t=25



t=30

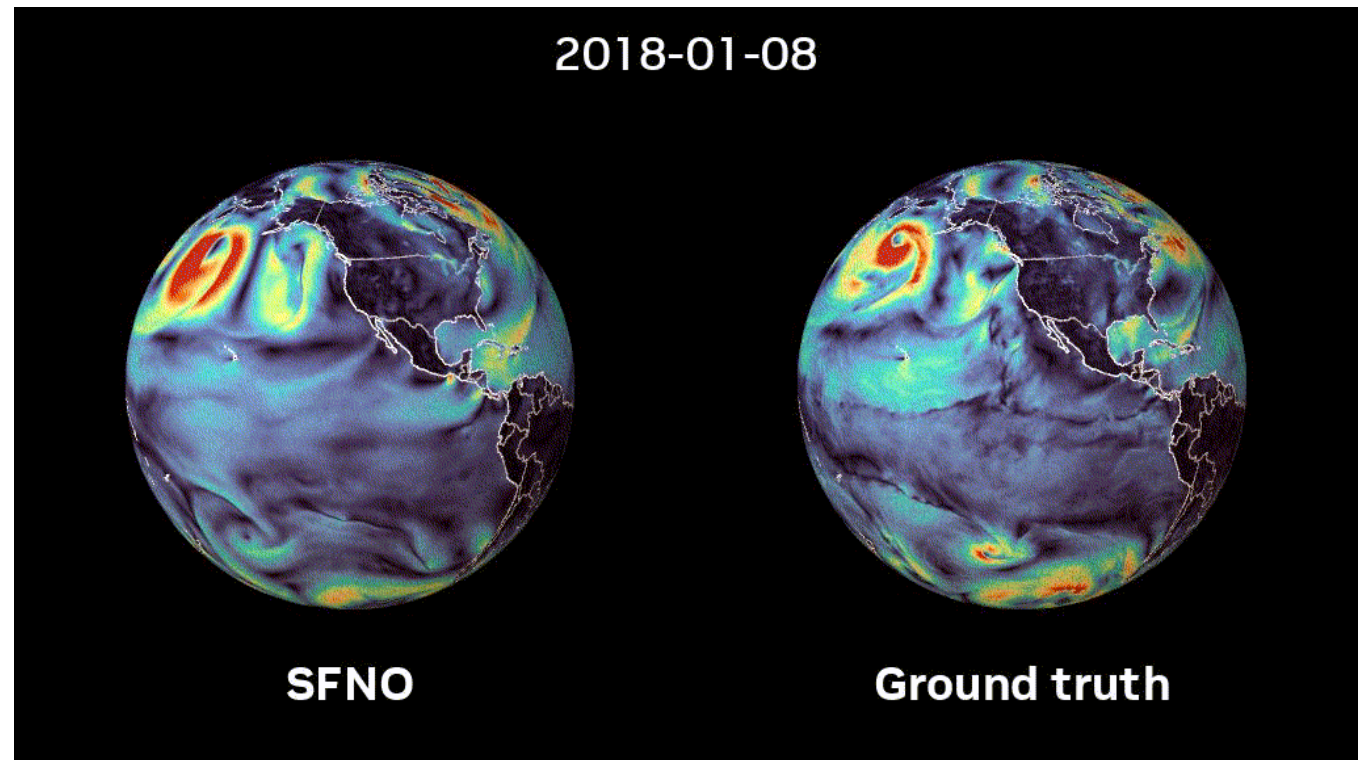



Prediction




Examples & Use Cases

- Spherical Fourier Neural Operators (SFNO)



Examples & Use Cases

- Neural operators can be combined with physics-informed loss



Takeaways

- Operator learning than solving
 - Learn a family of PDE instead of solving one instance at a time.
 - Doesn't require the explicit form, although a similar physics-loss could also be introduced (e.g. PINO)
- Learning in the function space
 - Map the values to e.g. Fourier space for continuous inputs and outputs.
 - Resolution invariant, mesh-invariant
- More accurate than other deep learning methods (discuss why?)